

▼ EECS 498-007/598-005 Assignment 1-1: PyTorch 101

Before we start, please put your name and UMID in following format

Your Answer:

Hello Artem KARPOV

Introduction

Python 3 and [PyTorch](https://cs231n.github.io/python-numpy-tutorial/) will be used throughout the semester, so it is important to be familiar with them. This material in this notebook draws from <http://cs231n.github.io/python-numpy-tutorial/> and <https://github.com/kuleshov/cs228-material/blob/master/tutorials/python/cs228-python-tutorial.ipynb>. This material focuses mainly on PyTorch.

This notebook will walk you through many of the important features of PyTorch that you will need to use throughout the semester. In some cells you will see code blocks that look like this:

```
#####  
# TODO: Write the equation for a line  
#####  
pass  
#####  
#                               END OF YOUR CODE                               #  
#####
```

You should replace the `pass` statement with your own code and leave the blocks intact, like this:

```
#####
```

```
#####  
y = m * x + b  
#####  
#                               END OF YOUR CODE                               #  
#####
```

When completing the notebook, please adhere to the following rules:

- Do not write or modify any code outside of code blocks
- Do not add or delete any cells from the notebook. You may add new cells to perform scratch work, but delete them before submitting.
- Run all cells before submitting. You will only get credit for code that has been run.

▼ Python 3

If you're unfamiliar with Python 3, here are some of the most common changes from Python 2 to look out for.

▼ Print is a function

```
print("Hello!")
```

Hello!

Without parentheses, printing will not work.

▼ Floating point division by default

```
5 / 2
```

```
2.5
```

To do integer division, we use two backslashes:

```
5 // 2
```

```
2
```

No xrange

The xrange from Python 2 is now merged into "range" for Python 3 and there is no xrange in Python 3. In Python 3, range(3) does not create a list of 3 elements as it would in Python 2, rather just creates a more memory efficient iterator.

Hence,

xrange in Python 3: Does not exist

range in Python 3: Has very similar behavior to Python 2's xrange

```
for i in range(3):  
    print(i)
```

```
0
```

```
1
```

```
2
```

```
range(3)
```

```
range(0, 3)
```

```
# If need be, can use the following to get a similar behavior to Python 2's range:  
print(list(range(3)))
```

```
[0, 1, 2]
```

Double-click (or enter) to edit

PyTorch

[PyTorch](#) is an open source machine learning framework. At its core, PyTorch provides a few key features:

- A multidimensional **Tensor** object, similar to [numpy](#) but with GPU acceleration.
- An optimized **autograd** engine for automatically computing derivatives
- A clean, modular API for building and deploying **deep learning models**

We will use PyTorch for all programming assignments throughout the semester. This notebook will focus on the **Tensor API**, as it is the main part of PyTorch that we will use for the first few assignments.

You can find more information about PyTorch by following one of the [official tutorials](#) or by [reading the documentation](#).

To use PyTorch, we first need to import the `torch` package.

We also check the version; the assignments in this course will use PyTorch version 1.1.0, since this is the default version in Google Colab.

```
import torch  
print(torch.__version__)
```

```
1.12.0+cu113
```

Tensor Basics

Creating and Accessing tensors

A `torch tensor` is a multidimensional grid of values, all of the same type, and is indexed by a tuple of nonnegative integers. The number of dimensions is the **rank** of the tensor; the **shape** of a tensor is a tuple of integers giving the size of the array along each dimension.

We can initialize `torch tensor` from nested Python lists. We can access or mutate elements of a PyTorch tensor using square brackets.

Accessing an element from a PyTorch tensor returns a PyTorch scalar; we can convert this to a Python scalar using the `.item()` method:

```
# Create a rank 1 tensor from a Python list
a = torch.tensor([1, 2, 3])
print('Here is a:')
print(a)
print('type(a): ', type(a))
print('rank of a: ', a.dim())
print('a.shape: ', a.shape)

# Access elements using square brackets
print()
print('a[0]: ', a[0])
print('type(a[0]): ', type(a[0]))
print('type(a[0].item()): ', type(a[0].item()))

# Mutate elements using square brackets
a[1] = 10
print()
print('a after mutating:')
print(a)
```

```
Here is a:
tensor([1, 2, 3])
```

```
type(a): <class 'torch.Tensor'>
rank of a: 1
a.shape: torch.Size([3])

a[0]: tensor(1)
type(a[0]): <class 'torch.Tensor'>
type(a[0].item()): <class 'int'>

a after mutating:
tensor([ 1, 10,  3])
```

The example above shows a one-dimensional tensor; we can similarly create tensors with two or more dimensions:

```
# Create a two-dimensional tensor
b = torch.tensor([[1, 2, 3], [4, 5, 5]])
print('Here is b:')
print(b)
print('rank of b:', b.dim())
print('b.shape: ', b.shape)

# Access elements from a multidimensional tensor
print()
print('b[0, 1]:', b[0, 1])
print('b[1, 2]:', b[1, 2])

# Mutate elements of a multidimensional tensor
b[1, 1] = 100
print()
print('b after mutating:')
print(b)
```

```
Here is b:
tensor([[1, 2, 3],
        [4, 5, 5]])
rank of b: 2
b.shape: torch.Size([2, 3])
```

Now it's your turn:

- ```

c = None
#####
TODO: Construct a tensor c filled with all zeros, initializing from nested
Python lists.
#####
c = torch.tensor([[0] * 2 for _ in range(3)])
#####
END OF YOUR CODE
#####
print('c is a tensor: ', torch.is_tensor(c))
print('Correct shape: ', c.shape == (3, 2))
print('All zeros: ', (c == 0).all().item() == 1)

#####
TODO: Set element (0, 1) of c to 10, and element (1, 0) of c to 100.
#####
Replace "pass" statement with your code
c[0, 1] = 10
c[1, 0] = 100
#####
END OF YOUR CODE
#####
print('\nAfter mutating:')

```

```
print('Correct shape: ', c.shape == (3, 2))
print('c[0, 1] correct: ', c[0, 1] == 10)
print('c[1, 0] correct: ', c[1, 0] == 100)
print('Rest of c is still zero: ', (c == 0).sum().item() == 4)
```

```
c is a tensor: True
Correct shape: True
All zeros: True
```

```
After mutating:
Correct shape: True
c[0, 1] correct: tensor(True)
c[1, 0] correct: tensor(True)
Rest of c is still zero: True
```

## Tensor constructors

PyTorch provides many convenience methods for constructing tensors; this avoids the need to use Python lists. For example:

- [torch.zeros](#): Creates a tensor of all zeros
- [torch.ones](#): Creates a tensor of all ones
- [torch.rand](#): Creates a tensor with uniform random numbers

You can find a full list of tensor creation operations [in the documentation](#).

```
Create a tensor of all zeros
a = torch.zeros(2, 3)
print('tensor of zeros:')
print(a)
```

```
Create a tensor of all ones
b = torch.ones(1, 2)
print('\ntensor of ones:')
print(b)
```



```
Create a 3x3 identity matrix
c = torch.eye(3)
print('\nidentity matrix:')
print(c)
```

```
Tensor of random values
d = torch.rand(4, 5)
print('\nrandom tensor:')
print(d)
```

```
tensor of zeros:
tensor([[0., 0., 0.],
 [0., 0., 0.]])
```

```
tensor of ones:
tensor([[1., 1.]])
```

```
identity matrix:
tensor([[1., 0., 0.],
 [0., 1., 0.],
 [0., 0., 1.]])
```

```
random tensor:
tensor([[0.0110, 0.8727, 0.4110, 0.7232, 0.6802],
 [0.6435, 0.5256, 0.3165, 0.8558, 0.7167],
 [0.5427, 0.5988, 0.0018, 0.5504, 0.0842],
 [0.4856, 0.9843, 0.9796, 0.5000, 0.2103]])
```

Your turn: Use a tensor creation function to create a tensor of shape (2, 3, 4) filled entirely with 7.

Hint: [torch.full](#)

```
e = None
#####
TODO: Create a tensor of shape (2, 3, 4) filled entirely with 7, stored in e
#####
```

```
Replace "pass" statement with your code
e = torch.full((2,3,4), 7)
#####
END OF YOUR CODE
#####
print('e is a tensor:', torch.is_tensor(e))
print('e has correct shape: ', e.shape == (2, 3, 4))
print('e is filled with sevens: ', (e == 7).all().item() == 1)
```

```
e is a tensor: True
e has correct shape: True
e is filled with sevens: True
```

## Datatypes

In the examples above, you may have noticed that some of our tensors contained floating-point values, while others contained integer values.

PyTorch provides a [large set of numeric datatypes](#) that you can use to construct tensors. PyTorch tries to guess a datatype when you create a tensor; functions that construct tensors typically have a `dtype` argument that you can use to explicitly specify a datatype.

Each tensor has a `dtype` attribute that you can use to check its data type:

```
Let torch choose the datatype
x0 = torch.tensor([1, 2]) # List of integers
x1 = torch.tensor([1., 2.]) # List of floats
x2 = torch.tensor([1., 2]) # Mixed list
print('dtype when torch chooses for us:')
print('List of integers:', x0.dtype)
print('List of floats:', x1.dtype)
print('Mixed list:', x2.dtype)

Force a particular datatype
x0 = torch.tensor([1, 2], dtype=torch.float32) # 32-bit float
```

```

y0 = torch.tensor([1, 2], dtype=torch.float32) # 32-bit float
y1 = torch.tensor([1, 2], dtype=torch.int32) # 32-bit (signed) integer
y2 = torch.tensor([1, 2], dtype=torch.int64) # 64-bit (signed) integer
print('\ndtype when we force a datatype:')
print('32-bit float: ', y0.dtype)
print('32-bit integer: ', y1.dtype)
print('64-bit integer: ', y2.dtype)

Other creation ops also take a dtype argument
z0 = torch.ones(1, 2) # Let torch choose for us
z1 = torch.ones(1, 2, dtype=torch.int16) # 16-bit (signed) integer
z2 = torch.ones(1, 2, dtype=torch.uint8) # 8-bit (unsigned) integer
print('\ntorch.ones with different dtypes')
print('default dtype:', z0.dtype)
print('16-bit integer:', z1.dtype)
print('8-bit unsigned integer:', z2.dtype)

```

```

dtype when torch chooses for us:
List of integers: torch.int64
List of floats: torch.float32
Mixed list: torch.float32

```

```

dtype when we force a datatype:
32-bit float: torch.float32
32-bit integer: torch.int32
64-bit integer: torch.int64

```

```

torch.ones with different dtypes
default dtype: torch.float32
16-bit integer: torch.int16
8-bit unsigned integer: torch.uint8

```

We can **cast** a tensor to another datatype using the `.to()` method; there are also convenience methods like `.float()` and `.long()` that cast to particular datatypes:

```

x0 = torch.eye(3, dtype=torch.int64)
x1 = x0.float() # Cast to 32-bit float

```

```

x1 = x0.float() # Cast to 32-bit float
x2 = x0.double() # Cast to 64-bit float
x3 = x0.to(torch.float32) # Alternate way to cast to 32-bit float
x4 = x0.to(torch.float64) # Alternate way to cast to 64-bit float
print('x0:', x0.dtype)
print('x1:', x1.dtype)
print('x2:', x2.dtype)
print('x3:', x3.dtype)
print('x4:', x4.dtype)

```

```

x0: torch.int64
x1: torch.float32
x2: torch.float64
x3: torch.float32
x4: torch.float64

```

PyTorch provides several ways to create a tensor with the same datatype as another tensor:

- PyTorch provides tensor constructors such as [torch.new\\_zeros\(\)](#) that create new tensors with the same shape and type as a given tensor
- Tensor objects have instance methods such as [.new\\_zeros\(\)](#) that create tensors the same type but possibly different shapes
- The tensor instance method [.to\(\)](#) can take a tensor as an argument, in which case it casts to the datatype of the argument.

```

x0 = torch.eye(3, dtype=torch.float64) # Shape (3, 3), dtype torch.float64
x1 = torch.zeros_like(x0) # Shape (3, 3), dtype torch.float64
x2 = x0.new_zeros(4, 5) # Shape (4, 5), dtype torch.float64
x3 = torch.ones(6, 7).to(x0) # Shape (6, 7), dtype torch.float64
print('x0 shape is %r, dtype is %r' % (x0.shape, x0.dtype))
print('x1 shape is %r, dtype is %r' % (x1.shape, x1.dtype))
print('x2 shape is %r, dtype is %r' % (x2.shape, x2.dtype))
print('x3 shape is %r, dtype is %r' % (x3.shape, x3.dtype))

```

```

x0 shape is torch.Size([3, 3]), dtype is torch.float64
x1 shape is torch.Size([3, 3]), dtype is torch.float64
x2 shape is torch.Size([4, 5]), dtype is torch.float64
x3 shape is torch.Size([6, 7]), dtype is torch.float64

```

```
x.shape is torch.Size([6, 1]), dtype is torch.float64
```

Your turn: Create a 64-bit floating-point tensor of shape (6,) (six-element vector) filled with evenly-spaced values between 10 and 20.

Hint: [torch.linspace](#)

```
x = None
#####
TODO: Make x contain a six-element vector of 64-bit floating-bit values,
evenly spaced between 10 and 20.
#####
Replace "pass" statement with your code
x = torch.linspace(10, 20, 6).to(torch.float64)
#####
END OF YOUR CODE
#####
print('x is a tensor: ', torch.is_tensor(x))
print('x has correct shape: ', x.shape == (6,))
print('x has correct dtype: ', x.dtype == torch.float64)
y = [10, 12, 14, 16, 18, 20]
correct_vals = all(a.item() == b for a, b in zip(x, y))
print('x has correct values: ', correct_vals)
```

```
x is a tensor: True
x has correct shape: True
x has correct dtype: True
x has correct values: True
```

Even though PyTorch provides a large number of numeric datatypes, the most commonly used datatypes are:

- `torch.float32`: Standard floating-point type; used to store learnable parameters, network activations, etc. Nearly all arithmetic is done using this type.
- `torch.int64`: Typically used to store indices
- `torch.uint8`: Typically used to store boolean values, where 0 is false and 1 is true.

```
100% | 64.0MB / 64.0MB | 0:00:00.000000 | 100% | 64.0MB / 64.0MB | 0:00:00.000000
```

- `torch.float16`: Used for mixed-precision arithmetic, usually on NVIDIA GPUs with [tensor cores](#). You won't need to worry about this datatype in this course.

Note that PyTorch version 1.2.0 introduced a new `torch.bool` datatype for holding boolean values. However for earlier versions (including 1.1.0 which we use in this course) you will see `torch.uint8` used to hold boolean values instead.

## Tensor indexing

We have already seen how to get and set individual elements of PyTorch tensors. PyTorch also provides many other ways of indexing into tensors. Getting comfortable with these different options makes it easy to modify different parts of tensors with ease.

### Slice indexing

Similar to Python lists and numpy arrays, PyTorch tensors can be **sliced** using the syntax `start:stop` or `start:stop:step`. The `stop` index is always non-inclusive: it is the first element not to be included in the slice.

Start and stop indices can be negative, in which case they count backward from the end of the tensor.

```
a = torch.tensor([0, 11, 22, 33, 44, 55, 66])
print(0, a) # (0) Original tensor
print(1, a[2:5]) # (1) Elements between index 2 and 5
print(2, a[2:]) # (2) Elements after index 2
print(3, a[:5]) # (3) Elements before index 5
print(4, a[:]) # (4) All elements
print(5, a[1:5:2]) # (5) Every second element between indices 1 and 5
print(6, a[:-1]) # (6) All but the last element
print(7, a[-4::2]) # (7) Every second element, starting from the fourth-last
```

```
0 tensor([0, 11, 22, 33, 44, 55, 66])
1 tensor([22, 33, 44])
```

```

2 tensor([22, 33, 44, 55, 66])
3 tensor([0, 11, 22, 33, 44])
4 tensor([0, 11, 22, 33, 44, 55, 66])
5 tensor([11, 33])
6 tensor([0, 11, 22, 33, 44, 55])
7 tensor([33, 55])

```

For multidimensional tensors, you can provide a slice or integer for each dimension of the tensor in order to extract different types of subtensors:

```

Create the following rank 2 tensor with shape (3, 4)
[[1 2 3 4]
[5 6 7 8]
[9 10 11 12]]
a = torch.tensor([[1,2,3,4], [5,6,7,8], [9,10,11,12]])
print('Original tensor:')
print(a)
print('shape: ', a.shape)

Get row 1, and all columns.
print('\nSingle row:')
print(a[1, :])
print(a[1]) # Gives the same result; we can omit : for trailing dimensions
print('shape: ', a[1].shape)

print('\nSingle column:')
print(a[:, 1])
print('shape: ', a[:, 1].shape)

Get the first two rows and the last three columns
print('\nFirst two rows, last two columns:')
print(a[:2, -3:])
print('shape: ', a[:2, -3:].shape)

Get every other row, and columns at index 1 and 2
print('\nEvery other row, middle columns:')

```

```
print('Every other row, middle columns: ')
print(a[::2, 1:3])
print('shape: ', a[::2, 1:3].shape)
```

```
Original tensor:
tensor([[1, 2, 3, 4],
 [5, 6, 7, 8],
 [9, 10, 11, 12]])
shape: torch.Size([3, 4])
```

```
Single row:
tensor([5, 6, 7, 8])
tensor([5, 6, 7, 8])
shape: torch.Size([4])
```

```
Single column:
tensor([2, 6, 10])
shape: torch.Size([3])
```

```
First two rows, last two columns:
tensor([[2, 3, 4],
 [6, 7, 8]])
shape: torch.Size([2, 3])
```

```
Every other row, middle columns:
tensor([[2, 3],
 [10, 11]])
shape: torch.Size([2, 2])
```

There are two common ways to access a single row or column of a tensor: using an integer will reduce the rank by one, and using a length-one slice will keep the same rank. Note that this is different behavior from MATLAB.

```
Create the following rank 2 tensor with shape (3, 4)
a = torch.tensor([[1,2,3,4], [5,6,7,8], [9,10,11,12]])
print('Original tensor')
print(a)
```

```
tensor([[1, 2, 3, 4],
 [5, 6, 7, 8],
 [9, 10, 11, 12]]) # Rank 2 view of the second row of a
```



```

row_r1 = a[1, :] # Rank 1 view of the second row of a
row_r2 = a[1:2, :] # Rank 2 view of the second row of a
print('\nTwo ways of accessing a single row:')
print(row_r1, row_r1.shape)
print(row_r2, row_r2.shape)

We can make the same distinction when accessing columns::
col_r1 = a[:, 1]
col_r2 = a[:, 1:2]
print('\nTwo ways of accessing a single column:')
print(col_r1, col_r1.shape)
print(col_r2, col_r2.shape)

```

```

Original tensor
tensor([[1, 2, 3, 4],
 [5, 6, 7, 8],
 [9, 10, 11, 12]])

```

```

Two ways of accessing a single row:
tensor([5, 6, 7, 8]) torch.Size([4])
tensor([[5, 6, 7, 8]]) torch.Size([1, 4])

```

```

Two ways of accessing a single column:
tensor([2, 6, 10]) torch.Size([3])
tensor([[2],
 [6],
 [10]]) torch.Size([3, 1])

```

Slicing a tensor returns a **view** into the same data, so modifying it will also modify the original tensor. To avoid this, you can use the `clone()` method to make a copy of a tensor.

```

Create a tensor, a slice, and a clone of a slice
a = torch.tensor([[1, 2, 3, 4], [5, 6, 7, 8]])
b = a[0, 1:]
c = a[0, 1:].clone()
print('Before mutating:')
print(a)

```

```

print(a)
print(b)
print(c)

a[0, 1] = 20 # a[0, 1] and b[0] point to the same element
b[1] = 30 # b[1] and a[0, 2] point to the same element
c[2] = 40 # c is a clone, so it has its own data
print('\nAfter mutating:')
print(a)
print(b)
print(c)

print(a.storage().data_ptr() == c.storage().data_ptr())

```

```

Before mutating:
tensor([[1, 2, 3, 4],
 [5, 6, 7, 8]])
tensor([2, 3, 4])
tensor([2, 3, 4])

After mutating:
tensor([[1, 20, 30, 4],
 [5, 6, 7, 8]])
tensor([20, 30, 4])
tensor([2, 3, 40])
False

```

Your turn: practice indexing tensors with slices

```

We will use this helper function to check your results
def check(orig, actual, expected):
 expected = torch.tensor(expected)
 same_elements = (actual == expected).all().item() == 1
 same_storage = (orig.storage().data_ptr() == actual.storage().data_ptr())
 return same_elements and same_storage

```

```

Create the following rank 2 tensor of shape (3, 5)
[[1 2 3 4 5]
[6 7 8 9 10]
[11 12 13 14 15]]
a = torch.tensor([[1, 2, 3, 4, 5], [6, 7, 8, 8, 10], [11, 12, 13, 14, 15]])

b, c, d, e = None, None, None, None
#####
TODO: Extract the last row of a, and store it in b; it should have rank 1.
#####
Replace "pass" statement with your code
b = a[-1,:]
#####
END OF YOUR CODE
#####
print('b correct:', check(a, b, [11, 12, 13, 14, 15]))

#####
TODO: Extract the third col of a, and store it in c; it should have rank 2
#####
Replace "pass" statement with your code
c = a[:,2:3]
#####
END OF YOUR CODE
#####
print('c correct:', check(a, c, [[3], [8], [13]]))

#####
TODO: Use slicing to extract the first two rows and first three columns
from a; store the result into d.
#####
Replace "pass" statement with your code
d = a[0:2, 0:3]
#####
END OF YOUR CODE
#####
print('d correct:', check(a, d, [[1, 2, 3], [6, 7, 8]]))

```

```
#####
TODO: Use slicing to extract a subtensor of a consisting of rows 0 and 2
and columns 1 and 4.
#####
Replace "pass" statement with your code
e = a[0:3:2, 1:5:3]
#####
END OF YOUR CODE
#####
print('e correct:', check(a, e, [[2, 5], [12, 15]]))
```

```
b correct: True
c correct: True
d correct: True
e correct: True
```

So far we have used slicing to **access** subtensors; we can also use slicing to **modify** subtensors by writing assignment expressions where the left-hand side is a slice expression, and the right-hand side is a constant or a tensor of the correct shape:

```
a = torch.zeros(2, 4, dtype=torch.int64)
a[:, :2] = 1
a[:, 2:] = torch.tensor([[2, 3], [4, 5]])
print(a)
```

```
tensor([[1, 1, 2, 3],
 [1, 1, 4, 5]])
```

Your turn: use slicing assignment to modify a tensor:

```
x = torch.zeros(4, 6, dtype=torch.int64)
#####
TODO: Use slicing to modify the tensor x so it has the following contents:
```

```

TODO: Use slicing to modify the tensor x so it has the following contents.
[[1, 0, 2, 2, 2, 2],
[0, 1, 2, 2, 2, 2],
[3, 4, 3, 4, 5, 5],
[3, 4, 3, 4, 5, 5]]
This can be achieved using five slicing assignment operations.
#####
Replace "pass" statement with your code
x[:2, :2] = torch.eye(2)
x[0:2, 2:] = 2
x[2:, 0:3:2] = 3
x[2:, 1:4:2] = 4
x[-2:, -2:] = 5
#####
END OF YOUR CODE
#####

expected = [
 [1, 0, 2, 2, 2, 2],
 [0, 1, 2, 2, 2, 2],
 [3, 4, 3, 4, 5, 5],
 [3, 4, 3, 4, 5, 5],
]
print('correct:', x.tolist() == expected)

```

```
correct: True
```

## Integer tensor indexing

When you index into torch tensor using slicing, the resulting tensor view will always be a subarray of the original tensor. This is powerful, but can be restrictive.

We can also use **index arrays** to index tensors; this lets us construct new tensors with a lot more flexibility than using slices.

As an example, we can use index arrays to reorder the rows or columns of a tensor:

```

Create the following rank 2 tensor with shape (3, 4)
[[1 2 3 4]
[5 6 7 8]
[9 10 11 12]]
a = torch.tensor([[1, 2, 3, 4], [5, 6, 7, 8], [9, 10, 11, 12]])
print('Original tensor:')
print(a)

Create a new tensor of shape (5, 4) by reordering rows from a:
- First two rows same as the first row of a
- Third row is the same as the last row of a
- Fourth and fifth rows are the same as the second row from a
idx = [0, 0, 2, 1, 1] # index arrays can be Python lists of integers
print('\nReordered rows:')
print(a[idx])

Create a new tensor of shape (3, 4) by reversing the columns from a
idx = torch.tensor([3, 2, 1, 0]) # Index arrays can be int64 torch tensors
print('\nReordered columns:')
print(a[:, idx])

```

```

Original tensor:
tensor([[1, 2, 3, 4],
 [5, 6, 7, 8],
 [9, 10, 11, 12]])

```

```

Reordered rows:
tensor([[1, 2, 3, 4],
 [1, 2, 3, 4],
 [9, 10, 11, 12],
 [5, 6, 7, 8],
 [5, 6, 7, 8]])

```

```

Reordered columns:
tensor([[4, 3, 2, 1],
 [8, 7, 6, 5],
 [12, 11, 10, 9]])

```

```
[12, 11, 10, 9]]
```

More generally, given index arrays `idx0` and `idx1` with `N` elements each, `a[idx0, idx1]` is equivalent to:

```
torch.tensor([
 a[idx0[0], idx1[0]],
 a[idx0[1], idx1[1]],
 ...,
 a[idx0[N - 1], idx1[N - 1]]
])
```

(A similar pattern extends to tensors with more than two dimensions)

We can for example use this to get or set the diagonal of a tensor:

```
a = torch.tensor([[1, 2, 3], [4, 5, 6], [7, 8, 9]])
print('Original tensor:')
print(a)

idx = [0, 1, 2]
print('\nGet the diagonal:')
print(a[idx, idx])

Modify the diagonal
a[idx, idx] = torch.tensor([11, 22, 33])
print('\nAfter setting the diagonal:')
print(a)
```

```
Original tensor:
tensor([[1, 2, 3],
 [4, 5, 6],
 [7, 8, 9]])
```

```
Get the diagonal:
tensor([1, 5, 9])
```

```
tensor([1, 5, 9])
```

After setting the diagonal:

```
tensor([[11, 2, 3],
 [4, 22, 6],
 [7, 8, 33]])
```

One useful trick with integer array indexing is selecting or mutating one element from each row or column of a matrix:

```
Create a new tensor from which we will select elements
a = torch.tensor([[1, 2, 3], [4, 5, 6], [7, 8, 9], [10, 11, 12]])
print('Original tensor:')
print(a)
```

```
Take one element from each row of a:
from row 0, take element 1;
from row 1, take element 2;
from row 2, take element 1;
from row 3, take element 0
idx0 = torch.arange(a.shape[0]) # Quick way to build [0, 1, 2, 3]
idx1 = torch.tensor([1, 2, 1, 0])
print('\nSelect one element from each row:')
print(a[idx0, idx1])
```

```
Now set each of those elements to zero
a[idx0, idx1] = 0
print('\nAfter modifying one element from each row:')
print(a)
```

```
Original tensor:
tensor([[1, 2, 3],
 [4, 5, 6],
 [7, 8, 9],
 [10, 11, 12]])
```

```
Select one element from each row:
tensor([2, 6, 8, 10])
```



After modifying one element from each row:

```
tensor([[1, 0, 3],
 [4, 5, 0],
 [7, 0, 9],
 [0, 11, 12]])
```

Your turn: practice with integer array indexing

```
Build a tensor of shape (4, 3):
[[1, 2, 3],
[4, 5, 6],
[7, 8, 9],
[10, 11, 12]]
a = torch.tensor([[1, 2, 3], [4, 5, 6], [7, 8, 9], [10, 11, 12]])
print('Here is a:')
print(a)

b, c, d = None, None, None
#####
TODO: Use integer array indexing to create a tensor of shape (4, 4) where:
- The first two columns are the same as the first column of a
- The next column is the same as the third column of a
- The last column is the same as the second column of a
Store the resulting tensor in b.
#####
Replace "pass" statement with your code
inx = torch.tensor([0, 0, 2, 1])
b = a[:, inx]
#####
END OF YOUR CODE
#####
print('\nHere is b:')
print(b)
expected = [[1, 1, 3, 2], [4, 4, 6, 5], [7, 7, 9, 8], [10, 10, 12, 11]]
print('b correct:', b.tolist() == expected)
```

```

print('b correct:', b.tolist() == expected)

#####
TODO: Use integer array indexing to create a new tensor which is the same
as a, but has its rows reversed. Store the result in c.
#####
Replace "pass" statement with your code
inx = torch.tensor([3, 2, 1, 0])
c = a[inx, :]
#####
END OF YOUR CODE
#####
print('\nHere is c:')
print(c)
expected = [[10, 11, 12], [7, 8, 9], [4, 5, 6], [1, 2, 3]]
print('c correct:', c.tolist() == expected)

#####
TODO: Use integer array indexing to create a new tensor by selecting one
element from each column of a:
- From the first column, take the second element.
- From the second column, take the first element.
- From the third column, take the fourth element.
Store the result in d.
#####
Replace "pass" statement with your code
cinx = torch.tensor([0,1,2])
rinx = torch.tensor([1, 0, 3])
d = a[rinx, cinx]
#####
END OF YOUR CODE
#####
print('\nHere is d:')
print(d)
expected = [4, 2, 12]
print('d correct:', d.tolist() == expected)

```

Here is a:

```
here is a:
tensor([[1, 2, 3],
 [4, 5, 6],
 [7, 8, 9],
 [10, 11, 12]])
```

```
Here is b:
tensor([[1, 1, 3, 2],
 [4, 4, 6, 5],
 [7, 7, 9, 8],
 [10, 10, 12, 11]])
b correct: True
```

```
Here is c:
tensor([[10, 11, 12],
 [7, 8, 9],
 [4, 5, 6],
 [1, 2, 3]])
c correct: True
```

```
Here is d:
tensor([4, 2, 12])
d correct: True
```

## Boolean tensor indexing

Boolean tensor indexing lets you pick out arbitrary elements of a tensor according to a boolean mask. Frequently this type of indexing is used to select or modify the elements of a tensor that satisfy some condition.

In PyTorch, we use tensors of dtype `torch.uint8` to hold boolean masks; 0 means false and 1 means true.

(PyTorch version 1.2.0 introduces a `torch.bool` type for tensors, which is used instead of `torch.uint8` for boolean masks. However in this class we are using PyTorch 1.1.0)

```
a = torch.tensor([[1,2], [3, 4], [5, 6]])
print('Original tensor:')
```

```

print('Original tensor:',
print(a)

Find the elements of a that are bigger than 3. The mask has the same shape as
a, where each element of mask tells whether the corresponding element of a
is greater than three.
mask = (a > 3)
print('\nMask tensor:')
print(mask)

We can use the mask to construct a rank-1 tensor containing the elements of a
that are selected by the mask
print('\nSelecting elements with the mask:')
print(a[mask])

We can also use boolean masks to modify tensors; for example this sets all
elements <= 3 to zero:
a[a <= 3] = 0
print('\nAfter modifying with a mask:')
print(a)

```

```

Original tensor:
tensor([[1, 2],
 [3, 4],
 [5, 6]])

```

```

Mask tensor:
tensor([[False, False],
 [False, True],
 [True, True]])

```

```

Selecting elements with the mask:
tensor([4, 5, 6])

```

```

After modifying with a mask:
tensor([[0, 0],
 [0, 4],
 [5, 6]])

```

Your turn: practice with boolean masks by implementing the following function:

```
def num_negative(x):
 """
 Return the number of negative values in the tensor x

 Inputs:
 - x: A tensor of any shape

 Returns:
 - num_neg: Number of negative values in x
 """
 num_neg = 0
 #####
 # TODO: Use boolean masks to count the number of negative elements in x. #
 #####
 # Replace "pass" statement with your code
 num_neg = (x < 0).sum()
 #####
 # END OF YOUR CODE #
 #####
 return num_neg

Make a few test cases
torch.manual_seed(598)
x0 = torch.tensor([-1, -1, 0], [0, 1, 2], [3, 4, 5])
x1 = torch.tensor([0, 1, 2, 3])
x2 = torch.randn(100, 100)
assert num_negative(x0) == 2
assert num_negative(x1) == 0
assert num_negative(x2) == 4984
print('num_negative seems to be correct!')
```

num\_negative seems to be correct!

Now implement a function that creates a matrix of **one-hot vectors** from a list of Python integers.

A one-hot vector for an integer  $n$  is a vector that has a one in its  $n$ th slot, and zeros in all other slots. One-hot vectors are commonly used to represent categorical variables in machine learning models.

For example, given a list `[1, 4, 3, 2]` of integers, your function should produce the tensor:

```
[[0 1 0 0 0],
 [0 0 0 0 1],
 [0 0 0 1 0],
 [0 0 1 0 0]]
```

Here the first row corresponds to the first element of the list: it has a one at index 1, and zeros at all other indices. The second row corresponds to the second element of the list: it has a one at index 4, and zeros at all other indices. The other rows follow the same pattern.

```
def make_one_hot(x):
 """
 Construct a tensor of one-hot-vectors from a list of Python integers.

 Input:
 - x: A list of N ints

 Returns:
 - y: A tensor of shape (N, C) where C = 1 + max(x) is one more than the max
 value in x. The nth row of y is a one-hot-vector representation of x[n];
 In other words, if x[n] = c then y[n, c] = 1; all other elements of y are
 zeros.
 """
 y = None
 #####
 # TODO: Complete the implementation of this function. #
 #####
 # Replace "pass" statement with your code
 # y = torch.zeros(x, max(x) + 1)
```

```

n, c = len(x), max(x) + 1
y = torch.zeros(n, c)
y[torch.arange(n), torch.tensor(x)] = 1
#####
END OF YOUR CODE
#####
return y

```

Now check your implementation:

```

def check_one_hot(x, y):
 C = y.shape[1]
 for i, n in enumerate(x):
 if n >= C: return False
 for j in range(C):
 expected = 1.0 if j == n else 0.0
 if y[i, j].item() != expected: return False
 return True

x0 = [1, 4, 3, 2]
y0 = make_one_hot(x0)
print('Here is y0:')
print(y0)
assert check_one_hot(x0, y0), 'y0 is wrong'

x1 = [1, 3, 5, 7, 6, 2]
y1 = make_one_hot(x1)
print('\nHere is y1:')
print(y1)
assert check_one_hot(x1, y1), 'y1 is wrong'

print('all checks pass!')

```

```

Here is y0:
tensor([[0., 1., 0., 0., 0.],
 [0., 0., 0., 0., 1.],

```

```
[0., 0., 0., 1., 0.],
[0., 0., 1., 0., 0.]])
```

Here is y1:

```
tensor([[0., 1., 0., 0., 0., 0., 0., 0.],
 [0., 0., 0., 1., 0., 0., 0., 0.],
 [0., 0., 0., 0., 0., 1., 0., 0.],
 [0., 0., 0., 0., 0., 0., 0., 1.],
 [0., 0., 0., 0., 0., 0., 1., 0.],
 [0., 0., 1., 0., 0., 0., 0., 0.]])
```

all checks pass!

## Reshaping operations

### View

PyTorch provides many ways to manipulate the shapes of tensors. The simplest example is [.view\(\)](#): This returns a new tensor with the same number of elements as its input, but with a different shape.

We can use `.view()` to flatten matrices into vectors, and to convert rank-1 vectors into rank-2 row or column matrices:

```
x0 = torch.tensor([[1, 2, 3, 4], [5, 6, 7, 8]])
print('Original tensor:')
print(x0)
print('shape:', x0.shape)

Flatten x0 into a rank 1 vector of shape (8,)
x1 = x0.view(8)
print('\nFlattened tensor:')
print(x1)
print('shape:', x1.shape)
```

```
"""
 Copyright (c) 2019-present, Facebook, Inc.
 All rights reserved.
 This source code is licensed under the Creative Commons
 Attribution 4.0 International License (CC BY 4.0).
"""
```



```
Convert x1 to a rank 2 "row vector" of shape (1, 8)
x2 = x1.view(1, 8)
print('\nRow vector:')
print(x2)
print('shape:', x2.shape)

Convert x1 to a rank 2 "column vector" of shape (8, 1)
x3 = x1.view(8, 1)
print('\nColumn vector:')
print(x3)
print('shape:', x3.shape)

Convert x1 to a rank 3 tensor of shape (2, 2, 2):
x4 = x1.view(2, 2, 2)
print('\nRank 3 tensor:')
print(x4)
print('shape:', x4.shape)
```

Original tensor:

```
tensor([[1, 2, 3, 4],
 [5, 6, 7, 8]])
shape: torch.Size([2, 4])
```

Flattened tensor:

```
tensor([1, 2, 3, 4, 5, 6, 7, 8])
shape: torch.Size([8])
```

Row vector:

```
tensor([[1, 2, 3, 4, 5, 6, 7, 8]])
shape: torch.Size([1, 8])
```

Column vector:

```
tensor([[1],
 [2],
 [3],
 [4],
 [5],
 [6],
 [7],
```

```

 [8]])
shape: torch.Size([8, 1])

Rank 3 tensor:
tensor([[[1, 2],
 [3, 4]],

 [[5, 6],
 [7, 8]]])
shape: torch.Size([2, 2, 2])

```

As a convenience, calls to `.view()` may include a single `-1` argument; this puts enough elements on that dimension so that the output has the same shape as the input. This makes it easy to write some reshape operations in a way that is agnostic to the shape of the tensor:

```

We can reuse these functions for tensors of different shapes
def flatten(x):
 return x.view(-1)

def make_row_vec(x):
 return x.view(1, -1)

x0 = torch.tensor([[1, 2, 3], [4, 5, 6]])
x0_flat = flatten(x0)
x0_row = make_row_vec(x0)
print('x0:')
print(x0)
print('x0_flat:')
print(x0_flat)
print('x0_row:')
print(x0_row)

x1 = torch.tensor([[1, 2], [3, 4]])
x1_flat = flatten(x1)
x1_row = make_row_vec(x1)
print('\nx1:')
print(x1)

```

```

print(x1,
print('x1_flat:')
print(x1_flat)
print('x1_row:')
print(x1_row)

```

```

x0:
tensor([[1, 2, 3],
 [4, 5, 6]])
x0_flat:
tensor([1, 2, 3, 4, 5, 6])
x0_row:
tensor([[1, 2, 3, 4, 5, 6]])

x1:
tensor([[1, 2],
 [3, 4]])
x1_flat:
tensor([1, 2, 3, 4])
x1_row:
tensor([[1, 2, 3, 4]])

```

As its name implies, a tensor returned by `.view()` shares the same data as the input, so changes to one will affect the other and vice-versa:

```

x = torch.tensor([[1, 2, 3], [4, 5, 6]])
x_flat = x.view(-1)
print('x before modifying:')
print(x)
print('x_flat before modifying:')
print(x_flat)

x[0, 0] = 10 # x[0, 0] and x_flat[0] point to the same data
x_flat[1] = 20 # x_flat[1] and x[0, 1] point to the same data

print('\nx after modifying:')
print(x)
print('x_flat after modifying:')

```

```
print(x_flat)
```

```
x before modifying:
tensor([[1, 2, 3],
 [4, 5, 6]])
x_flat before modifying:
tensor([1, 2, 3, 4, 5, 6])

x after modifying:
tensor([[10, 20, 3],
 [4, 5, 6]])
x_flat after modifying:
tensor([10, 20, 3, 4, 5, 6])
```

## Swapping axes

Another common reshape operation you might want to perform is transposing a matrix. You might be surprised if you try to transpose a matrix with `.view()`: The `view()` function takes elements in row-major order, so **you cannot transpose matrices with `.view()`**.

In general, you should only use `.view()` to add new dimensions to a tensor, or to collapse adjacent dimensions of a tensor.

For other types of reshape operations, you usually need to use a function that can swap axes of a tensor. The simplest such function is `.t()`, specifically for transposing matrices. It is available both as a [function in the torch module](#), and as a [tensor instance method](#):

```
x = torch.tensor([[1, 2, 3], [4, 5, 6]])
print('Original matrix:')
print(x)
print('\nTransposing with view DOES NOT WORK!')
print(x.view(3, 2))
print('\nTransposed matrix:')
print(torch.t(x))
print(x.t())
```

```
Original matrix:
tensor([[1, 2, 3],
```

```
tensor([[1, 2, 3],
 [4, 5, 6]])
```

Transposing with view DOES NOT WORK!

```
tensor([[1, 2],
 [3, 4],
 [5, 6]])
```

Transposed matrix:

```
tensor([[1, 4],
 [2, 5],
 [3, 6]])
tensor([[1, 4],
 [2, 5],
 [3, 6]])
```

For tensors with more than two dimensions, we can use the function [torch.transpose](#) to swap arbitrary dimensions, or the [.permute](#) method to arbitrarily permute dimensions:

```
Create a tensor of shape (2, 3, 4)
x0 = torch.tensor([
 [[1, 2, 3, 4],
 [5, 6, 7, 8],
 [9, 10, 11, 12]],
 [[13, 14, 15, 16],
 [17, 18, 19, 20],
 [21, 22, 23, 24]]])
print('Original tensor:')
print(x0)
print('shape:', x0.shape)

Swap axes 1 and 2; shape is (2, 4, 3)
x1 = x0.transpose(1, 2)
print('\nSwap axes 1 and 2:')
print(x1)
print(x1.shape)
```

```
Permute axes; the argument (1, 2, 0) means:
- Make the old dimension 1 appear at dimension 0;
- Make the old dimension 2 appear at dimension 1;
- Make the old dimension 0 appear at dimension 2
This results in a tensor of shape (3, 4, 2)
x2 = x0.permute(1, 2, 0)
print('\nPermute axes')
print(x2)
print('shape:', x2.shape)
```

Original tensor:

```
tensor([[[1, 2, 3, 4],
 [5, 6, 7, 8],
 [9, 10, 11, 12]],
```

```
 [[13, 14, 15, 16],
 [17, 18, 19, 20],
 [21, 22, 23, 24]])])
```

shape: torch.Size([2, 3, 4])

Swap axes 1 and 2:

```
tensor([[[1, 5, 9],
 [2, 6, 10],
 [3, 7, 11],
 [4, 8, 12]],
```

```
 [[13, 17, 21],
 [14, 18, 22],
 [15, 19, 23],
 [16, 20, 24]])])
```

torch.Size([2, 4, 3])

Permute axes

```
tensor([[[1, 13],
 [2, 14],
 [3, 15],
 [4, 16]],
```

```
 [[5, 17],
```

```

 [6, 18],
 [7, 19],
 [8, 20]],
 [[9, 21],
 [10, 22],
 [11, 23],
 [12, 24]]])
shape: torch.Size([3, 4, 2])

```

## Contiguous tensors

Some combinations of reshaping operations will fail with cryptic errors. The exact reasons for this have to do with the way that tensors and views of tensors are implemented, and are beyond the scope of this assignment. However if you're curious, [this blog post by Edward Yang](#) gives a clear explanation of the problem.

What you need to know is that you can typically overcome these sorts of errors by either by calling `.contiguous()` before `.view()`, or by using `.reshape()` instead of `.view()`.

```

x0 = torch.randn(2, 3, 4)

try:
 # This sequence of reshape operations will crash
 x1 = x0.transpose(1, 2).view(8, 3)
except RuntimeError as e:
 print(type(e), e)

We can solve the problem using either .contiguous() or .reshape()
x1 = x0.transpose(1, 2).contiguous().view(8, 3)
x2 = x0.transpose(1, 2).reshape(8, 3)
print('x1 shape: ', x1.shape)
print('x2 shape: ', x2.shape)

```

```

<class 'RuntimeError'> view size is not compatible with input tensor's size and stride (at least one dimension s

```

```
RuntimeError: view size is not compatible with input tensor's size and stride (at least one dimension spans across two nonzero strides)
x1 shape: torch.Size([8, 3])
x2 shape: torch.Size([8, 3])
```

## Your turn

Given the 1-dimensional input tensor `x0` containing the numbers 0 through 23 in order, apply a sequence of reshaping operations to `x0` to create the following tensor:

```
x1 = tensor([[0, 1, 2, 3, 12, 13, 14, 15],
 [4, 5, 6, 7, 16, 17, 18, 19],
 [8, 9, 10, 11, 20, 21, 22, 23]])
```

Hint: You will need to create an intermediate tensor of rank 3

```
x0 = torch.arange(24)
print('Here is x0:')
print(x0)

x1 = None
#####
TODO: Use reshape operations to create x1 from x0
#####
Replace "pass" statement with your code
x1 = torch.zeros(3, 8)
x0 = x0.view(6, 4)
x1[:, 0:4] = x0[0:3, :]
x1[:, 4:] = x0[3:, :]
#####
END OF YOUR CODE
#####
print('\nHere is x1:')
print(x1)
```



```
print(x1)

expected = [
 [0, 1, 2, 3, 12, 13, 14, 15],
 [4, 5, 6, 7, 16, 17, 18, 19],
 [8, 9, 10, 11, 20, 21, 22, 23]]
print('Correct:', x1.tolist() == expected)
```

```
--NORMAL--
```

<Esc>

```
Here is x0:
tensor([0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17,
 18, 19, 20, 21, 22, 23])

Here is x1:
tensor([[0., 1., 2., 3., 12., 13., 14., 15.],
 [4., 5., 6., 7., 16., 17., 18., 19.],
 [8., 9., 10., 11., 20., 21., 22., 23.]])
Correct: True
```

## Tensor operations

### Elementwise operations

Basic mathematical functions operate elementwise on tensors, and are available as operator overloads, as functions in the `torch` module, and as instance methods on torch objects; all produce the same results:

```
x = torch.tensor([[1, 2, 3, 4]], dtype=torch.float32)
y = torch.tensor([[5, 6, 7, 8]], dtype=torch.float32)

Elementwise sum; all give the same result
print('Elementwise sum:')
print(x + y)
print(torch.add(x, y))
```

```

print(torch.add(x, y))
print(x.add(y))

Elementwise difference
print('\nElementwise difference:')
print(x - y)
print(torch.sub(x, y))
print(x.sub(y))

Elementwise product
print('\nElementwise product:')
print(x * y)
print(torch.mul(x, y))
print(x.mul(y))

Elementwise division
print('\nElementwise division')
print(x / y)
print(torch.div(x, y))
print(x.div(y))

Elementwise power
print('\nElementwise power')
print(x ** y)
print(torch.pow(x, y))
print(x.pow(y))

```

```

Elementwise sum:
tensor([[6., 8., 10., 12.]])
tensor([[6., 8., 10., 12.]])
tensor([[6., 8., 10., 12.]])

```

```

Elementwise difference:
tensor([[-4., -4., -4., -4.]])
tensor([[-4., -4., -4., -4.]])
tensor([[-4., -4., -4., -4.]])

```

```

Elementwise product:
tensor([[36., 64., 100., 144.]])

```

```
tensor([[5., 12., 21., 32.]])
tensor([[5., 12., 21., 32.]])
tensor([[5., 12., 21., 32.]])
```

Elementwise division

```
tensor([[0.2000, 0.3333, 0.4286, 0.5000]])
tensor([[0.2000, 0.3333, 0.4286, 0.5000]])
tensor([[0.2000, 0.3333, 0.4286, 0.5000]])
```

Elementwise power

```
tensor([[1.0000e+00, 6.4000e+01, 2.1870e+03, 6.5536e+04]])
tensor([[1.0000e+00, 6.4000e+01, 2.1870e+03, 6.5536e+04]])
tensor([[1.0000e+00, 6.4000e+01, 2.1870e+03, 6.5536e+04]])
```

Torch also provides many standard mathematical functions; these are available both as functions in the `torch` module and as instance methods on tensors:

You can find a full list of all available mathematical functions [in the documentation](#); many functions in the `torch` module have corresponding instance methods [on tensor objects](#).

```
x = torch.tensor([1, 2, 3, 4], dtype=torch.float32)
```

```
print('Square root:')
print(torch.sqrt(x))
print(x.sqrt())
```

```
print('\nTrig functions:')
print(torch.sin(x))
print(x.sin())
print(torch.cos(x))
print(x.cos())
```

Square root:

```
tensor([[1.0000, 1.4142, 1.7321, 2.0000]])
tensor([[1.0000, 1.4142, 1.7321, 2.0000]])
```

Trig functions:

```

 trig_functions.
tensor([[0.8415, 0.9093, 0.1411, -0.7568]])
tensor([[0.8415, 0.9093, 0.1411, -0.7568]])
tensor([[0.5403, -0.4161, -0.9900, -0.6536]])
tensor([[0.5403, -0.4161, -0.9900, -0.6536]])

```

## Reduction operations

So far we've seen basic arithmetic operations on tensors that operate elementwise. We may sometimes want to perform operations that aggregate over part or all of a tensor, such as a summation; these are called **reduction** operations.

Like the elementwise operations above, most reduction operations are available both as functions in the `torch` module and as instance methods on `tensor` objects.

The simplest reduction operation is summation. We can use the `.sum()` function to reduce either an entire tensor, or to reduce along only one dimension of the tensor using the `dim` argument:

```

x = torch.tensor([[1, 2, 3],
 [4, 5, 6]], dtype=torch.float32)
print('Original tensor:')
print(x)

print('\nSum over entire tensor:')
print(torch.sum(x))
print(x.sum())

We can sum over each row:
print('\nSum of each row:')
print(torch.sum(x, dim=0))
print(x.sum(dim=0))

Sum over each column:
print('\nSum of each column:')
print(torch.sum(x, dim=1))

```

```
print(x.sum(dim=1))
```

```
Original tensor:
tensor([[1., 2., 3.],
 [4., 5., 6.]])
```

```
Sum over entire tensor:
tensor(21.)
tensor(21.)
```

```
Sum of each row:
tensor([5., 7., 9.])
tensor([5., 7., 9.])
```

```
Sum of each column:
tensor([6., 15.])
tensor([6., 15.])
```

Other useful reduction operations include [mean](#), [min](#), and [max](#). You can find a full list of all available reduction operations [in the documentation](#).

Some reduction operations return more than one value; for example `min` returns both the minimum value over the specified dimension, as well as the index where the minimum value occurs:

```
x = torch.tensor([[2, 4, 3, 5], [3, 3, 5, 2]], dtype=torch.float32)
print('Original tensor:')
print(x, x.shape)

Finding the overall minimum only returns a single value
print('\nOverall minimum: ', x.min())

Compute the minimum along each column; we get both the value and location:
The minimum of the first column is 2, and it appears at index 0;
the minimum of the second column is 3 and it appears at index 1; etc
col_min_vals, col_min_idxes = x.min(dim=0)
print('\nMinimum along each column:')
```

```

print('values:', col_min_vals)
print('idxs:', col_min_idx)

Compute the minimum along each row; we get both the value and the minimum
row_min_vals, row_min_idx = x.min(dim=1)
print('\nMinimum along each row:')
print('values:', row_min_vals)
print('idxs:', row_min_idx)

```

```

Original tensor:
tensor([[2., 4., 3., 5.],
 [3., 3., 5., 2.]]) torch.Size([2, 4])

```

```

Overall minimum: tensor(2.)

```

```

Minimum along each column:
values: tensor([2., 3., 3., 2.])
idxs: tensor([0, 1, 0, 1])

```

```

Minimum along each row:
values: tensor([2., 2.])
idxs: tensor([0, 3])

```

Reduction operations *reduce* the rank of tensors: the dimension over which you perform the reduction will be removed from the shape of the output. If you pass `keepdim=True` to a reduction operation, the specified dimension will not be removed; the output tensor will instead have a shape of 1 in that dimension.

When you are working with multidimensional tensors, thinking about rows and columns can become confusing; instead it's more useful to think about the shape that will result from each operation. For example:

```

Create a tensor of shape (128, 10, 3, 64, 64)
x = torch.randn(128, 10, 3, 64, 64)
print(x.shape)

Take the mean over dimension 1; shape is now (128, 3, 64, 64)
y = x.mean(dim=1)

```

```

x = x.mean(dim=1)
print(x.shape)

Take the sum over dimension 2; shape is now (128, 3, 64)
x = x.sum(dim=2)
print(x.shape)

Take the mean over dimension 1, but keep the dimension from being eliminated
by passing keepdim=True; shape is now (128, 1, 64)
x = x.mean(dim=1, keepdim=True)
print(x.shape)

torch.Size([128, 10, 3, 64, 64])
torch.Size([128, 3, 64, 64])
torch.Size([128, 3, 64])
torch.Size([128, 1, 64])

```

Your turn: use reduction and indexing operations to implement a function that sets the minimum value along each row of a tensor to zero.

Hint: [torch.argmax](#)

```

def zero_row_min(x):
 """
 Return a copy of x, where the minimum value along each row has been set to 0.

 For example, if x is:
 x = torch.tensor([[
 [10, 20, 30],
 [2, 5, 1]
]])

 Then y = zero_row_min(x) should be:
 torch.tensor([
 [0, 20, 30],
 [2, 5, 0]
])

```

Inputs:

- x: Tensor of rank 2 with shape (N, M)

Returns:

- y: Tensor of rank 2 that is a copy of x, except the minimum value along each row is replaced with 0.

"""

y = x.clone()

#####  
# TODO: Complete the implementation of this function. #

#####  
# Replace "pass" statement with your code

inx = y.argmax(dim=1)

n,m = x.shape

y[torch.arange(n), inx] = 0

#####  
# END OF YOUR CODE #

#####  
return y

Now test your implementation with a few small test cases:

```
x0 = torch.tensor([[10, 20, 30], [2, 5, 1]])
```

```
print('Here is x0:')
```

```
print(x0)
```

```
y0 = zero_row_min(x0)
```

```
print('Here is y0:')
```

```
print(y0)
```

```
assert y0.tolist() == [[0, 20, 30], [2, 5, 0]]
```

```
x1 = torch.tensor([[2, 5, 10, -1], [1, 3, 2, 4], [5, 6, 2, 10]])
```

```
print('\nHere is x1:')
```

```
print(x1)
```

```
y1 = zero_row_min(x1)
```

```
print('Here is y1:')
```



```

print('here is y1:')
print(y1)
assert y1.tolist() == [[2, 5, 10, 0], [0, 3, 2, 4], [5, 6, 0, 10]]

print('\nSimple tests pass!')

```

```

Here is x0:
tensor([[10, 20, 30],
 [2, 5, 1]])
Here is y0:
tensor([[0, 20, 30],
 [2, 5, 0]])

Here is x1:
tensor([[2, 5, 10, -1],
 [1, 3, 2, 4],
 [5, 6, 2, 10]])
Here is y1:
tensor([[2, 5, 10, 0],
 [0, 3, 2, 4],
 [5, 6, 0, 10]])

Simple tests pass!

```

## Matrix operations

Note that unlike MATLAB, `*` is elementwise multiplication, not matrix multiplication. PyTorch provides a number of linear algebra functions that compute different types of vector and matrix products. The most commonly used are:

- [torch.dot](#): Computes inner product of vectors
- [torch.mm](#): Computes matrix-matrix products
- [torch.mv](#): Computes matrix-vector products
- [torch.addmm](#) / [torch.addmv](#): Computes matrix-matrix and matrix-vector multiplications plus a bias
- [torch.bmm](#) / [torch.baddmm](#): Batched versions of `torch.mm` and `torch.addmm`, respectively

- `torch.matmul`: General matrix product that performs different operations depending on the rank of the inputs; this is similar to `np.dot` in numpy.

You can find a full list of the available linear algebra operators [in the documentation](#).

Here is an example of using `torch.dot` to compute inner products. Like the other mathematical operators we've seen, most linear algebra operators are available both as functions in the `torch` module and as instance methods of tensors:

```
v = torch.tensor([9,10], dtype=torch.float32)
w = torch.tensor([11, 12], dtype=torch.float32)

Inner product of vectors
print('Dot products:')
print(torch.dot(v, w))
print(v.dot(w))

dot only works for vectors -- it will give an error for tensors of rank > 1
x = torch.tensor([[1,2],[3,4]], dtype=torch.float32)
y = torch.tensor([[5,6],[7,8]], dtype=torch.float32)
try:
 print(x.dot(y))
except RuntimeError as e:
 print(e)

Instead we use mm for matrix-matrix products:
print('\nMatrix-matrix product:')
print(torch.mm(x, y))
print(x.mm(y))
```

```
Dot products:
tensor(219.)
tensor(219.)
1D tensors expected, but got 2D and 2D tensors
```

```
Matrix-matrix product:
tensor([[19., 22.],
 [43., 50.]])
```

```
tensor([[19., 22.],
 [43., 50.]])
```

With all the different linear algebra operators that PyTorch provides, there is usually more than one way to compute something. For example to compute matrix-vector products we can use `torch.mv`; we can reshape the vector to have rank 2 and use `torch.mm`; or we can use `torch.matmul`. All give the same results, but the outputs might have different ranks:

```
print('Here is x (rank 2):')
print(x)
print('\nHere is v (rank 1):')
print(v)

Matrix-vector multiply with torch.mv produces a rank-1 output
print('\nMatrix-vector product with torch.mv (rank 1 output)')
print(torch.mv(x, v))
print(x.mv(v))

We can reshape the vector to have rank 2 and use torch.mm to perform
matrix-vector products, but the result will have rank 2
print('\nMatrix-vector product with torch.mm (rank 2 output)')
print(torch.mm(x, v.view(2, 1)))
print(x.mm(v.view(2, 1)))

print('\nMatrix-vector product with torch.matmul (rank 1 output)')
print(torch.matmul(x, v))
print(x.matmul(v))
```

```
Here is x (rank 2):
tensor([[1., 2.],
 [3., 4.]])
```

```
Here is v (rank 1):
tensor([9., 10.])
```

```
Matrix-vector product with torch.mv (rank 1 output)
```

```

tensor([29., 67.])
tensor([29., 67.])

Matrix-vector product with torch.mm (rank 2 output)
tensor([[29.],
 [67.]])
tensor([[29.],
 [67.]])

Matrix-vector product with torch.matmul (rank 1 output)
tensor([29., 67.])
tensor([29., 67.])

```

Your turn: use [torch.bmm](#) to perform a batched matrix multiply.

```

B, N, M, P = 3, 2, 5, 4
x = torch.rand(B, N, M) # Random tensor of shape (B, N, M)
y = torch.rand(B, M, P) # Random tensor of shape (B, M, P)

We can use a for loop to (inefficiently) compute a batch of matrix multiply
operations
z1 = torch.empty(B, N, P) # Empty tensor of shape (B, N, P)
for i in range(B):
 z1[i] = x[i].mm(y[i])
print('Here is the result of batched matrix multiply with a loop:')
print(z1)

z2 = None
#####
TODO: Use bmm to compute a batched matrix multiply between x and y; store
the result in z2.
#####
Replace "pass" statement with your code
z2 = x.bmm(y)
#####
END OF YOUR CODE
.....

```

```
#####
print('\nHere is the result of batched matrix multiply with bmm:')
print(z2)

The two may not return exactly the same result; different linear algebra
routines often return slightly different results due to the fact that
floating-point math is non-exact and non-associative.
diff = (z1 - z2).abs().max().item()
print('\nDifference:', diff)
print('Difference within threshold:', diff < 1e-6)
```

Here is the result of batched matrix multiply with a loop:

```
tensor([[[[1.1885, 2.0566, 0.5310, 2.0098],
 [1.1136, 1.9369, 0.4811, 1.7487]],

 [[1.6451, 1.6084, 1.5285, 1.3382],
 [0.8577, 1.4708, 0.9948, 1.0112]],

 [[0.9115, 1.5527, 1.3150, 1.7780],
 [0.5416, 1.3144, 0.9242, 1.2828]]]])
```

Here is the result of batched matrix multiply with bmm:

```
tensor([[[[1.1885, 2.0566, 0.5310, 2.0098],
 [1.1136, 1.9369, 0.4811, 1.7487]],

 [[1.6451, 1.6084, 1.5285, 1.3382],
 [0.8577, 1.4708, 0.9948, 1.0112]],

 [[0.9115, 1.5527, 1.3150, 1.7780],
 [0.5416, 1.3144, 0.9242, 1.2828]]]])
```

```
Difference: 1.1920928955078125e-07
Difference within threshold: True
```

## Broadcasting

Broadcasting is a powerful mechanism that allows PyTorch to work with arrays of different shapes when performing arithmetic operations. Frequently we have a smaller tensor and a larger tensor, and we want to use the smaller tensor multiple times to perform some operation on the larger tensor.

For example, suppose that we want to add a constant vector to each row of a tensor. We could do it like this:

```
We will add the vector v to each row of the matrix x,
storing the result in the matrix y
x = torch.tensor([[1,2,3], [4,5,6], [7,8,9], [10, 11, 12]])
v = torch.tensor([1, 0, 1])
y = torch.zeros_like(x) # Create an empty matrix with the same shape as x

Add the vector v to each row of the matrix x with an explicit loop
for i in range(4):
 y[i, :] = x[i, :] + v

print(y)
```

```
tensor([[2, 2, 4],
 [5, 5, 7],
 [8, 8, 10],
 [11, 11, 13]])
```

This works; however when the tensor  $x$  is very large, computing an explicit loop in Python could be slow. Note that adding the vector  $v$  to each row of the tensor  $x$  is equivalent to forming a tensor  $vv$  by stacking multiple copies of  $v$  vertically, then performing elementwise summation of  $x$  and  $vv$ . We could implement this approach like this:

```
vv = v.repeat((4, 1)) # Stack 4 copies of v on top of each other
print(vv) # Prints "[[1 0 1]
 # [1 0 1]
 # [1 0 1]
 # [1 0 1]]"
```

```
tensor([[1, 0, 1],
 [1, 0, 1],
 [1, 0, 1],
 [1, 0, 1]])
```

```
[1, 0, 1],
[1, 0, 1],
[1, 0, 1]])
```

```
y = x + vv # Add x and vv elementwise
print(y)
```

```
tensor([[2, 2, 4],
 [5, 5, 7],
 [8, 8, 10],
 [11, 11, 13]])
```

PyTorch broadcasting allows us to perform this computation without actually creating multiple copies of `v`. Consider this version, using broadcasting:

```
We will add the vector v to each row of the matrix x,
storing the result in the matrix y
x = torch.tensor([[1,2,3], [4,5,6], [7,8,9], [10, 11, 12]])
v = torch.tensor([1, 0, 1])
y = x + v # Add v to each row of x using broadcasting
print(y)
```

```
tensor([[2, 2, 4],
 [5, 5, 7],
 [8, 8, 10],
 [11, 11, 13]])
```

The line `y = x + v` works even though `x` has shape `(4, 3)` and `v` has shape `(3,)` due to broadcasting; this line works as if `v` actually had shape `(4, 3)`, where each row was a copy of `v`, and the sum was performed elementwise.

Broadcasting two tensors together follows these rules:

1. If the tensors do not have the same rank, prepend the shape of the lower rank array with 1s until both shapes have the same length.
2. The two tensors are said to be *compatible* in a dimension if they have the same size in the dimension, or if one of the tensors has size

- 1 in that dimension.
- 3. The tensors can be broadcast together if they are compatible in all dimensions.
- 4. After broadcasting, each tensor behaves as if it had shape equal to the elementwise maximum of shapes of the two input tensors.
- 5. In any dimension where one tensor had size 1 and the other tensor had size greater than 1, the first tensor behaves as if it were copied along that dimension

If this explanation does not make sense, try reading the explanation from the [documentation](#).

Not all functions support broadcasting. You can find functions that does not support broadcasting from the official docs. (e.g. [torch.mm](#) does not support broadcasting, but [torch.matmul](#) does)

Broadcasting can let us easily implement many different operations. For example we can compute an outer product of vectors:

```
Compute outer product of vectors
v = torch.tensor([1, 2, 3]) # v has shape (3,)
w = torch.tensor([4, 5]) # w has shape (2,)
To compute an outer product, we first reshape v to be a column
vector of shape (3, 1); we can then broadcast it against w to yield
an output of shape (3, 2), which is the outer product of v and w:
print(v.view(3, 1) * w)

 tensor([[4, 5],
 [8, 10],
 [12, 15]])
```

We can add a vector to each row of a matrix:

```
x = torch.tensor([[1, 2, 3], [4, 5, 6]]) # x has shape (2, 3)
v = torch.tensor([1, 2, 3]) # v has shape (3,)
print('Here is the matrix:')
print(x)
print('\nHere is the vector:')
print(v)
```



```
x has shape (2, 3) and v has shape (3,) so they broadcast to (2, 3),
giving the following matrix:
print('\nAdd the vector to each row of the matrix:')
print(x + v)
```

```
Here is the matrix:
tensor([[1, 2, 3],
 [4, 5, 6]])
```

```
Here is the vector:
tensor([1, 2, 3])
```

```
Add the vector to each row of the matrix:
tensor([[2, 4, 6],
 [5, 7, 9]])
```

We can add a vector to each column of a matrix:

```
x = torch.tensor([[1, 2, 3], [4, 5, 6]]) # x has shape (2, 3)
w = torch.tensor([4, 5]) # w has shape (2,)
print('Here is the matrix:')
print(x)
print('\nHere is the vector:')
print(w)

x has shape (2, 3) and w has shape (2,). We reshape w to (2, 1);
then when we add the two the result broadcasts to (2, 3):
print('\nAdd the vector to each column of the matrix:')
print(x + w.view(-1, 1))

Another solution is the following:
1. Transpose x so it has shape (3, 2)
2. Since w has shape (2,), adding will broadcast to (3, 2)
3. Transpose the result, resulting in a shape (2, 3)
print((x.t() + w).t())
```

Here is the matrix:

```
tensor([[1, 2, 3],
 [4, 5, 6]])
```

Here is the vector:

```
tensor([4, 5])
```

Add the vector to each column of the matrix:

```
tensor([[5, 6, 7],
 [9, 10, 11]])
tensor([[5, 6, 7],
 [9, 10, 11]])
```

Multiply a tensor by a set of constants:

```
x = torch.tensor([[1, 2, 3], [4, 5, 6]]) # x has shape (2, 3)
c = torch.tensor([1, 10, 11, 100]) # c has shape (4)
print('Here is the matrix:')
print(x)
print('\nHere is the vector:')
print(c)

We do the following:
1. Reshape c from (4,) to (4, 1, 1)
2. x has shape (2, 3). Since they have different ranks, when we multiply the
two, x behaves as if its shape were (1, 2, 3)
3. The result of the broadcast multiplication between tensor of shape
(4, 1, 1) and (1, 2, 3) has shape (4, 2, 3)
4. The result y has shape (4, 2, 3), and y[i] (shape (2, 3)) is equal to
c[i] * x
y = c.view(-1, 1, 1) * x
print('\nMultiply x by a set of constants:')
print(y)
```

Here is the matrix:

```
tensor([[1, 2, 3],
```

```
tensor([[[1, 2, 3],
 [4, 5, 6]])
```

Here is the vector:

```
tensor([1, 10, 11, 100])
```

Multiply x by a set of constants:

```
tensor([[[1, 2, 3],
 [4, 5, 6]],

 [[10, 20, 30],
 [40, 50, 60]],

 [[11, 22, 33],
 [44, 55, 66]],

 [[100, 200, 300],
 [400, 500, 600]]])
```

Your turn: write a function that normalizes the columns of a matrix. It should compute the mean and standard deviation of each column, then subtract the mean and divide by the standard deviation for each element in the column.

Example:

```
x = [[0, 30, 600],
 [1, 10, 200],
 [-1, 20, 400]]
```

- The first column has mean 0 and std 1
- The second column has mean 20 and std 10
- The third column has mean 400 and std 200

After normalizing the columns, the result should be:

```
y = [[0, 1, 1],
```

```
[1, -1, -1],
[-1, 0, 0]]
```

```
def normalize_columns(x):
 """
 Normalize the columns of a matrix by subtracting the mean and dividing by the
 standard deviation.

 Inputs:
 - x: Tensor of shape (N, M)

 Returns:
 - y: Tensor of shape (N, M) which is a copy of x with normalized columns.
 """
 y = x.clone()
 #####
 # TODO: Complete the implementation of this function. Do not modify x. #
 # Your implementation should not use any loops; instead you should use #
 # reduction and broadcasting operations. #
 #####
 # Replace "pass" statement with your code
 y = (y - y.mean(dim=0)) / y.std(dim=0)
 #####
 # END OF YOUR CODE #
 #####
 return y
```

Now test your implementation with a simple test:

```
x0 = torch.tensor([[0., 30., 600.], [1., 10., 200.], [-1., 20., 400.]])
y0 = normalize_columns(x0)
print('Here is x0:')
print(x0)
```

```
print('Here is y0:')
print(y0)
assert y0.tolist() == [[0., 1., 1.], [1., -1., -1.], [-1., 0., 0.]]
assert x0.tolist() == [[0., 30., 600.], [1., 10., 200.], [-1., 20., 400.]]
```

```
Here is x0:
tensor([[0., 30., 600.],
 [1., 10., 200.],
 [-1., 20., 400.]])
Here is y0:
tensor([[0., 1., 1.],
 [1., -1., -1.],
 [-1., 0., 0.]])
```

## Running on GPU

One of the most important features of PyTorch is that it can use graphics processing units (GPUs) to accelerate its tensor operations.

We can easily check whether PyTorch is configured to use GPUs:

Tensors can be moved onto any device using the `.to` method.

```
import torch

if torch.cuda.is_available:
 print('PyTorch can use GPUs!')
else:
 print('PyTorch cannot use GPUs.')
```

```
PyTorch can use GPUs!
```

You can enable GPUs in Colab via Runtime -> Change Runtime Type -> Hardware Accelerator -> GPU.

This may cause the Colab runtime to restart, so we will re-import torch in the next cell.

We have already seen that PyTorch tensors have a `dtype` attribute specifying their datatype. All PyTorch tensors also have a `device` attribute that specifies the device where the tensor is stored -- either CPU, or CUDA (for NVIDIA GPUs). A tensor on a CUDA device will automatically use that device to accelerate all of its operations.

Just as with datatypes, we can use the `.to()` method to change the device of a tensor. We can also use the convenience methods `.cuda()` and `.cpu()` methods to move tensors between CPU and GPU.

```
Construct a tensor on the CPU
x0 = torch.tensor([[1, 2], [3, 4]], dtype=torch.float32)
print('x0 device:', x0.device)

Move it to the GPU using .to()
x1 = x0.to('cuda')
print('x1 device:', x1.device)

Move it to the GPU using .cuda()
x2 = x0.cuda()
print('x2 device:', x2.device)

Move it back to the CPU using .to()
x3 = x1.to('cpu')
print('x3 device:', x3.device)

Move it back to the CPU using .cpu()
x4 = x2.cpu()
print('x4 device:', x4.device)

We can construct tensors directly on the GPU as well
y = torch.tensor([[1, 2, 3], [4, 5, 6]], dtype=torch.float64, device='cuda')
print('y device / dtype:', y.device, y.dtype)

Calling x.to(y) where y is a tensor will return a copy of x with the same
device and dtype as y
x5 = x0.to(y)
```

```
print('x5 device / dtype:', x5.device, x5.dtype)
```

```
x0 device: cpu
x1 device: cuda:0
x2 device: cuda:0
x3 device: cpu
x4 device: cpu
y device / dtype: cuda:0 torch.float64
x5 device / dtype: cuda:0 torch.float64
```

Performing large tensor operations on a GPU can be **a lot faster** than running the equivalent operation on CPU.

Here we compare the speed of adding two tensors of shape (10000, 10000) on CPU and GPU:

(Note that GPU code may run asynchronously with CPU code, so when timing the speed of operations on the GPU it is important to use `torch.cuda.synchronize` to synchronize the CPU and GPU.)

```
import time

a_cpu = torch.randn(10000, 10000, dtype=torch.float32)
b_cpu = torch.randn(10000, 10000, dtype=torch.float32)

a_gpu = a_cpu.cuda()
b_gpu = b_cpu.cuda()
torch.cuda.synchronize()

t0 = time.time()
c_cpu = a_cpu + b_cpu
t1 = time.time()
c_gpu = a_gpu + b_gpu
torch.cuda.synchronize()
t2 = time.time()

Check that they computed the same thing
diff = (c_gpu.cpu() - c_cpu).abs().max().item()
print('Max difference between c_gpu and c_cpu:', diff)
```

```

cpu_time = 1000.0 * (t1 - t0)
gpu_time = 1000.0 * (t2 - t1)
print('CPU time: %.2f ms' % cpu_time)
print('GPU time: %.2f ms' % gpu_time)
print('GPU speedup: %.2f x' % (cpu_time / gpu_time))

```

```

 Max difference between c_gpu and c_cpu: 0.0
 CPU time: 138.55 ms
 GPU time: 5.09 ms
 GPU speedup: 27.22 x

```

You should see that running the same computation on the GPU was more than 30 times faster than on the CPU! Due to the massive speedups that GPUs offer, we will use GPUs to accelerate much of our machine learning code starting in Assignment 2.

Your turn: Use the GPU to accelerate the following matrix multiplication operation. You should see ~10x speedup by using the GPU.

```

import time

x = torch.rand(512, 4096)
w = torch.rand(4096, 4096)

t0 = time.time()
y0 = x.mm(w)
t1 = time.time()

y1 = None
#####
TODO: Write a bit of code that performs matrix multiplication of x and w
on the GPU, and then moves the result back to the CPU. Store the result
in y1.#####
Replace "pass" statement with your code
y1 = x.cuda().mm(w.cuda()).cpu()
#####
END OF YOUR CODE

```



```
#####
torch.cuda.synchronize()
t2 = time.time()

print('y1 on CPU:', y1.device == torch.device('cpu'))
diff = (y0 - y1).abs().max().item()
print('Max difference between y0 and y1:', diff)
print('Difference within tolerance:', diff < 5e-2)

cpu_time = 1000.0 * (t1 - t0)
gpu_time = 1000.0 * (t2 - t1)
print('CPU time: %.2f ms' % cpu_time)
print('GPU time: %.2f ms' % gpu_time)
print('GPU speedup: %.2f x' % (cpu_time / gpu_time))
```

```
y1 on CPU: True
Max difference between y0 and y1: 0.001220703125
Difference within tolerance: True
CPU time: 220.52 ms
GPU time: 28.95 ms
GPU speedup: 7.62 x
```

