

Install starter code

We will continue using the utility functions that we've used for Assignment 1: [cutils package](#). Run this cell to download and install it.

```
1 !pip install git+https://github.com/deepvision-class/starter-code
```

Looking in indexes: <https://pypi.org/simple>, <https://us-python.pkg.dev/colab-wheels/public/simple/>

Collecting git+<https://github.com/deepvision-class/starter-code>

Cloning <https://github.com/deepvision-class/starter-code> to /tmp/pip-req-build-9c3z_bm3

Running command git clone -q <https://github.com/deepvision-class/starter-code> /tmp/pip-req-build-9c3z_bm3

Requirement already satisfied: pydrive in /usr/local/lib/python3.7/dist-packages (from Colab-Utils==0.1.dev0) (1.3.1)

Requirement already satisfied: oauth2client>=4.0.0 in /usr/local/lib/python3.7/dist-packages (from pydrive->Colab-Utils==0.1.dev0) (4.1.3)

Requirement already satisfied: google-api-python-client>=1.2 in /usr/local/lib/python3.7/dist-packages (from pydrive->Colab-Utils==0.1.dev0) (1.2.0)

Requirement already satisfied: PyYAML>=3.0 in /usr/local/lib/python3.7/dist-packages (from pydrive->Colab-Utils==0.1.dev0) (3.13)

Requirement already satisfied: google-auth<3dev,>=1.16.0 in /usr/local/lib/python3.7/dist-packages (from google-api-python-client>=1.2->pydrive) (1.16.0)

Requirement already satisfied: httplib2<1dev,>=0.15.0 in /usr/local/lib/python3.7/dist-packages (from google-api-python-client>=1.2->pydrive) (0.15.0)

Requirement already satisfied: google-api-core<3dev,>=1.21.0 in /usr/local/lib/python3.7/dist-packages (from google-api-python-client>=1.2->pydrive) (1.21.0)

Requirement already satisfied: uritemplate<4dev,>=3.0.0 in /usr/local/lib/python3.7/dist-packages (from google-api-python-client>=1.2->pydrive) (3.0.0)

Requirement already satisfied: six<2dev,>=1.13.0 in /usr/local/lib/python3.7/dist-packages (from google-api-python-client>=1.2->pydrive) (1.13.0)

Requirement already satisfied: google-auth-httplib2>=0.0.3 in /usr/local/lib/python3.7/dist-packages (from google-api-python-client>=1.2->pydrive) (0.0.3)

Requirement already satisfied: setuptools>=40.3.0 in /usr/local/lib/python3.7/dist-packages (from google-api-core<3dev,>=1.21.0->google-api-python-client) (40.3.0)

Requirement already satisfied: packaging>=14.3 in /usr/local/lib/python3.7/dist-packages (from google-api-core<3dev,>=1.21.0->google-api-python-client) (14.3)

Requirement already satisfied: pytz in /usr/local/lib/python3.7/dist-packages (from google-api-core<3dev,>=1.21.0->google-api-python-client) (2017.4.17)

Requirement already satisfied: googleapis-common-protos<2.0dev,>=1.6.0 in /usr/local/lib/python3.7/dist-packages (from google-api-core<3dev,>=1.21.0->google-api-python-client) (1.6.0)

Requirement already satisfied: protobuf<4.0.0dev,>=3.12.0 in /usr/local/lib/python3.7/dist-packages (from google-api-core<3dev,>=1.21.0->google-api-python-client) (3.12.0)

Requirement already satisfied: requests<3.0.0dev,>=2.18.0 in /usr/local/lib/python3.7/dist-packages (from google-api-core<3dev,>=1.21.0->google-api-python-client) (2.18.0)

Requirement already satisfied: rsa<5,>=3.1.4 in /usr/local/lib/python3.7/dist-packages (from google-auth<3dev,>=1.16.0->google-api-python-client) (3.1.4)

Requirement already satisfied: cachetools<5.0,>=2.0.0 in /usr/local/lib/python3.7/dist-packages (from google-auth<3dev,>=1.16.0->google-api-python-client) (2.0.0)

Requirement already satisfied: pyasn1-modules>=0.2.1 in /usr/local/lib/python3.7/dist-packages (from google-auth<3dev,>=1.16.0->google-api-python-client) (0.2.1)

Requirement already satisfied: pyasn1>=0.1.7 in /usr/local/lib/python3.7/dist-packages (from google-auth<3dev,>=1.16.0->google-api-python-client) (0.1.7)

Requirement already satisfied: pyparsing!=3.0.5,>=2.0.2 in /usr/local/lib/python3.7/dist-packages (from packaging>=14.3->google-api-core<3dev,>=1.21.0->google-api-python-client) (2.0.2)

Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.7/dist-packages (from requests<3.0.0dev,>=2.18.0->google-api-core<3dev,>=1.21.0->google-api-python-client) (2017.4.17)

Requirement already satisfied: urllib3!=1.25.0,!1.25.1,<1.26,>=1.21.1 in /usr/local/lib/python3.7/dist-packages (from requests<3.0.0dev,>=2.18.0->google-api-core<3dev,>=1.21.0->google-api-python-client) (1.25.1)

Requirement already satisfied: idna<3,>=2.5 in /usr/local/lib/python3.7/dist-packages (from requests<3.0.0dev,>=2.18.0->google-api-core<3dev,>=1.21.0->google-api-python-client) (2.5)

Requirement already satisfied: chardet<4,>=3.0.2 in /usr/local/lib/python3.7/dist-packages (from requests<3.0.0dev,>=2.18.0->google-api-core<3dev,>=1.21.0->google-api-python-client) (3.0.2)

Building wheels for collected packages: Colab-Utils

Building wheel for Colab-Utils (setup.py) ... done

Created wheel for Colab-Utils: filename=Colab_Utils-0.1.dev0-py3-none-any.whl size=10306 sha256=a718a4a78f2e38fa855d6bdacbb8f7687740738

Stored in directory: /tmp/pip-ephem-wheel-cache-sve0uwoi/wheels/eb/3c/88/465b0d78ef4a63d1f487c4208bd4691a448f05923eda0ef5f6

Installing collected packages: Colab-Utills
Successfully installed Colab-Utills-0.1.dev0

✓ 0 сек. выполнено в 08:58



▼ Setup code

Run some setup code for this notebook: Import some useful packages and increase the default figure size.

```
1 from __future__ import print_function
2 from __future__ import division
3
4 import torch
5 import couils
6 import random
7 import time
8 import math
9 import matplotlib.pyplot as plt
10 from torchvision.utils import make_grid
11
12 %matplotlib inline
13 plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
14 plt.rcParams['image.interpolation'] = 'nearest'
15 plt.rcParams['image.cmap'] = 'gray'
```

Starting in this assignment, we will use the GPU to accelerate our computation. Run this cell to make sure you are using a GPU.

```
1 if torch.cuda.is_available:
2     print('Good to go!')
3 else:
4     print('Please set GPU via Edit -> Notebook Settings.')
```

Good to go!

Now, we will load CIFAR10 dataset, with normalization.

In this notebook we will use the **bias trick**: By adding an extra constant feature of ones to each image, we avoid the need to keep track of a bias vector; the bias will be encoded as the part of the weight matrix that interacts with the constant ones in the input.

In the `two_layer_net.ipynb` notebook that follows this one, we will not use the bias trick.

```

1 def get_CIFAR10_data(validation_ratio = 0.02):
2     """
3     Load the CIFAR-10 dataset from disk and perform preprocessing to prepare
4     it for the linear classifier. These are the same steps as we used for the
5     SVM, but condensed to a single function.
6     """
7     X_train, y_train, X_test, y_test = coutils.data.cifar10()
8
9     # Move all the data to the GPU
10    X_train = X_train.cuda()
11    y_train = y_train.cuda()
12    X_test = X_test.cuda()
13    y_test = y_test.cuda()
14
15    # 0. Visualize some examples from the dataset.
16    class_names = [
17        'plane', 'car', 'bird', 'cat', 'deer',
18        'dog', 'frog', 'horse', 'ship', 'truck'
19    ]
20    img = coutils.utils.visualize_dataset(X_train, y_train, 12, class_names)
21    plt.imshow(img)
22    plt.axis('off')
23    plt.show()
24
25    # 1. Normalize the data: subtract the mean RGB (zero mean)
26    mean_image = X_train.mean(dim=0, keepdim=True).mean(dim=2, keepdim=True).mean(dim=3, keepdim=True)
27    X_train -= mean_image
28    X_test -= mean_image
29
30    # 2. Reshape the image data into rows
31    X_train = X_train.reshape(X_train.shape[0], -1)
32    X_test = X_test.reshape(X_test.shape[0], -1)
33    # print(X_test.shape)
34
35    # 3. Add bias dimension and transform into columns
36    ones_train = torch.ones(X_train.shape[0], 1, device=X_train.device)
37    X_train = torch.cat([X_train, ones_train], dim=1)
38    ones_test = torch.ones(X_test.shape[0], 1, device=X_test.device)
39    X_test = torch.cat([X_test, ones_test], dim=1)
40    # print(X_test.shape)
41
42    # 4. Carve out part of the training set to use for validation.
43    # For random permutation, you can use torch.randperm or torch.randint
44    # But, for this homework, we use slicing instead.
45    num_training = int( X_train.shape[0] * (1.0 - validation_ratio) )

```




```
Train data shape: torch.Size([49000, 3073])
Train labels shape: torch.Size([49000])
Validation data shape: torch.Size([1000, 3073])
Validation labels shape: torch.Size([1000])
Test data shape: torch.Size([10000, 3073])
Test labels shape: torch.Size([10000])
```

For Softmax and SVM, we will analytically compute the gradient, as a sanity check.

```
1 def grad_check_sparse(f, x, analytic_grad, num_checks=10, h=1e-7):
2     """
3     Utility function to perform numeric gradient checking. We use the centered
4     difference formula to compute a numeric derivative:
5
6      $f'(x) \approx (f(x + h) - f(x - h)) / (2h)$ 
7
8     Rather than computing a full numeric gradient, we sparsely sample a few
9     dimensions along which to compute numeric derivatives.
10
11     Inputs:
12     - f: A function that inputs a torch tensor and returns a torch scalar
13     - x: A torch tensor giving the point at which to evaluate the numeric gradient
14     - analytic_grad: A torch tensor giving the analytic gradient of f at x
15     - num_checks: The number of dimensions along which to check
16     - h: Step size for computing numeric derivatives
17     """
18     # fix random seed for
19     coutils.utils.fix_random_seed()
20
21     for i in range(num_checks):
22
23         ix = tuple([random.randrange(m) for m in x.shape])
24
25         oldval = x[ix].item()
26         x[ix] = oldval + h # increment by h
27         fxph = f(x).item() # evaluate f(x + h)
28         x[ix] = oldval - h # increment by h
29         fxmh = f(x).item() # evaluate f(x - h)
30         x[ix] = oldval # reset
```

```

31
32     grad_numerical = (fxph - fxmh) / (2 * h)
33     grad_analytic = analytic_grad[ix]
34     rel_error_top = abs(grad_numerical - grad_analytic)
35     rel_error_bot = (abs(grad_numerical) + abs(grad_analytic) + 1e-12)
36     rel_error = rel_error_top / rel_error_bot
37     msg = 'numerical: %f analytic: %f, relative error: %e'
38     print(msg % (grad_numerical, grad_analytic, rel_error))

```

SVM Classifier

In this section, you will:

- implement a fully-vectorized **loss function** for the SVM
- implement the fully-vectorized expression for its **analytic gradient**
- **check your implementation** using numerical gradient
- use a validation set to **tune the learning rate and regularization** strength
- **optimize** the loss function with **SGD**
- **visualize** the final learned weights

```

1 def svm_loss_naive(W, X, y, reg):
2     """
3     Structured SVM loss function, naive implementation (with loops).
4
5     Inputs have dimension D, there are C classes, and we operate on minibatches
6     of N examples. When you implment the regularization over W, please DO NOT
7     multiply the regularization term by 1/2 (no coefficient).
8
9     Inputs:
10    - W: A PyTorch tensor of shape (D, C) containing weights.
11    - X: A PyTorch tensor of shape (N, D) containing a minibatch of data.
12    - y: A PyTorch tensor of shape (N,) containing training labels; y[i] = c means
13        that X[i] has label c, where 0 <= c < C.
14    - reg: (float) regularization strength
15
16    Returns a tuple of:
17    - loss as torch scalar
18    - gradient of loss with respect to weights W; a tensor of same shape as W
19    """
20    dW = torch.zeros_like(W) # initialize the gradient as zero
21

```

```

22 # compute the loss and the gradient
23 num_classes = W.shape[1]
24 num_train = X.shape[0]
25 loss = 0.0
26 for i in range(num_train):
27     scores = W.t().mv(X[i])
28     correct_class_score = scores[y[i]]
29     for j in range(num_classes):
30         if j == y[i]:
31             continue
32         margin = scores[j] - correct_class_score + 1 # note delta = 1
33         if margin > 0:
34             loss += margin
35         #####
36         # TODO:                                     #
37         # Compute the gradient of the loss function and store it dW. (part 1) #
38         # Rather than first computing the loss and then computing the         #
39         # derivative, it is simple to compute the derivative at the same time #
40         # that the loss is being computed.                                     #
41         #####
42         # Replace "pass" statement with your code
43         dW[:,y[i]] += -X[i,:].t()
44         dW[:,j] += X[i,:].t()
45         #####
46         #                                     END OF YOUR CODE                 #
47         #####
48
49
50 # Right now the loss is a sum over all training examples, but we want it
51 # to be an average instead so we divide by num_train.
52 loss /= num_train
53
54 # Add regularization to the loss.
55 loss += reg * torch.sum(W * W)
56
57 #####
58 # TODO:                                     #
59 # Compute the gradient of the loss function and store it in dW. (part 2) #
60 #####
61 # Replace "pass" statement with your code
62 dW = dW / num_train + reg * 2 * W
63 #####
64 #                                     END OF YOUR CODE                 #
65 #####
66
67 return loss, dW

```

Evaluate the naive implementation of the loss we provided for you. You will get around 9.000175.

```
1 # generate a random SVM weight tensor of small numbers
2 coutils.utils.fix_random_seed()
3 W = torch.randn(3073, 10, device=data_dict['X_val'].device) * 0.0001
4
5 loss, grad = svm_loss_naive(W, data_dict['X_val'], data_dict['y_val'], 0.000005)
6 print('loss: %f' % (loss, ))
```

```
loss: 9.000433
```

The `grad` returned from the function above is right now all zero. Derive and implement the gradient for the SVM cost function and implement it inline inside the function `svm_loss_naive`. You will find it helpful to interleave your new code inside the existing function.

To check that you have implemented the gradient correctly, you can numerically estimate the gradient of the loss function and compare the numeric estimate to the gradient that you computed. We have provided code that does this for you (The relative errors should be less than $1e-6$).

```
1 # Once you've implemented the gradient, recompute it with the code below
2 # and gradient check it with the function we provided for you
3
4 # Use a random W and a minibatch of data from the val set for gradient checking
5 # For numeric gradient checking it is a good idea to use 64-bit floating point
6 # numbers for increased numeric precision; however when actually training models
7 # we usually use 32-bit floating point numbers for increased speed.
8 coutils.utils.fix_random_seed()
9 W = 0.0001 * torch.randn(3073, 10, device=data_dict['X_val'].device).double()
10 batch_size = 64
11 X_batch = data_dict['X_val'][:64].double()
12 y_batch = data_dict['y_val'][:64]
13
14 # Compute the loss and its gradient at W.
15 loss, grad = svm_loss_naive(W.double(), X_batch, y_batch, reg=0.0)
16
17 # Numerically compute the gradient along several randomly chosen dimensions, and
18 # compare them with your analytically computed gradient. The numbers should
19 # match almost exactly along all dimensions.
20 f = lambda w: svm_loss_naive(w, X_batch, y_batch, reg=0.0)[0]
21 grad_numerical = grad_check_sparse(f, W.double(), grad)
```



```
numerical: -0.034577 analytic: -0.034577, relative error: 2.372452e-07
numerical: 0.126951 analytic: 0.126951, relative error: 5.721190e-08
numerical: -0.068597 analytic: -0.068597, relative error: 2.249695e-07
numerical: 0.025717 analytic: 0.025717, relative error: 4.774223e-07
numerical: 0.048266 analytic: 0.048266, relative error: 2.668362e-07
numerical: 0.052260 analytic: 0.052260, relative error: 2.475153e-07
numerical: 0.096133 analytic: 0.096133, relative error: 4.690208e-09
numerical: 0.032702 analytic: 0.032702, relative error: 3.644517e-07
numerical: -0.117158 analytic: -0.117158, relative error: 4.006759e-08
numerical: -0.154093 analytic: -0.154093, relative error: 7.809949e-08
```

Let's do the gradient check once again with regularization turned on. (You didn't forget the regularization gradient, did you?)

You should see relative errors less than `1e-5`.

```
1 # Use a minibatch of data from the val set for gradient checking
2 coutils.utils.fix_random_seed()
3 W = 0.0001 * torch.randn(3073, 10, device=data_dict['X_val'].device).double()
4 batch_size = 64
5 X_batch = data_dict['X_val'][:64].double()
6 y_batch = data_dict['y_val'][:64]
7
8 # Compute the loss and its gradient at W.
9 loss, grad = svm_loss_naive(W.double(), X_batch, y_batch, reg=1e3)
10
11 # Numerically compute the gradient along several randomly chosen dimensions, and
12 # compare them with your analytically computed gradient. The numbers should
13 # match almost exactly along all dimensions.
14 f = lambda w: svm_loss_naive(w, X_batch, y_batch, reg=1e3)[0]
15 grad_numerical = grad_check_sparse(f, W.double(), grad)
```

```
numerical: -0.121624 analytic: -0.121624, relative error: 7.160875e-08
numerical: 0.020923 analytic: 0.020923, relative error: 3.331613e-07
numerical: -0.076604 analytic: -0.076604, relative error: 1.800879e-07
numerical: 0.256069 analytic: 0.256069, relative error: 6.121908e-08
numerical: -0.330089 analytic: -0.330089, relative error: 4.165267e-08
numerical: 0.004713 analytic: 0.004713, relative error: 3.512375e-06
numerical: 0.101968 analytic: 0.101968, relative error: 1.802643e-08
numerical: 0.097235 analytic: 0.097235, relative error: 1.618033e-07
numerical: -0.117112 analytic: -0.117112, relative error: 2.948518e-08
numerical: -0.257260 analytic: -0.257260, relative error: 4.474401e-08
```

Now, let's implement vectorized version of SVM: `svm_loss_vectorized`. It should compute the same inputs and outputs as the naive version

```
41 #####
```

```

44 #####
45 # TODO:                                                                    #
46 # Implement a vectorized version of the gradient for the structured SVM    #
47 # loss, storing the result in dW.                                          #
48 #                                                                           #
49 # Hint: Instead of computing the gradient from scratch, it may be easier   #
50 # to reuse some of the intermediate values that you used to compute the   #
51 # loss.                                                                     #
52 #####
53 # Replace "pass" statement with your code
54 scores_diff[scores_diff > 0] = 1
55 correct_label_vals = torch.sum(scores_diff , 1) * -1
56 scores_diff[correct_label_score_idxes] = correct_label_vals
57
58 dW = X.t().mm(scores_diff)
59 dW /= num_train
60 # add the regularization contribution to the gradient
61 dW += reg * 2* W
62 #####
63 #                               END OF YOUR CODE                           #
64 #####
65
66 return loss, dW

```

Let's first check the speed and performance between the non-vectorized and the vectorized version. You should see a speedup of more than 100x.

(Note: It may have some difference, but should be less than 1e-6)

```

1 # Next implement the function svm_loss_vectorized; for now only compute the loss;
2 # we will implement the gradient in a moment.
3
4 # Use random weights and a minibatch of val data for gradient checking
5 coutils.utils.fix_random_seed()
6 W = 0.0001 * torch.randn(3073, 10, device=data_dict['X_val'].device).double()
7 X_batch = data_dict['X_val'][:128].double()
8 y_batch = data_dict['y_val'][:128]
9 reg = 0.000005
10
11 # Run and time the naive version
12 torch.cuda.synchronize()
13 tic = time.time()
14 loss_naive, grad_naive = svm_loss_naive(W, X_batch, y_batch, reg)
15 torch.cuda.synchronize()
16 toc = time.time()

```

```

16 toc = time.time()
17 ms_naive = 1000.0 * (toc - tic)
18 print('Naive loss: %e computed in %.2fms' % (loss_naive, ms_naive))
19
20 # Run and time the vectorized version
21 torch.cuda.synchronize()
22 tic = time.time()
23 loss_vec, _ = svm_loss_vectorized(W, X_batch, y_batch, reg)
24 torch.cuda.synchronize()
25 toc = time.time()
26 ms_vec = 1000.0 * (toc - tic)
27 print('Vectorized loss: %e computed in %.2fms' % (loss_vec, ms_vec))
28
29 # The losses should match but your vectorized implementation should be much faster.
30 print('Difference: %.2e' % (loss_naive - loss_vec))
31 print('Speedup: %.2fX' % (ms_naive / ms_vec))

```

```

Naive loss: 9.000144e+00 computed in 225.62ms
Vectorized loss: 9.000144e+00 computed in 25.49ms
Difference: -1.07e-14
Speedup: 8.85X

```

Then, let's compute the gradient of the loss function. We can check the difference of gradient as well. (The error should be less than $1e-6$)

Now implement a vectorized version of the gradient computation in `svm_loss_vectorize` above. Run the cell below to compare the gradient of your naive and vectorized implementations. The difference between the gradients should be less than $1e-6$, and the vectorized version should run at least 100x faster.

```

1 # The naive implementation and the vectorized implementation should match, but
2 # the vectorized version should still be much faster.
3
4 # Use random weights and a minibatch of val data for gradient checking
5 coutils.utils.fix_random_seed()
6 W = 0.0001 * torch.randn(3073, 10, device=data_dict['X_val'].device).double()
7 X_batch = data_dict['X_val'][:128].double()
8 y_batch = data_dict['y_val'][:128]
9 reg = 0.000005
10
11 # Run and time the naive version
12 torch.cuda.synchronize()
13 tic = time.time()
14 _, grad_naive = svm_loss_naive(W, X_batch, y_batch, 0.000005)
15 torch.cuda.synchronize()
16 toc = time.time()

```

```

17 ms_naive = 1000.0 * (toc - tic)
18 print('Naive loss and gradient: computed in %.2fms' % ms_naive)
19
20 # Run and time the vectorized version
21 torch.cuda.synchronize()
22 tic = time.time()
23 _, grad_vec = svm_loss_vectorized(W, X_batch, y_batch, 0.000005)
24 torch.cuda.synchronize()
25 toc = time.time()
26 print('Vectorized loss and gradient: computed in %fs' % (toc - tic))
27
28 # The loss is a single number, so it is easy to compare the values computed
29 # by the two implementations. The gradient on the other hand is a tensor, so
30 # we use the Frobenius norm to compare them.
31 grad_difference = torch.norm(grad_naive - grad_vec, p='fro')
32 print('Gradient difference: %.2e' % grad_difference)
33 print('Speedup: %.2fX' % (ms_naive / ms_vec))

```

```

Naive loss and gradient: computed in 228.97ms
Vectorized loss and gradient: computed in 0.002930s
Gradient difference: 0.00e+00
Speedup: 8.98X

```

Now that we have an efficient vectorized implementation of the SVM loss and its gradient, we can implement a training pipeline for linear classifiers.

Complete the implementation of the following function:

```

1 def train_linear_classifier(loss_func, W, X, y, learning_rate=1e-3,
2                             reg=1e-5, num_iters=100, batch_size=200, verbose=False):
3     """
4     Train this linear classifier using stochastic gradient descent.
5
6     Inputs:
7     - loss_func: loss function to use when training. It should take W, X, y
8       and reg as input, and output a tuple of (loss, dW)
9     - W: A PyTorch tensor of shape (D, C) giving the initial weights of the
10       classifier. If W is None then it will be initialized here.
11     - X: A PyTorch tensor of shape (N, D) containing training data; there are N
12       training samples each of dimension D.
13     - y: A PyTorch tensor of shape (N,) containing training labels; y[i] = c
14       means that X[i] has label 0 ≤ c < C for C classes.
15     - learning_rate: (float) learning rate for optimization.
16     - reg: (float) regularization strength.

```

```

17 - num_iters: (integer) number of steps to take when optimizing
18 - batch_size: (integer) number of training examples to use at each step.
19 - verbose: (boolean) If true, print progress during optimization.
20
21 Returns: A tuple of:
22 - W: The final value of the weight matrix and the end of optimization
23 - loss_history: A list of Python scalars giving the values of the loss at each
24   training iteration.
25 """
26 # assume y takes values 0...K-1 where K is number of classes
27 num_classes = torch.max(y) + 1
28 num_train, dim = X.shape
29 if W is None:
30     # lazily initialize W
31     W = 0.000001 * torch.randn(dim, num_classes, device=X.device, dtype=X.dtype)
32
33 # Run stochastic gradient descent to optimize W
34 loss_history = []
35 for it in range(num_iters):
36     X_batch = None
37     y_batch = None
38     #####
39     # TODO:                                     #
40     # Sample batch_size elements from the training data and their       #
41     # corresponding labels to use in this round of gradient descent.     #
42     # Store the data in X_batch and their corresponding labels in       #
43     # y_batch; after sampling, X_batch should have shape (batch_size, dim) #
44     # and y_batch should have shape (batch_size,)                       #
45     #                                                                     #
46     # Hint: Use torch.randint to generate indices.                       #
47     #####
48     # Replace "pass" statement with your code
49     batch_idxes = torch.randint(num_train, (batch_size,))
50     X_batch = X[batch_idxes, :]
51     y_batch = y[batch_idxes]
52     #####
53     #                                     END OF YOUR CODE                 #
54     #####
55
56     # evaluate loss and gradient
57     loss, grad = loss_func(W, X_batch, y_batch, reg)
58     loss_history.append(loss.item())
59
60     # perform parameter update
61     #####
62     # TODO:                                     #
63     # Update the weights using the gradient and the learning rate.     #

```

```

64 #####
65 # Replace "pass" statement with your code
66 W -= learning_rate * grad
67 #####
68 #                               END OF YOUR CODE                               #
69 #####
70
71 if verbose and it % 100 == 0:
72     print('iteration %d / %d: loss %f' % (it, num_iters, loss))
73
74 return W, loss_history

```

Once you have implemented the training function, run the following cell to train a linear classifier using some default hyperparameters:

(You should see a final loss close to 9.0, and your training loop should run in about two seconds)

```

1 # fix random seed before we perform this operation
2 coutils.utils.fix_random_seed()
3
4 torch.cuda.synchronize()
5 tic = time.time()
6
7 W, loss_hist = train_linear_classifier(svm_loss_vectorized, None,
8                                     data_dict['X_train'],
9                                     data_dict['y_train'],
10                                    learning_rate=3e-11, reg=2.5e4,
11                                    num_iters=1500, verbose=True)
12
13 torch.cuda.synchronize()
14 toc = time.time()
15 print('That took %fs' % (toc - tic))

```

```


iteration 0 / 1500: loss 9.000759
iteration 100 / 1500: loss 9.000790
iteration 200 / 1500: loss 9.000771
iteration 300 / 1500: loss 9.000764
iteration 400 / 1500: loss 9.000766
iteration 500 / 1500: loss 9.000767
iteration 600 / 1500: loss 9.000770
iteration 700 / 1500: loss 9.000748
iteration 800 / 1500: loss 9.000766
iteration 900 / 1500: loss 9.000761
iteration 1000 / 1500: loss 9.000764
iteration 1100 / 1500: loss 9.000766

```

```
iteration 1200 / 1500: loss 9.000779  
iteration 1300 / 1500: loss 9.000757  
iteration 1400 / 1500: loss 9.000749  
That took 1.221067s
```

A useful debugging strategy is to plot the loss as a function of iteration number. In this case it seems our hyperparameters are not good, since the training loss is not decreasing very fast.

```
1 plt.plot(loss_hist, 'o')  
2 plt.xlabel('Iteration number')  
3 plt.ylabel('Loss value')  
4 plt.show()
```



Let's move on to the prediction stage.

```
1 def predict_linear_classifier(W, X):
2     """
3     Use the trained weights of this linear classifier to predict labels for
4     data points.
5
6     Inputs:
7     - W: A PyTorch tensor of shape (D, C), containing weights of a model
8     - X: A PyTorch tensor of shape (N, D) containing training data; there are N
9       training samples each of dimension D.
10
11     Returns:
12     - y_pred: PyTorch int64 tensor of shape (N,) giving predicted labels for each
13       element of X. Each element of y_pred should be between 0 and C - 1.
14     """
15     y_pred = torch.zeros(X.shape[0])
16     #####
17     # TODO:                                     #
18     # Implement this method. Store the predicted labels in y_pred.             #
19     #####
20     # Replace "pass" statement with your code
21     class_preds = X.mm(W)
22     y_pred = torch.argmax(class_preds, dim=1)
23     #####
24     #                                     END OF YOUR CODE                                     #
25     #####
26     return y_pred
```

Then, let's evaluate the performance our trained model on both the training and validation set. You should see validation accuracy less than 10%.

```
1 # evaluate the performance on both the training and validation set
2 y_train_pred = predict_linear_classifier(W, data_dict['X_train'])
3 train_acc = 100.0 * (data_dict['y_train'] == y_train_pred).float().mean().item()
4 print('Training accuracy: %.2f%%' % train_acc)
5 y_val_pred = predict_linear_classifier(W, data_dict['X_val'])
6 val_acc = 100.0 * (data_dict['y_val'] == y_val_pred).float().mean().item()
7 print('Validation accuracy: %.2f%%' % val_acc)
```

```
Training accuracy: 10.24%
Validation accuracy: 10.20%
```

Unfortunately, the performance of our initial model is quite bad. To find a better hyperparameters, let's first modularize the functions that we've implemented.

```
1 # Note: We will re-use `LinearClassifier` in Softmax section
2 class LinearClassifier(object):
3
4     def __init__(self):
5         self.W = None
6
7     def train(self, X_train, y_train, learning_rate=1e-3, reg=1e-5, num_iters=100,
8              batch_size=200, verbose=False):
9         train_args = (self.loss, self.W, X_train, y_train, learning_rate, reg,
10                      num_iters, batch_size, verbose)
11         self.W, loss_history = train_linear_classifier(*train_args)
12         return loss_history
13
14     def predict(self, X):
15         return predict_linear_classifier(self.W, X)
16
17     def loss(self, W, X_batch, y_batch, reg):
18         """
19         Compute the loss function and its derivative.
20         Subclasses will override this.
21
22         Inputs:
23         - W: A PyTorch tensor of shape (D, C) containing (trained) weight of a model.
24         - X_batch: A PyTorch tensor of shape (N, D) containing a minibatch of N
25           data points; each point has dimension D.
26         - y_batch: A PyTorch tensor of shape (N,) containing labels for the minibatch.
27         - reg: (float) regularization strength.
28
29         Returns: A tuple containing:
30         - loss as a single float
31         - gradient with respect to self.W; an tensor of the same shape as W
32         """
33         pass
34     def _loss(self, X_batch, y_batch, reg):
35         self.loss(self.W, X_batch, y_batch, reg)
36
37
38 class LinearSVM(LinearClassifier):
39     """ A subclass that uses the Multiclass SVM loss function """
40
41     def loss(self, W, X_batch, y_batch, reg):
```

```
42     return svm_loss_vectorized(W, X_batch, y_batch, reg)
```

Now, please use the validation set to tune hyperparameters (regularization strength and learning rate). You should experiment with different ranges for the learning rates and regularization strengths.

To get full credit for the assignment your best model found through cross-validation should achieve an accuracy of at least 37% on the validation set.

(Our best model got over 40% -- did you beat us?)

```
1 # results is dictionary mapping tuples of the form
2 # (learning_rate, regularization_strength) to tuples of the form
3 # (training_accuracy, validation_accuracy). The accuracy is simply the fraction
4 # of data points that are correctly classified.
5 results = {}
6 best_val = -1 # The highest validation accuracy that we have seen so far.
7 best_svm = None # The LinearSVM object that achieved the highest validation rate.
8 learning_rates = [1e-1, 1e-2, 1e-3, 1e-4, 1e-5] # learning rate candidates, e.g. [1e-3, 1e-2, ...]
9 regularization_strengths = [1.0, 1e-1, 1e-2, 1e-3] # regularization strengths candidates e.g. [1e0, 1e1, ...]
10
11 #####
12 # TODO: #
13 # Write code that chooses the best hyperparameters by tuning on the validation #
14 # set. For each combination of hyperparameters, train a linear SVM on the #
15 # training set, compute its accuracy on the training and validation sets, and #
16 # store these numbers in the results dictionary. In addition, store the best #
17 # validation accuracy in best_val and the LinearSVM object that achieves this #
18 # accuracy in best_svm. #
19 # #
20 # Hint: You should use a small value for num_iters as you develop your #
21 # validation code so that the SVMs don't take much time to train; once you are #
22 # confident that your validation code works, you should rerun the validation #
23 # code with a larger value for num_iters. #
24 #####
25 # Replace "pass" statement with your code
26 from scipy.stats import loguniform
27 #grid_search = [ (lr,rg) for lr in learning_rates for rg in regularization_strengths ]
28 X_train=data_dict['X_train'].double()
29 y_train=data_dict['y_train']
30 X_val=data_dict['X_val'].double()
31 y_val=data_dict['y_val']
32 tries = 50
33 lrs = loguniform(1, 1000).rvs(size=tries) / 1e4
34 rgs = loguniform(1, 1000).rvs(size=tries) / 1e3
```

```

34 rgs = logarithmic(1, 1000).rs(size=100) / 100
35 for lr, rg in zip(lrs, rgs):
36     # Create a new SVM instance
37     svm = LinearSVM()
38     # Train the model with current parameters
39     train_loss = svm.train(X_train, y_train, learning_rate=lr, reg=rg,
40                             num_iters=5000, batch_size=500, verbose=False)
41     # Predict values for training set
42     y_train_pred = svm.predict(X_train)
43     # Calculate accuracy
44     train_accuracy = torch.mean((y_train_pred == y_train).float())
45     # Predict values for validation set
46     y_val_pred = svm.predict(X_val)
47     # Calculate accuracy
48     val_accuracy = torch.mean((y_val_pred == y_val).float())
49     # Save results
50     results[(lr, rg)] = (train_accuracy.cpu(), val_accuracy.cpu())
51     if best_val < val_accuracy:
52         best_val = val_accuracy
53         best_svm = svm
54
55 #####
56 #                               END OF YOUR CODE                               #
57 #####
58
59 # Print out results.
60 for lr, reg in sorted(results):
61     train_accuracy, val_accuracy = results[(lr, reg)]
62     print('lr %e reg %e train accuracy: %f val accuracy: %f' % (
63         lr, reg, train_accuracy, val_accuracy))
64
65 print('best validation accuracy achieved during cross-validation: %f' % best_val)

```

```

lr 1.127664e-04 reg 9.421768e-01 train accuracy: 0.336265 val accuracy: 0.345000
lr 1.408724e-04 reg 7.210319e-02 train accuracy: 0.336816 val accuracy: 0.350000
lr 1.513240e-04 reg 4.319945e-02 train accuracy: 0.338061 val accuracy: 0.344000
lr 1.778549e-04 reg 6.947590e-02 train accuracy: 0.336939 val accuracy: 0.343000
lr 2.198666e-04 reg 3.582803e-02 train accuracy: 0.339816 val accuracy: 0.351000
lr 2.691032e-04 reg 6.231922e-03 train accuracy: 0.338796 val accuracy: 0.349000
lr 2.776062e-04 reg 4.874833e-02 train accuracy: 0.339796 val accuracy: 0.352000
lr 3.442511e-04 reg 3.622188e-02 train accuracy: 0.338551 val accuracy: 0.347000
lr 3.683242e-04 reg 1.439797e-03 train accuracy: 0.337041 val accuracy: 0.346000
lr 3.722480e-04 reg 4.319655e-03 train accuracy: 0.337878 val accuracy: 0.351000
lr 5.163707e-04 reg 2.847976e-01 train accuracy: 0.342959 val accuracy: 0.351000
lr 9.213078e-04 reg 2.835140e-02 train accuracy: 0.340592 val accuracy: 0.347000
lr 1.256022e-03 reg 2.996293e-03 train accuracy: 0.338163 val accuracy: 0.342000
lr 1.355538e-03 reg 1.208743e-03 train accuracy: 0.339122 val accuracy: 0.337000

```

```

lr 1.553542e-03 reg 4.525897e-01 train accuracy: 0.343306 val accuracy: 0.347000
lr 1.620250e-03 reg 1.341086e-02 train accuracy: 0.338775 val accuracy: 0.339000
lr 2.406219e-03 reg 2.411564e-02 train accuracy: 0.341531 val accuracy: 0.340000
lr 4.099733e-03 reg 3.040770e-03 train accuracy: 0.340122 val accuracy: 0.328000
lr 5.069754e-03 reg 4.442841e-01 train accuracy: 0.342163 val accuracy: 0.355000
lr 5.402859e-03 reg 6.148845e-02 train accuracy: 0.340122 val accuracy: 0.344000
lr 5.481807e-03 reg 9.280047e-01 train accuracy: 0.334980 val accuracy: 0.339000
lr 6.080236e-03 reg 2.748448e-01 train accuracy: 0.340878 val accuracy: 0.354000
lr 6.819868e-03 reg 9.243849e-02 train accuracy: 0.341878 val accuracy: 0.349000
lr 8.544718e-03 reg 2.181268e-01 train accuracy: 0.336367 val accuracy: 0.351000
lr 1.207872e-02 reg 2.890302e-03 train accuracy: 0.341061 val accuracy: 0.339000
lr 1.258856e-02 reg 1.746623e-03 train accuracy: 0.336122 val accuracy: 0.336000
lr 1.469194e-02 reg 4.662686e-02 train accuracy: 0.334061 val accuracy: 0.340000
lr 1.966603e-02 reg 7.831266e-02 train accuracy: 0.340775 val accuracy: 0.358000
lr 2.022643e-02 reg 2.445158e-03 train accuracy: 0.328163 val accuracy: 0.330000
lr 2.110290e-02 reg 6.550421e-03 train accuracy: 0.332551 val accuracy: 0.341000
lr 2.235903e-02 reg 4.692708e-03 train accuracy: 0.338082 val accuracy: 0.332000
lr 2.343567e-02 reg 2.524965e-02 train accuracy: 0.335816 val accuracy: 0.338000
lr 2.461006e-02 reg 2.805650e-03 train accuracy: 0.336163 val accuracy: 0.349000
lr 2.528027e-02 reg 4.057072e-03 train accuracy: 0.339673 val accuracy: 0.333000
lr 2.578877e-02 reg 3.821987e-03 train accuracy: 0.335102 val accuracy: 0.341000
lr 2.664693e-02 reg 1.479884e-03 train accuracy: 0.340143 val accuracy: 0.354000
lr 2.861773e-02 reg 1.401769e-02 train accuracy: 0.337000 val accuracy: 0.345000
lr 3.098276e-02 reg 2.312447e-03 train accuracy: 0.337020 val accuracy: 0.345000
lr 3.317256e-02 reg 8.643469e-03 train accuracy: 0.332653 val accuracy: 0.347000
lr 3.547755e-02 reg 2.923620e-03 train accuracy: 0.331633 val accuracy: 0.358000
lr 4.168228e-02 reg 3.144347e-03 train accuracy: 0.337265 val accuracy: 0.347000
lr 4.770873e-02 reg 4.543722e-01 train accuracy: 0.276388 val accuracy: 0.286000
lr 5.213762e-02 reg 9.900882e-02 train accuracy: 0.296735 val accuracy: 0.303000
lr 6.022935e-02 reg 7.677022e-01 train accuracy: 0.258143 val accuracy: 0.277000
lr 6.236284e-02 reg 3.373487e-03 train accuracy: 0.312571 val accuracy: 0.332000
lr 6.445422e-02 reg 2.047562e-01 train accuracy: 0.289429 val accuracy: 0.306000
lr 7.973036e-02 reg 1.747381e-03 train accuracy: 0.321367 val accuracy: 0.344000
lr 8.915393e-02 reg 5.124749e-01 train accuracy: 0.262388 val accuracy: 0.279000
lr 9.185618e-02 reg 2.480028e-03 train accuracy: 0.250163 val accuracy: 0.247000
lr 9.665097e-02 reg 3.895687e-03 train accuracy: 0.308122 val accuracy: 0.316000
best validation accuracy achieved during cross-validation: 0.358000

```

Visualize the cross-validation results. You can use this as a debugging tool – after examining the cross-validation results here, you may want to go back and rerun your cross-validation from above.

```

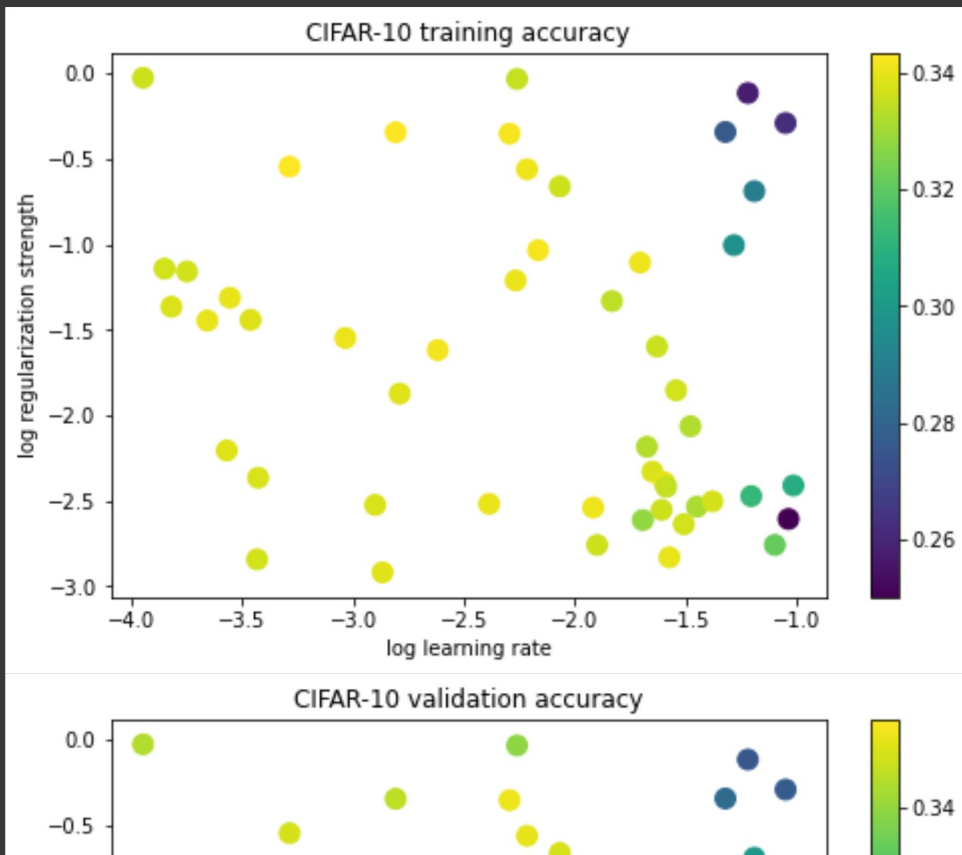
1 x_scatter = [math.log10(x[0]) for x in results]
2 y_scatter = [math.log10(x[1]) for x in results]

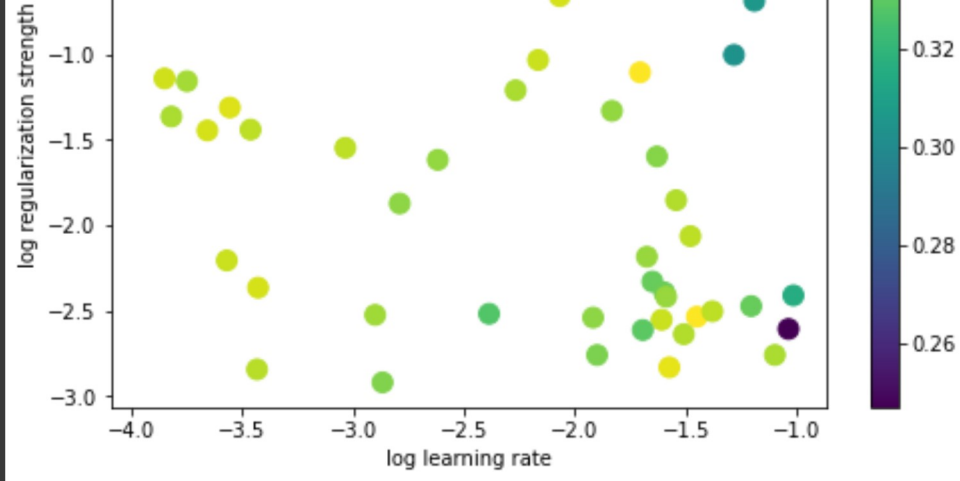
```

```

3
4 # plot training accuracy
5 marker_size = 100
6 colors = [results[x][0] for x in results]
7 plt.scatter(x_scatter, y_scatter, marker_size, c=colors, cmap='viridis')
8 plt.colorbar()
9 plt.xlabel('log learning rate')
10 plt.ylabel('log regularization strength')
11 plt.title('CIFAR-10 training accuracy')
12 plt.gcf().set_size_inches(8, 5)
13 plt.show()
14
15 # plot validation accuracy
16 colors = [results[x][1] for x in results] # default size of markers is 20
17 plt.scatter(x_scatter, y_scatter, marker_size, c=colors, cmap='viridis')
18 plt.colorbar()
19 plt.xlabel('log learning rate')
20 plt.ylabel('log regularization strength')
21 plt.title('CIFAR-10 validation accuracy')
22 plt.gcf().set_size_inches(8, 5)
23 plt.show()

```





Evaluate the best svm on test set. To get full credit for the assignment you should achieve a test-set accuracy above 35%.

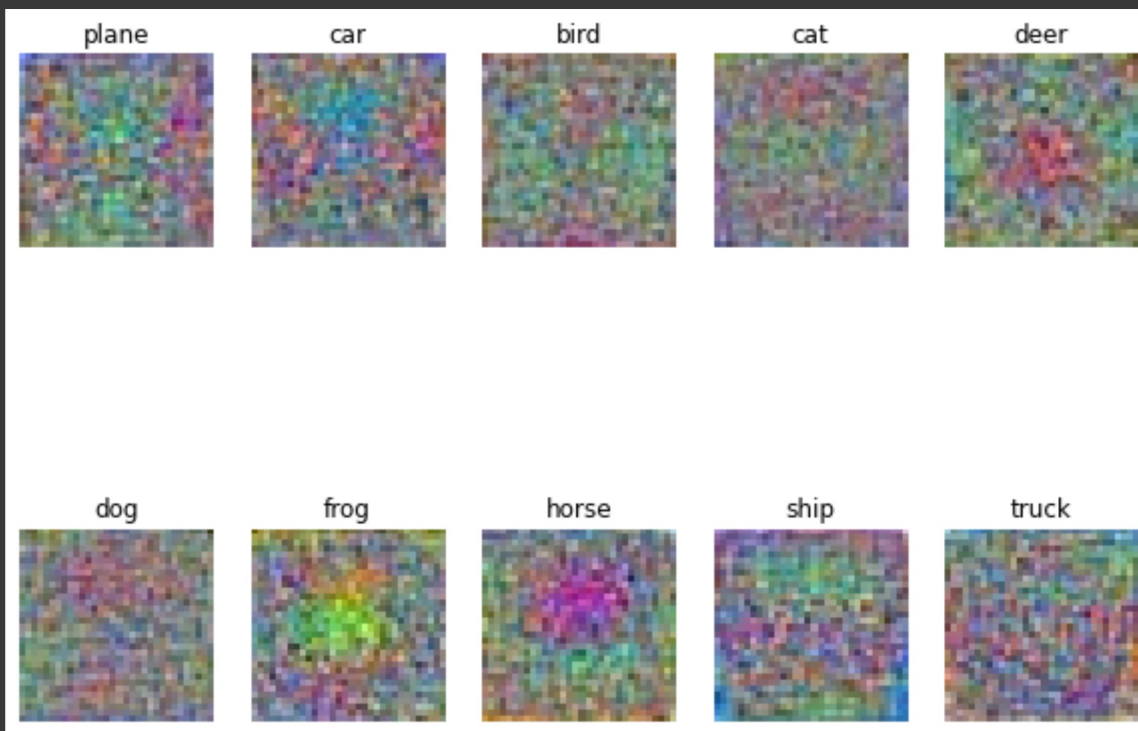
(Our best was over 38% -- did you beat us?)

```
1 y_test_pred = best_svm.predict(data_dict['X_test'].double())
2 test_accuracy = torch.mean((data_dict['y_test'] == y_test_pred).float())
3 print('linear SVM on raw pixels final test set accuracy: %f' % test_accuracy)
```

linear SVM on raw pixels final test set accuracy: 0.319600

Visualize the learned weights for each class. Depending on your choice of learning rate and regularization strength, these may or may not be nice to look at.

```
1 w = best_svm.W[:-1,:] # strip out the bias
2 w = w.reshape(3, 32, 32, 10)
3 w = w.transpose(0, 2).transpose(1, 0)
4
5 w_min, w_max = torch.min(w), torch.max(w)
6 classes = ['plane', 'car', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']
7 for i in range(10):
8     plt.subplot(2, 5, i + 1)
9
10    # Rescale the weights to be between 0 and 255
11    wimg = 255.0 * (w[:, :, :, i].squeeze() - w_min) / (w_max - w_min)
12    plt.imshow(wimg.type(torch.uint8).cpu())
13    plt.axis('off')
14    plt.title(classes[i])
```



Softmax Classifier

Similar to the SVM, you will:

- implement a fully-vectorized **loss function** for the Softmax classifier
- implement the fully-vectorized expression for its **analytic gradient**
- **check your implementation** with numerical gradient
- use a validation set to **tune the learning rate and regularization** strength
- **optimize** the loss function with **SGD**
- **visualize** the final learned weights

First, let's start from implementing the naive softmax loss function with nested loops.

```
1 def softmax_loss_naive(W, X, y, reg):
2     """
3     Softmax loss function, naive implementation (with loops). When you implement
```



```

4 the regularization over W, please DO NOT multiply the regularization term by
5 1/2 (no coefficient).
6
7 Inputs have dimension D, there are C classes, and we operate on minibatches
8 of N examples.
9
10 Inputs:
11 - W: A PyTorch tensor of shape (D, C) containing weights.
12 - X: A PyTorch tensor of shape (N, D) containing a minibatch of data.
13 - y: A PyTorch tensor of shape (N,) containing training labels; y[i] = c means
14     that X[i] has label c, where 0 <= c < C.
15 - reg: (float) regularization strength
16
17 Returns a tuple of:
18 - loss as single float
19 - gradient with respect to weights W; an tensor of same shape as W
20 """
21 # Initialize the loss and gradient to zero.
22 loss = 0.0
23 grad = torch.zeros_like(W)
24
25 #####
26 # TODO: Compute the softmax loss and its gradient using explicit loops.      #
27 # Store the loss in loss and the gradient in grad. If you are not careful    #
28 # here, it is easy to run into numeric instability (Check Numeric Stability  #
29 # in http://cs231n.github.io/linear-classify/). Plus, don't forget the      #
30 # regularization!                                                            #
31 #####
32 # Replace "pass" statement with your code
33 D, C = W.shape
34 N, D = X.shape
35 f = X.mm(W) # Score, (N,C).
36 for i in range(N):
37     f[i] -= torch.max(f[i]) # Numeric stability.
38     fi_e_s = f[i].exp().sum()
39     loss += -f[i][y[i]] + fi_e_s.log()
40     for j in range(C):
41         grad[:, j] += X[i,:] * f[i,j].exp() / fi_e_s
42         if y[i] == j:
43             grad[:, j] -= X[i,:]
44 loss /= N # Mean.
45 loss += reg * W.pow(2).sum()
46 grad /= N
47 grad += reg * 2 * W
48
49 #####
50 #                                     END OF YOUR CODE                                     #

```

```

51 #####
52
53 return loss, grad

```

As a sanity check to see whether we have implemented the loss correctly, run the softmax classifier with a small random weight matrix and no regularization. You should see loss near $\log(10) = 2.3$

```

1 # Generate a random softmax weight tensor and use it to compute the loss.
2 coutils.utils.fix_random_seed()
3 W = 0.0001 * torch.randn(3073, 10, device=data_dict['X_val'].device).double()
4
5 X_batch = data_dict['X_val'][:128].double()
6 y_batch = data_dict['y_val'][:128]
7
8 # Complete the implementation of softmax_loss_naive and implement a (naive)
9 # version of the gradient that uses nested loops.
10 loss, grad = softmax_loss_naive(W, X_batch, y_batch, reg=0.0)
11
12 # As a rough sanity check, our loss should be something close to log(10.0).
13 print('loss: %f' % loss)
14 print('sanity check: %f' % (math.log(10.0)))

```

```

    loss: 2.302600
    sanity check: 2.302585

```

Next, we use gradient checking to debug the analytic gradient of our naive softmax loss function. If you've implemented the gradient correctly, you should see relative errors less than `1e-6`.

```

1 coutils.utils.fix_random_seed()
2 W = 0.0001 * torch.randn(3073, 10, device=data_dict['X_val'].device).double()
3 X_batch = data_dict['X_val'][:128].double()
4 y_batch = data_dict['y_val'][:128]
5
6 loss, grad = softmax_loss_naive(W, X_batch, y_batch, reg=0.0)
7
8 f = lambda w: softmax_loss_naive(w, X_batch, y_batch, reg=0.0)[0]
9 grad_check_sparse(f, W, grad, 10)

```

```

    numerical: 0.008387 analytic: 0.008387, relative error: 2.690277e-07
    numerical: 0.009227 analytic: 0.009227, relative error: 1.339656e-07
    numerical: -0.002471 analytic: -0.002471, relative error: 1.718331e-07

```

```
numerical: -0.003144 analytic: -0.003144, relative error: 2.403080e-06
numerical: 0.006011 analytic: 0.006011, relative error: 6.813253e-08
numerical: 0.005936 analytic: 0.005936, relative error: 2.473992e-07
numerical: 0.015703 analytic: 0.015703, relative error: 2.149831e-08
numerical: 0.006452 analytic: 0.006452, relative error: 2.068055e-09
numerical: -0.015533 analytic: -0.015533, relative error: 1.855205e-07
numerical: -0.010170 analytic: -0.010170, relative error: 4.657956e-07
```

Let's perform another gradient check with regularization enabled. Again you should see relative errors less than `1e-6`.

```
1 cutils.utils.fix_random_seed()
2 W = 0.0001 * torch.randn(3073, 10, device=data_dict['X_val'].device).double()
3 reg = 10.0
4
5 X_batch = data_dict['X_val'][:128].double()
6 y_batch = data_dict['y_val'][:128]
7
8 loss, grad = softmax_loss_naive(W, X_batch, y_batch, reg)
9
10 f = lambda w: softmax_loss_naive(w, X_batch, y_batch, reg)[0]
11 grad_check_sparse(f, W, grad, 10)
```

```
numerical: 0.007517 analytic: 0.007517, relative error: 3.763147e-07
numerical: 0.008167 analytic: 0.008167, relative error: 1.238116e-07
numerical: -0.002551 analytic: -0.002551, relative error: 3.147037e-08
numerical: -0.000841 analytic: -0.000841, relative error: 8.976250e-06
numerical: 0.002228 analytic: 0.002228, relative error: 4.070641e-07
numerical: 0.005460 analytic: 0.005460, relative error: 1.891176e-07
numerical: 0.015762 analytic: 0.015762, relative error: 7.076879e-08
numerical: 0.007097 analytic: 0.007097, relative error: 1.474077e-07
numerical: -0.015532 analytic: -0.015532, relative error: 1.434386e-07
numerical: -0.011201 analytic: -0.011201, relative error: 3.391712e-07
```

Then, let's move on to the vectorized form

```
1 def softmax_loss_vectorized(W, X, y, reg):
2     """
3     Softmax loss function, vectorized version. When you implement the
4     regularization over W, please DO NOT multiply the regularization term by 1/2
5     (no coefficient).
6
7     Inputs and outputs are the same as softmax loss naive.
```

```

8  """
9  # Initialize the loss and gradient to zero.
10 loss = 0.0
11 grad = torch.zeros_like(W)
12
13 #####
14 # TODO: Compute the softmax loss and its gradient using no explicit loops. #
15 # Store the loss in loss and the gradient in grad. If you are not careful #
16 # here, it is easy to run into numeric instability (Check Numeric Stability #
17 # in http://cs231n.github.io/linear-classify/). Don't forget the #
18 # regularization! #
19 #####
20 # Replace "pass" statement with your code
21 D, C = W.shape
22 N, D = X.shape
23
24 XW = X.mm(W) # Score, (N,C).
25 XW -= XW.max(dim=1, keepdim=True).values # Numeric stability.
26 Y = torch.zeros((N,C), device='cuda').to(XW) # (N,C)
27 Y[range(N), y] = 1
28 YXW = Y * XW # (N, C)
29 YXW_sums = YXW.sum(dim=1, keepdim=True) # (N,1)
30 XWexp = XW.exp() # (N,C)
31 XWexp_s = XWexp.sum(dim=1, keepdim=True) # (N,1)
32 XWexp_s_l = XWexp_s.log()
33 loss = (XWexp_s_l - YXW_sums).sum() / N # (N,1)
34 loss += reg * (W*W).sum()
35
36 XWexp_s_div = 1.0 / ( XWexp_s + 1e-9) # (N,1)
37 grad = XWexp * XWexp_s_div # (N,C)
38 grad = X.T.mm(grad)
39 grad -= X.T.mm(Y)
40 grad /= N
41 grad += 2 * reg * W
42
43 #####
44 #                               END OF YOUR CODE                               #
45 #####
46
47 return loss, grad

```

Now that we have a naive implementation of the softmax loss function and its gradient, implement a vectorized version in `softmax_loss_vectorized`. The two versions should compute the same results, but the vectorized version should be much faster.

The differences between the naive and vectorized losses and gradients should both be less than `1e-6`, and your vectorized implementation

should be at least 100x faster than the naive implementation.

```
1 cutils.utils.fix_random_seed()
2 W = 0.0001 * torch.randn(3073, 10, device=data_dict['X_val'].device)
3 reg = 0.05
4
5 X_batch = data_dict['X_val'][:128]
6 y_batch = data_dict['y_val'][:128]
7
8 # Run and time the naive version
9 torch.cuda.synchronize()
10 tic = time.time()
11 loss_naive, grad_naive = softmax_loss_naive(W, X_batch, y_batch, reg)
12 torch.cuda.synchronize()
13 toc = time.time()
14 ms_naive = 1000.0 * (toc - tic)
15 print('naive loss: %e computed in %fs' % (loss_naive, ms_naive))
16
17 # Run and time the vectorized version
18 torch.cuda.synchronize()
19 tic = time.time()
20 loss_vec, grad_vec = softmax_loss_vectorized(W, X_batch, y_batch, reg)
21 torch.cuda.synchronize()
22 toc = time.time()
23 ms_vec = 1000.0 * (toc - tic)
24 print('vectorized loss: %e computed in %fs' % (loss_vec, ms_vec))
25
26 # we use the Frobenius norm to compare the two versions of the gradient.
27 loss_diff = (loss_naive - loss_vec).abs().item()
28 grad_diff = torch.norm(grad_naive - grad_vec, p='fro')
29 print('Loss difference: %.2e' % loss_diff)
30 print('Gradient difference: %.2e' % grad_diff)
31 print('Speedup: %.2fX' % (ms_naive / ms_vec))
```

```
naive loss: 2.302615e+00 computed in 141.087055s
vectorized loss: 2.302616e+00 computed in 0.938177s
Loss difference: 4.77e-07
Gradient difference: 3.45e-07
Speedup: 150.38X
```

Let's check that your implementation of the softmax loss is numerically stable.

If either of the following print `nan` then you should double-check the numeric stability of your implementations.

```

1 device = data_dict['X_train'].device
2 dtype = torch.float32
3 D = data_dict['X_train'].shape[1]
4 C = 10
5
6 W_ones = torch.ones(D, C, device=device, dtype=dtype)
7 W, loss_hist = train_linear_classifier(softmax_loss_naive, W_ones,
8                                     data_dict['X_train'],
9                                     data_dict['y_train'],
10                                    learning_rate=1e-8, reg=2.5e4,
11                                    num_iters=1, verbose=True)
12
13
14 W_ones = torch.ones(D, C, device=device, dtype=dtype)
15 W, loss_hist = train_linear_classifier(softmax_loss_vectorized, W_ones,
16                                     data_dict['X_train'],
17                                     data_dict['y_train'],
18                                    learning_rate=1e-8, reg=2.5e4,
19                                    num_iters=1, verbose=True)
20

```

```

iteration 0 / 1: loss 768249984.000000
iteration 0 / 1: loss 768249984.000000

```

Now lets train a softmax classifier with some default hyperparameters:

```

1 # fix random seed before we perform this operation
2 coutils.utils.fix_random_seed(10)
3
4 torch.cuda.synchronize()
5 tic = time.time()
6
7 W, loss_hist = train_linear_classifier(softmax_loss_vectorized, None,
8                                     data_dict['X_train'],
9                                     data_dict['y_train'],
10                                    learning_rate=1e-10, reg=2.5e4,
11                                    num_iters=1500, verbose=True)
12
13 torch.cuda.synchronize()
14 toc = time.time()
15 print('That took %fs' % (toc - tic))

```

```


iteration 0 / 1500: loss 2.303346


```

```
iteration 100 / 1500: loss 2.303344
iteration 200 / 1500: loss 2.303344
iteration 300 / 1500: loss 2.303342
iteration 400 / 1500: loss 2.303342
iteration 500 / 1500: loss 2.303342
iteration 600 / 1500: loss 2.303342
iteration 700 / 1500: loss 2.303341
iteration 800 / 1500: loss 2.303338
iteration 900 / 1500: loss 2.303339
iteration 1000 / 1500: loss 2.303339
iteration 1100 / 1500: loss 2.303336
iteration 1200 / 1500: loss 2.303335
iteration 1300 / 1500: loss 2.303337
iteration 1400 / 1500: loss 2.303334
That took 1.048374s
```

Plot the loss curve:

```
1 plt.plot(loss_hist, 'o')
2 plt.xlabel('Iteration number')
3 plt.ylabel('Loss value')
4 plt.show()
```





Let's compute the accuracy of current model. It should be less than 10%.

```
1 # evaluate the performance on both the training and validation set
2 y_train_pred = predict_linear_classifier(W, data_dict['X_train'])
3 train_acc = 100.0 * (data_dict['y_train'] == y_train_pred).float().mean().item()
4 print('training accuracy: %.2f%%' % train_acc)
5 y_val_pred = predict_linear_classifier(W, data_dict['X_val'])
6 val_acc = 100.0 * (data_dict['y_val'] == y_val_pred).float().mean().item()
7 print('validation accuracy: %.2f%%' % val_acc)
```

```
training accuracy: 8.43%
validation accuracy: 8.00%
```

Now use the validation set to tune hyperparameters (regularization strength and learning rate). You should experiment with different ranges for the learning rates and regularization strengths.

To get full credit for the assignment, your best model found through cross-validation should achieve an accuracy above 0.37 on the validation set.

(Our best model was above 0.40 -- did you beat us?)

```
1 class Softmax(LinearClassifier):
2     """ A subclass that uses the Softmax + Cross-entropy loss function """
3     def loss(self, W, X_batch, y_batch, reg):
4         return softmax_loss_vectorized(W, X_batch, y_batch, reg)
```

```
1 results = {}
2 best_val = -1
3 best_softmax = None
4
5 learning_rates = [1e-1, 1e-2, 1e-3, 1e-4] # learning rate candidates, e.g. [1e-3, 1e-2, ...]
6 regularization_strengths = [1e-1, 1e-2, 1e-3] # regularization strengths candidates e.g. [1e0, 1e1, ...]
```



```

7
8 # As before, store your cross-validation results in this dictionary.
9 # The keys should be tuples of (learning_rate, regularization_strength) and
10 # the values should be tuples (train_accuracy, val_accuracy)
11 results = {}
12
13 #####
14 # TODO: #
15 # Use the validation set to set the learning rate and regularization strength. #
16 # This should be similar to the cross-validation that you used for the SVM, #
17 # but you may need to select different hyperparameters to achieve good #
18 # performance with the softmax classifier. Save your best trained softmax #
19 # classifier in best_softmax. #
20 #####
21 # Replace "pass" statement with your code
22 from scipy.stats import loguniform
23 X_train=data_dict['X_train'].double()
24 y_train=data_dict['y_train']
25 X_val=data_dict['X_val'].double()
26 y_val=data_dict['y_val']
27 tries = 50
28 lrs = loguniform(1, 1000).rvs(size=tries) / 1e4
29 rgs = loguniform(1, 1000).rvs(size=tries) / 1e3
30 for lr, rg in zip(lrs, rgs):
31     softmax = Softmax()
32     # Train the model with current parameters
33     train_loss = softmax.train(X_train, y_train, learning_rate=lr, reg=rg,
34                               num_iters=5000, batch_size=200, verbose=False)
35     # Predict values for training set
36     y_train_pred = softmax.predict(X_train)
37     # Calculate accuracy
38     train_accuracy = torch.mean((y_train_pred == y_train).float())
39     # Predict values for validation set
40     y_val_pred = softmax.predict(X_val)
41     # Calculate accuracy
42     val_accuracy = torch.mean((y_val_pred == y_val).float())
43     # Save results
44     results[(lr,rg)] = (train_accuracy.cpu(), val_accuracy.cpu())
45     if best_val < val_accuracy:
46         best_val = val_accuracy
47         best_softmax = softmax
48
49 #####
50 #                               END OF YOUR CODE #
51 #####
52
53 # Print out results.

```

```
54 for lr, reg in sorted(results):
55     train_accuracy, val_accuracy = results[(lr, reg)]
56     print('lr %e reg %e train accuracy: %f val accuracy: %f' % (
57         lr, reg, train_accuracy, val_accuracy))
58
59 print('best validation accuracy achieved during cross-validation: %f' % best_val)
```

```
lr 1.279019e-04 reg 3.018058e-03 train accuracy: 0.294735 val accuracy: 0.302000
lr 1.775872e-04 reg 1.846863e-01 train accuracy: 0.303714 val accuracy: 0.314000
lr 1.963025e-04 reg 2.672647e-01 train accuracy: 0.305469 val accuracy: 0.313000
lr 2.028232e-04 reg 4.778375e-03 train accuracy: 0.313020 val accuracy: 0.322000
lr 2.509903e-04 reg 7.782972e-01 train accuracy: 0.297204 val accuracy: 0.306000
lr 2.603718e-04 reg 6.294445e-01 train accuracy: 0.302796 val accuracy: 0.310000
lr 2.632870e-04 reg 3.464881e-01 train accuracy: 0.310490 val accuracy: 0.319000
lr 2.701229e-04 reg 1.136813e-01 train accuracy: 0.321347 val accuracy: 0.327000
lr 3.154744e-04 reg 7.653344e-03 train accuracy: 0.331061 val accuracy: 0.338000
lr 3.465328e-04 reg 4.853213e-01 train accuracy: 0.313918 val accuracy: 0.320000
lr 4.477644e-04 reg 3.506859e-01 train accuracy: 0.324347 val accuracy: 0.327000
lr 5.373772e-04 reg 1.219566e-01 train accuracy: 0.342347 val accuracy: 0.342000
lr 6.382205e-04 reg 3.134151e-01 train accuracy: 0.327306 val accuracy: 0.336000
lr 7.051506e-04 reg 1.780770e-03 train accuracy: 0.361245 val accuracy: 0.374000
lr 9.481275e-04 reg 1.274656e-03 train accuracy: 0.369265 val accuracy: 0.379000
lr 1.085796e-03 reg 4.729068e-03 train accuracy: 0.372633 val accuracy: 0.377000
lr 1.217774e-03 reg 1.777550e-02 train accuracy: 0.373796 val accuracy: 0.383000
lr 1.514552e-03 reg 9.742901e-01 train accuracy: 0.296469 val accuracy: 0.306000
lr 1.758840e-03 reg 2.924202e-03 train accuracy: 0.384775 val accuracy: 0.395000
lr 1.895918e-03 reg 6.022217e-02 train accuracy: 0.374327 val accuracy: 0.383000
lr 1.951465e-03 reg 5.981318e-02 train accuracy: 0.375980 val accuracy: 0.388000
lr 2.397706e-03 reg 9.393601e-02 train accuracy: 0.369041 val accuracy: 0.373000
lr 2.866916e-03 reg 6.235058e-01 train accuracy: 0.307959 val accuracy: 0.316000
lr 3.099653e-03 reg 7.842023e-03 train accuracy: 0.395673 val accuracy: 0.392000
lr 3.161284e-03 reg 5.110237e-03 train accuracy: 0.397041 val accuracy: 0.402000
lr 3.745480e-03 reg 1.302576e-02 train accuracy: 0.396449 val accuracy: 0.407000
lr 4.105544e-03 reg 1.710288e-01 train accuracy: 0.355245 val accuracy: 0.362000
lr 4.396318e-03 reg 2.506823e-03 train accuracy: 0.403918 val accuracy: 0.404000
lr 6.328190e-03 reg 2.653711e-01 train accuracy: 0.337041 val accuracy: 0.346000
lr 6.804145e-03 reg 2.443296e-01 train accuracy: 0.343531 val accuracy: 0.349000
lr 8.221490e-03 reg 1.455287e-01 train accuracy: 0.360224 val accuracy: 0.370000
lr 9.341744e-03 reg 1.485272e-01 train accuracy: 0.355735 val accuracy: 0.362000
lr 1.235193e-02 reg 2.812736e-02 train accuracy: 0.397408 val accuracy: 0.399000
lr 1.305006e-02 reg 1.089116e-01 train accuracy: 0.363612 val accuracy: 0.374000
lr 1.488446e-02 reg 1.544398e-03 train accuracy: 0.422327 val accuracy: 0.416000
lr 1.531618e-02 reg 1.022489e-02 train accuracy: 0.413898 val accuracy: 0.405000
lr 1.703540e-02 reg 1.715669e-02 train accuracy: 0.404959 val accuracy: 0.410000
lr 1.858392e-02 reg 1.608448e-01 train accuracy: 0.354408 val accuracy: 0.366000
lr 2.128294e-02 reg 4.691005e-01 train accuracy: 0.321837 val accuracy: 0.327000
```

```

lr 2.666534e-02 reg 3.691153e-01 train accuracy: 0.321939 val accuracy: 0.340000
lr 2.713348e-02 reg 9.425434e-03 train accuracy: 0.414429 val accuracy: 0.407000
lr 4.342613e-02 reg 8.371512e-03 train accuracy: 0.415490 val accuracy: 0.403000
lr 4.498635e-02 reg 2.787238e-03 train accuracy: 0.427878 val accuracy: 0.415000
lr 4.661909e-02 reg 3.357158e-02 train accuracy: 0.379816 val accuracy: 0.395000
lr 4.719540e-02 reg 1.858477e-03 train accuracy: 0.431653 val accuracy: 0.408000
lr 5.850926e-02 reg 3.924808e-01 train accuracy: 0.313755 val accuracy: 0.318000
lr 6.513662e-02 reg 1.570185e-03 train accuracy: 0.432082 val accuracy: 0.410000
lr 6.554669e-02 reg 5.608152e-02 train accuracy: 0.373020 val accuracy: 0.369000
lr 7.526105e-02 reg 4.398645e-01 train accuracy: 0.321796 val accuracy: 0.316000
lr 8.655925e-02 reg 1.328476e-02 train accuracy: 0.406551 val accuracy: 0.395000
best validation accuracy achieved during cross-validation: 0.416000


```


Run the following to visualize your cross-validation results:

```

1 x_scatter = [math.log10(x[0]) for x in results]
2 y_scatter = [math.log10(x[1]) for x in results]
3
4 # plot training accuracy
5 marker_size = 100
6 colors = [results[x][0] for x in results]
7 plt.scatter(x_scatter, y_scatter, marker_size, c=colors, cmap='viridis')
8 plt.colorbar()
9 plt.xlabel('log learning rate')
10 plt.ylabel('log regularization strength')
11 plt.title('CIFAR-10 training accuracy')
12 plt.gcf().set_size_inches(8, 5)
13 plt.show()
14
15 # plot validation accuracy
16 colors = [results[x][1] for x in results] # default size of markers is 20
17 plt.scatter(x_scatter, y_scatter, marker_size, c=colors, cmap='viridis')
18 plt.colorbar()
19 plt.xlabel('log learning rate')
20 plt.ylabel('log regularization strength')
21 plt.title('CIFAR-10 validation accuracy')
22 plt.gcf().set_size_inches(8, 5)
23 plt.show()

```





Then, evaluate the performance of your best model on test set. To get full credit for this assignment you should achieve a test-set accuracy above 0.36.

(Our best was just over 0.40 -- did you beat us?)

```
1 y_test_pred = best_softmax.predict(data_dict['X_test'].double())
2 test_accuracy = torch.mean((data_dict['y_test'] == y_test_pred).float())
3 print('softmax on raw pixels final test set accuracy: %f' % (test_accuracy, ))
```

softmax on raw pixels final test set accuracy: 0.409800

Finally, let's visualize the learned weights for each class

```
1 w = best_softmax.W[:-1,:] # strip out the bias
2 w = w.reshape(3, 32, 32, 10)
3 w = w.transpose(0, 2).transpose(1, 0)
4
5 w_min, w_max = torch.min(w), torch.max(w)
6
7 classes = ['plane', 'car', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']
8 for i in range(10):
9     plt.subplot(2, 5, i + 1)
10
11     # Rescale the weights to be between 0 and 255
12     wimg = 255.0 * (w[:, :, :, i].squeeze() - w_min) / (w_max - w_min)
13     plt.imshow(wimg.type(torch.uint8).cpu())
14     plt.axis('off')
15     plt.title(classes[i])
```

