

▼ EECS 498-007/598-005 Assignment 2-2: Two Layer Neural Network

Before we start, please put your name and UMID in following format

: Firstname LASTNAME, #00000000 // e.g.) Justin JOHNSON, #12345678

Artyom KARPOV (www.artkpv.net)

Implementing a Neural Network

In this exercise we will develop a neural network with fully-connected layers to perform classification, and test it out on the CIFAR-10 dataset.

We train the network with a softmax loss function and L2 regularization on the weight matrices. The network uses a ReLU nonlinearity after the first fully connected layer.

In other words, the network has the following architecture:

input - fully connected layer - ReLU - fully connected layer - softmax

The outputs of the second fully-connected layer are the scores for each class.

Note: When you implement the regularization over W, please DO NOT multiply the regularization term by 1/2 (no coefficient).

▼ Install starter code

We will continue using the utility functions that we've used for Assignment 1: [coutils package](#). Run this cell to download and install it.

```
1 !pip install git+https://github.com/deepvision-class/starter-code
```

```
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/public/simple/
```

```
Collecting git+https://github.com/deepvision-class/starter-code
```

```
  Cloning https://github.com/deepvision-class/starter-code to /tmp/pip-req-build-tntt7vf3
```

```
    Running command git clone -q https://github.com/deepvision-class/starter-code /tmp/pip-req-build-tntt7vf3
```

```
Requirement already satisfied: pydrive in /usr/local/lib/python3.7/dist-packages (from Colab-Utils==0.1.dev0) (1.3.1)
```

```
Requirement already satisfied: google-api-python-client>=1.2 in /usr/local/lib/python3.7/dist-packages (from pydrive->Colab-Utils==0.1.de
```

```
Requirement already satisfied: oauth2client>=4.0.0 in /usr/local/lib/python3.7/dist-packages (from pydrive->Colab-Utils==0.1.dev0) (4.1.3)
```

✓ 0 сек. выполнено в 08:35

```
Requirement already satisfied: google-auth<http://lib>=0.0.3 in /usr/local/lib/python3/dist-packages (from google-api-python-client>=1.2->pydrive)
Requirement already satisfied: google-api-core<3dev,>=1.21.0 in /usr/local/lib/python3.7/dist-packages (from google-api-python-client>=1.2->pydrive)
Requirement already satisfied: google-auth<3dev,>=1.16.0 in /usr/local/lib/python3.7/dist-packages (from google-api-python-client>=1.2->pydrive)
Requirement already satisfied: uritemplate<4dev,>=3.0.0 in /usr/local/lib/python3.7/dist-packages (from google-api-python-client>=1.2->pydrive)
Requirement already satisfied: six<2dev,>=1.13.0 in /usr/local/lib/python3.7/dist-packages (from google-api-python-client>=1.2->pydrive)
Requirement already satisfied: httplib2<1dev,>=0.15.0 in /usr/local/lib/python3.7/dist-packages (from google-api-python-client>=1.2->pydrive)
Requirement already satisfied: protobuf<4.0.0dev,>=3.12.0 in /usr/local/lib/python3.7/dist-packages (from google-api-core<3dev,>=1.21.0->google-api-python-client)
Requirement already satisfied: googleapis-common-protos<2.0dev,>=1.6.0 in /usr/local/lib/python3.7/dist-packages (from google-api-core<3dev,>=1.21.0->google-api-python-client)
Requirement already satisfied: requests<3.0.0dev,>=2.18.0 in /usr/local/lib/python3.7/dist-packages (from google-api-core<3dev,>=1.21.0->google-api-python-client)
Requirement already satisfied: setuptools>=40.3.0 in /usr/local/lib/python3.7/dist-packages (from google-api-core<3dev,>=1.21.0->google-api-python-client)
Requirement already satisfied: pytz in /usr/local/lib/python3.7/dist-packages (from google-api-core<3dev,>=1.21.0->google-api-python-client)
Requirement already satisfied: packaging>=14.3 in /usr/local/lib/python3.7/dist-packages (from google-api-core<3dev,>=1.21.0->google-api-python-client)
Requirement already satisfied: cachetools<5.0,>=2.0.0 in /usr/local/lib/python3.7/dist-packages (from google-auth<3dev,>=1.16.0->google-api-python-client)
Requirement already satisfied: pyasn1-modules>=0.2.1 in /usr/local/lib/python3.7/dist-packages (from google-auth<3dev,>=1.16.0->google-api-python-client)
Requirement already satisfied: rsa<5,>=3.1.4 in /usr/local/lib/python3.7/dist-packages (from google-auth<3dev,>=1.16.0->google-api-python-client)
Requirement already satisfied: pyasn1>=0.1.7 in /usr/local/lib/python3.7/dist-packages (from oauth2client>=4.0.0->pydrive->Colab-Utils==0.0.1)
Requirement already satisfied: pyparsing!=3.0.5,>=2.0.2 in /usr/local/lib/python3.7/dist-packages (from packaging>=14.3->google-api-core<3dev,>=1.21.0->google-api-python-client)
Requirement already satisfied: chardet<4,>=3.0.2 in /usr/local/lib/python3.7/dist-packages (from requests<3.0.0dev,>=2.18.0->google-api-python-client)
Requirement already satisfied: urllib3!=1.25.0,!>=1.25.1,<1.26,>=1.21.1 in /usr/local/lib/python3.7/dist-packages (from requests<3.0.0dev,>=2.18.0->google-api-python-client)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.7/dist-packages (from requests<3.0.0dev,>=2.18.0->google-api-python-client)
Requirement already satisfied: idna<3,>=2.5 in /usr/local/lib/python3.7/dist-packages (from requests<3.0.0dev,>=2.18.0->google-api-core<3dev,>=1.21.0->google-api-python-client)
```

▼ Setup code

Run some setup code for this notebook: Import some useful packages and increase the default figure size.

```
1 from __future__ import print_function
2 from __future__ import division
3
4 import torch
5 import cutils
6 import random
7 import math
8 import matplotlib.pyplot as plt
9 from torchvision.utils import make_grid
10
11 # for plotting
12 %matplotlib inline
13 plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
14 plt.rcParams['image.interpolation'] = 'nearest'
15 plt.rcParams['image.cmap'] = 'gray'
```

We will use GPUs to accelerate our computation in this notebook. Run the following to make sure GPUs are enabled:

```
1 if torch.cuda.is_available():
2     print('Good to go!')
3 else:
4     print('Please set GPU via Edit -> Notebook Settings.')
```

Good to go!

The inputs to our network will be a batch of N (`num_inputs`) D -dimensional vectors (`input_size`); the hidden layer will have H hidden units (`hidden_size`), and we will predict classification scores for C categories (`num_classes`). This means that the learnable weights and biases of the network will have the following shapes:

- W_1 : First layer weights; has shape (D, H)
- b_1 : First layer biases; has shape $(H,)$
- W_2 : Second layer weights; has shape (H, C)
- b_2 : Second layer biases; has shape $(C,)$

We will use the following function to generate random weights for a small toy model while we implement the model:

```
1 def get_toy_data(num_inputs=5, input_size=4, hidden_size=10, num_classes=3,
2                  dtype=torch.float32):
3     N = num_inputs
4     D = input_size
5     H = hidden_size
6     C = num_classes
7
8     # We set the random seed for repeatable experiments.
9     coutils.utils.fix_random_seed()
10
11    # Generate some random parameters, storing them in a dict
12    params = {}
13    params['W1'] = 1e-4 * torch.randn(D, H, device='cuda').to(dtype)
14    params['b1'] = torch.zeros(H, device='cuda').to(dtype)
15    params['W2'] = 1e-4 * torch.randn(H, C, device='cuda').to(dtype)
16    params['b2'] = torch.zeros(C, device='cuda').to(dtype)
17
18    # Generate some random inputs and labels
19    toy_X = 10.0 * torch.randn(N, D, device='cuda').to(dtype)
20    toy_y = torch.tensor([0, 1, 2, 2, 1], dtype=torch.int64, device='cuda')
```

```
22     return toy_X, toy_y, params
```

Forward pass: compute scores

Like in the Linear Classifiers exercise, we want to write a function that takes as input the model weights and a batch of images and labels, and returns the loss and the gradient of the loss with respect to each model parameter.

However rather than attempting to implement the entire function at once, we will take a staged approach and ask you to implement the full forward and backward pass one step at a time.

First we will implement the forward pass of the network which uses the weights and biases to compute scores for all inputs:

```
1 def nn_loss_part1(params, X, y=None, reg=0.0):
2     """
3         The first stage of our neural network implementation: Run the forward pass
4         of the network to compute the hidden layer features and classification
5         scores. The network architecture should be:
6
7         FC layer -> ReLU (hidden) -> FC layer (scores)
8
9     Inputs:
10    - params: a dictionary of PyTorch Tensor that store the weights of a model.
11        It should have following keys with shape
12            W1: First layer weights; has shape (D, H)
13            b1: First layer biases; has shape (H,)
14            W2: Second layer weights; has shape (H, C)
15            b2: Second layer biases; has shape (C,)
16    - X: Input data of shape (N, D). Each X[i] is a training sample.
17    - y: Vector of training labels. y[i] is the label for X[i], and each y[i] is
18        an integer in the range 0 <= y[i] < C. This parameter is optional; if it
19        is not passed then we only return scores, and if it is passed then we
20        instead return the loss and gradients.
21    - reg: Regularization strength.
22
23    Returns a tuple of:
24    - scores: Tensor of shape (N, C) giving the classification scores for X
25    - hidden: Tensor of shape (N, H) giving the hidden layer representation
26        for each input value (after the ReLU).
27    """
28
29     # Unpack variables from the params dictionary
30     W1, b1 = params['W1'], params['b1']
31     W2, b2 = params['W2'], params['b2']
```

```

30 W2, b2 = params['W2'], params['b2']
31 N, D = X.shape
32
33 # Compute the forward pass
34 hidden = None
35 scores = None
36 #####
37 # TODO: Perform the forward pass, computing the class scores for the input. #
38 # Store the result in the scores variable, which should be an tensor of      #
39 # shape (N, C).                                #
40 #####
41 # Replace "pass" statement with your code
42 relu = torch.nn.ReLU()
43 hidden = relu(X.mm(W1) + b1)  # (N,H)
44 scores = hidden.mm(W2) + b2  # (N,C)
45
46 #####
47 #                                     END OF YOUR CODE          #
48 #####
49
50 return scores, hidden

```

Compute the scores and compare with the answer. The distance gap should be smaller than 1e-10.

```

1 toy_X, toy_y, params = get_toy_data()
2
3 scores, _ = nn_loss_part1(params, toy_X)
4 print('Your scores:')
5 print(scores)
6 print()
7 print('correct scores:')
8 correct_scores = torch.tensor([
9     [ 9.7003e-08, -1.1143e-07, -3.9961e-08],
10    [-7.4297e-08,  1.1502e-07,  1.5685e-07],
11    [-2.5860e-07,  2.2765e-07,  3.2453e-07],
12    [-4.7257e-07,  9.0935e-07,  4.0368e-07],
13    [-1.8395e-07,  7.9303e-08,  6.0360e-07]], dtype=torch.float32, device=scores.device)
14 print(correct_scores)
15 print()
16
17 # The difference should be very small. We get < 1e-10
18 scores_diff = (scores - correct_scores).abs().sum().item()
19 print('Difference between your scores and correct scores: %.2e' % scores_diff)

```

Your scores:

```
tensor([[ 9.7003e-08, -1.1143e-07, -3.9961e-08],
       [-7.4297e-08,  1.1502e-07,  1.5685e-07],
       [-2.5860e-07,  2.2765e-07,  3.2453e-07],
       [-4.7257e-07,  9.0935e-07,  4.0368e-07],
       [-1.8395e-07,  7.9303e-08,  6.0360e-07]], device='cuda:0')
```

correct scores:

```
tensor([[ 9.7003e-08, -1.1143e-07, -3.9961e-08],
       [-7.4297e-08,  1.1502e-07,  1.5685e-07],
       [-2.5860e-07,  2.2765e-07,  3.2453e-07],
       [-4.7257e-07,  9.0935e-07,  4.0368e-07],
       [-1.8395e-07,  7.9303e-08,  6.0360e-07]], device='cuda:0')
```

Difference between your scores and correct scores: 2.28e-11

Forward pass: compute loss

Now, we implement the first part of `nn_loss_part2` that computes the data and regularization loss.

For the data loss, we will use the softmax loss. For the regularization loss we will use L2 regularization on the weight matrices `w1` and `w2`; we will not apply regularization loss to the bias vectors `b1` and `b2`.

```
1 def nn_loss_part2(params, X, y=None, reg=0.0):
2     """
3         Compute the loss and gradients for a two layer fully connected neural
4         network.
5
6         Inputs: Same as nn_loss_part1
7
8         Returns:
9         If y is None, return a tensor scores of shape (N, C) where scores[i, c] is
10            the score for class c on input X[i].
11
12        If y is not None, instead return a tuple of:
13            - loss: Loss (data loss and regularization loss) for this batch of training
14                samples.
15            - grads: Dictionary mapping parameter names to gradients of those parameters
16                with respect to the loss function; has the same keys as self.params.
17        """
18
19        # Unpack variables from the params dictionary
20        W1, b1 = params['W1'], params['b1']
21        W2, b2 = params['W2'], params['b2']
22        N, D = X.shape
```

```
21 N, D = X.shape
22
23 scores, h1 = nn_loss_part1(params, X, y, reg)
24 # If the targets are not given then jump out, we're done
25 if y is None:
26     return scores
27
28 # Compute the loss
29 loss = None
30 #####
31 # TODO: Finish the forward pass, and compute the loss. This should include #
32 # both the data loss and L2 regularization for W1 and W2. Store the result #
33 # in the variable loss, which should be a scalar. Use the Softmax          #
34 # classifier loss. When you implement the regularization over W, please DO  #
35 # NOT multiply the regularization term by 1/2 (no coefficient). If you are #
36 # not careful here, it is easy to run into numeric instability (Check        #
37 # Numeric Stability in http://cs231n.github.io/linear-classify/).           #
38 #####
39 # Replace "pass" statement with your code
40
41 # W1: First layer weights; has shape (D, H)
42 # b1: First layer biases; has shape (H,)
43 # W2: Second layer weights; has shape (H, C)
44 # b2: Second layer biases; has shape (C,)
45 H, C = W2.shape
46 relu = torch.nn.ReLU()
47 #XdotW1 = X.mm(W1) + b1                      # 1) (N,D)mm(D,H) + (H,) = (N,H)
48 #hidden = relu(XdotW1)                         # 2) (N,H)
49 f = h1.mm(W2) + b2                          # 3) (N,H)mm(H,C) + (C,) = (N,C)
50 f_max = f.max(dim=1, keepdim=True).values # 3.1) Numeric instability.
51 f -= f_max                                  # 3.2)
52 f_exp = f.exp()                            # 4) (N,C)
53 f_exp_s = f_exp.sum(dim=1, keepdim=True) # 5) (N,1)
54 f_exp_s_l = f_exp_s.log()                  # 6) (N,1)
55 f_exp_s_l_s = f_exp_s_l.sum()            # 7) (1,)
56 f_exp_s_l_s_div = f_exp_s_l_s / float(N) # 8) (1,)
57 y_m = torch.zeros((N,C)).to(W1)           # (N,C). y matrix
58 y_m[range(N), y] = 1
59 f_y_m = (f*y_m)                           # 9) (N,C)
60 f_y_m_s = f_y_m.sum()                     # 9.1) (1,)
61 f_y_m_s_div = f_y_m_s / -float(N)        # 10) (1,)
62 regW1pow2sum = W1.pow(2).sum() * 2 * reg  # 11) ()
63 regW2pow2sum = W2.pow(2).sum() * 2 * reg  # 12)
64 rprod = regW1pow2sum + regW2pow2sum # 13)
65 loss = f_y_m_s_div + f_exp_s_l_s_div + rprod # 14)
66
67 dloss = 1.0
68 # 1.1)
```

```

68 # 14)
69 df_y_m_s_div = dloss
70 df_exp_s_l_s_div = dloss
71 drprod = dloss
72 # 13)
73 d2regW1pow2sum = drprod
74 d2regW2pow2sum = drprod
75 # 12)
76 dW2 = d2regW2pow2sum * 2 * reg * torch.ones_like(W2) * 2 * W2 # (H,C)
77 # 11)
78 dW1 = d2regW1pow2sum * 2 * reg * torch.ones_like(W1) * 2 * W1 # (D,H)
79 # 10)
80 df_y_m_s = df_exp_s_l_s_div * (1.0 / -float(N)) # (1,)
81 # 9.1)
82 df_y_m = torch.ones_like(f_y_m) * df_y_m_s # (N,C)
83 # 9
84 df = y_m * df_y_m # (N,C) * (N,C)
85 # 8
86 df_exp_s_l_s = df_exp_s_l_s_div * (1.0 / float(N)) # (1,)
87 # 7
88 df_exp_s_l = torch.ones_like(f_exp_s_l) * df_exp_s_l_s # (N, 1)
89 # 6
90 df_exp_s = (1.0 / f_exp_s ) * df_exp_s_l # (N,1) * (N,1)
91 # 5
92 df_exp = df_exp_s.mm(torch.ones((1,C)).to(f_exp)) # (N,C)
93 #df_exp = torch.ones_like(f_exp).sum(dim=1, keepdim=True) * df_exp_s # (N,1)
94 # 4
95 df = df + f_exp * df_exp # (N,C) + (N,C) * (N,1) = (N,C)
96 # 3.2
97 # df =
98 # 3
99 dW2 = dW2 + h1.T.mm(df) # (H,N) mm (N,C) -> (H,C)
100 db2 = df.T.sum(dim=1) # (C,N) mm (N,1) = (C, 1)
101 # 3
102 dhidden = df.mm(W2.T) # (N,C) x (C,H)-> (N,H)
103 # 2
104 drelu = torch.where(h1 <= 0, 0, dhidden) # (N,H)
105 # 1
106 dW1 = dW1 + X.T.mm(drelu) # (D,N) x (N,H) -> (D,H)
107 db1 = drelu.T.sum(dim=1) # -> (H, )
108 #####
109 # END OF YOUR CODE #
110 #####
111 #####
112 # Backward pass: compute gradients
113 grads = {}
114

```

```

115 #####
116 # TODO: Compute the backward pass, computing the derivatives of the weights #
117 # and biases. Store the results in the grads dictionary. For example,      #
118 # grads['W1'] should store the gradient on W1, and be a tensor of same size #
119 #####
120 # Replace "pass" statement with your code
121 grads['W1'] = dW1
122 grads['W2'] = dW2
123 grads['b1'] = db1
124 grads['b2'] = db2
125
126 #####
127 #           END OF YOUR CODE           #
128 #####
129
130 return loss, grads

```

First, implement the forward pass in the function `nn_loss_part2` above. Then run the following to check your implementation.

We compute the loss for the toy data, and compare with the answer computed by our implementation. The difference between the correct and computed loss should be less than `1e-4`.

```

1 toy_X, toy_y, params = get_toy_data()
2
3 loss, _ = nn_loss_part2(params, toy_X, toy_y, reg=0.05)
4 print('Your loss: ', loss.item())
5 correct_loss = 1.0986
6 print('Correct loss: ', correct_loss)
7 diff = (correct_loss - loss).item()
8
9 # should be very small, we get < 1e-4
10 print('Difference: %.4e' % diff)

```

```

Your loss: 1.0986123085021973
Correct loss: 1.0986
Difference: -1.2279e-05

```

Backward pass

Now implement the backward pass for the entire network in `nn_loss_part2`.

After doing so, we will use numeric gradient checking to see whether the analytic gradient computed by our backward pass matches a

numeric gradient.

First we define a couple utility functions for our numeric gradient check:

```
1 def compute_numeric_gradient(f, x, h=1e-7):
2     """
3         Compute the numeric gradient of f at x using a finite differences
4         approximation. We use the centered difference:
5
6         df/dx ~= (f(x + h) - f(x - h)) / (2 * h)
7
8     Inputs:
9         - f: A function that inputs a torch tensor and returns a torch scalar
10        - x: A torch tensor giving the point at which to compute the gradient
11
12    Returns:
13        - grad: A tensor of the same shape as x giving the gradient of f at x
14    """
15    fx = f(x) # evaluate function value at original point
16    flat_x = x.contiguous().view(-1)
17    grad = torch.zeros_like(x)
18    flat_grad = grad.view(-1)
19    # iterate over all indexes in x
20    for i in range(flat_x.shape[0]):
21        oldval = flat_x[i].item() # Store the original value
22        flat_x[i] = oldval + h    # Increment by h
23        fxph = f(x).item()       # Evaluate f(x + h)
24        flat_x[i] = oldval - h    # Decrement by h
25        fxmh = f(x).item()       # Evaluate f(x - h)
26        flat_x[i] = oldval       # Restore original value
27
28        # compute the partial derivative with centered formula
29        flat_grad[i] = (fxph - fxmh) / (2 * h)
30
31    return grad
32
33
34 def rel_error(x, y, eps=1e-10):
35     """ returns relative error between x and y """
36     top = (x - y).abs().max().item()
37     bot = (x.abs() + y.abs()).clamp(min=eps).max().item()
38     return top / bot
```

Now we will compute the gradient of the loss with respect to the variables `w1`, `b1`, `w2`, and `b2`. Now that you (hopefully!) have a correctly

implemented forward pass, you can debug your backward pass using a numeric gradient check.

You should see relative errors less than `1e-4` for all parameters.

```
1 reg = 0.05
2 toy_X, toy_y, params = get_toy_data(dtype=torch.float64)
3 loss, grads = nn_loss_part2(params, toy_X, toy_y, reg=reg)
4
5 for param_name, grad in grads.items():
6     param = params[param_name]
7     f = lambda w: nn_loss_part2(params, toy_X, toy_y, reg=reg)[0]
8     grad_numeric = compute_numeric_gradient(f, param)
9     error = rel_error(grad, grad_numeric)
10    print('%s max relative error: %e' % (param_name, error))
```

```
W1 max relative error: 1.402414e-06
W2 max relative error: 1.154466e-06
b1 max relative error: 2.666979e-05
b2 max relative error: 3.845737e-09
```

Train the network

To train the network we will use stochastic gradient descent (SGD), similar to the SVM and Softmax classifiers. Look at the function `nn_train` and fill in the missing sections to implement the training procedure. This should be very similar to the training procedure you used for the SVM and Softmax classifiers.

```
1 def nn_train(params, loss_func, pred_func, X, y, X_val, y_val,
2             learning_rate=1e-3, learning_rate_decay=0.95,
3             reg=5e-6, num_iters=100,
4             batch_size=200, verbose=False):
5     """
6     Train this neural network using stochastic gradient descent.
7
8     Inputs:
9     - params: a dictionary of PyTorch Tensor that store the weights of a model.
10    It should have following keys with shape
11        W1: First layer weights; has shape (D, H)
12        b1: First layer biases; has shape (H,)
13        W2: Second layer weights; has shape (H, C)
14        b2: Second layer biases; has shape (C,)
15    - loss_func: a loss function that computes the loss and the gradients.
```

```
16 It takes as input:  
17 - params: Same as input to nn_train  
18 - X_batch: A minibatch of inputs of shape (B, D)  
19 - y_batch: Ground-truth labels for X_batch  
20 - reg: Same as input to nn_train  
21 And it returns a tuple of:  
22     - loss: Scalar giving the loss on the minibatch  
23     - grads: Dictionary mapping parameter names to gradients of the loss with  
24         respect to the corresponding parameter.  
25 - pred_func: prediction function that im  
26 - X: A PyTorch tensor of shape (N, D) giving training data.  
27 - y: A PyTorch tensor f shape (N,) giving training labels; y[i] = c means that  
28     X[i] has label c, where 0 <= c < C.  
29 - X_val: A PyTorch tensor of shape (N_val, D) giving validation data.  
30 - y_val: A PyTorch tensor of shape (N_val,) giving validation labels.  
31 - learning_rate: Scalar giving learning rate for optimization.  
32 - learning_rate_decay: Scalar giving factor used to decay the learning rate  
33     after each epoch.  
34 - reg: Scalar giving regularization strength.  
35 - num_iters: Number of steps to take when optimizing.  
36 - batch_size: Number of training examples to use per step.  
37 - verbose: boolean; if true print progress during optimization.  
38  
39 Returns: A dictionary giving statistics about the training process  
40 """  
41 num_train = X.shape[0]  
42 iterations_per_epoch = max(num_train // batch_size, 1)  
43  
44 # Use SGD to optimize the parameters in self.model  
45 loss_history = []  
46 train_acc_history = []  
47 val_acc_history = []  
48  
49 for it in range(num_iters):  
50     X_batch = None  
51     y_batch = None  
52  
53 #####  
54 # TODO: Create a random minibatch of training data and labels, storing #  
55 # them in X_batch and y_batch respectively. #  
56 # hint: torch.randint #  
57 #####  
58 # Replace "pass" statement with your code  
59 batch_idxes = torch.randint(num_train, (batch_size,))  
60 X_batch = X[batch_idxes, :]  
61 y_batch = y[batch_idxes]  
62 #####
```

```

63 #                                     END OF YOUR CODE #
64 #####
65
66 # Compute loss and gradients using the current minibatch
67 loss, grads = loss_func(params, X_batch, y=y_batch, reg=reg)
68 loss_history.append(loss.item())
69
70 #####
71 # TODO: Use the gradients in the grads dictionary to update the      #
72 # parameters of the network (stored in the dictionary self.params)      #
73 # using stochastic gradient descent. You'll need to use the gradients   #
74 # stored in the grads dictionary defined above.                         #
75 #####
76 # Replace "pass" statement with your code
77 params['W1'] -= learning_rate * grads['W1']
78 params['W2'] -= learning_rate * grads['W2']
79 params['b1'] -= learning_rate * grads['b1']
80 params['b2'] -= learning_rate * grads['b2']
81 #####
82 #                                     END OF YOUR CODE #
83 #####
84
85 if verbose and it % 100 == 0:
86     print('iteration %d / %d: loss %f' % (it, num_iters, loss.item()))
87
88 # Every epoch, check train and val accuracy and decay learning rate.
89 if it % iterations_per_epoch == 0:
90     # Check accuracy
91     y_train_pred = pred_func(params, loss_func, X_batch)
92     train_acc = (y_train_pred == y_batch).float().mean().item()
93     y_val_pred = pred_func(params, loss_func, X_val)
94     val_acc = (y_val_pred == y_val).float().mean().item()
95     train_acc_history.append(train_acc)
96     val_acc_history.append(val_acc)
97
98     # Decay learning rate
99     learning_rate *= learning_rate_decay
100
101 return {
102     'loss_history': loss_history,
103     'train_acc_history': train_acc_history,
104     'val_acc_history': val_acc_history,
105 }

```

You will also have to implement `nn_predict`, as the training process periodically performs prediction to keep track of accuracy over time while

the network trains.

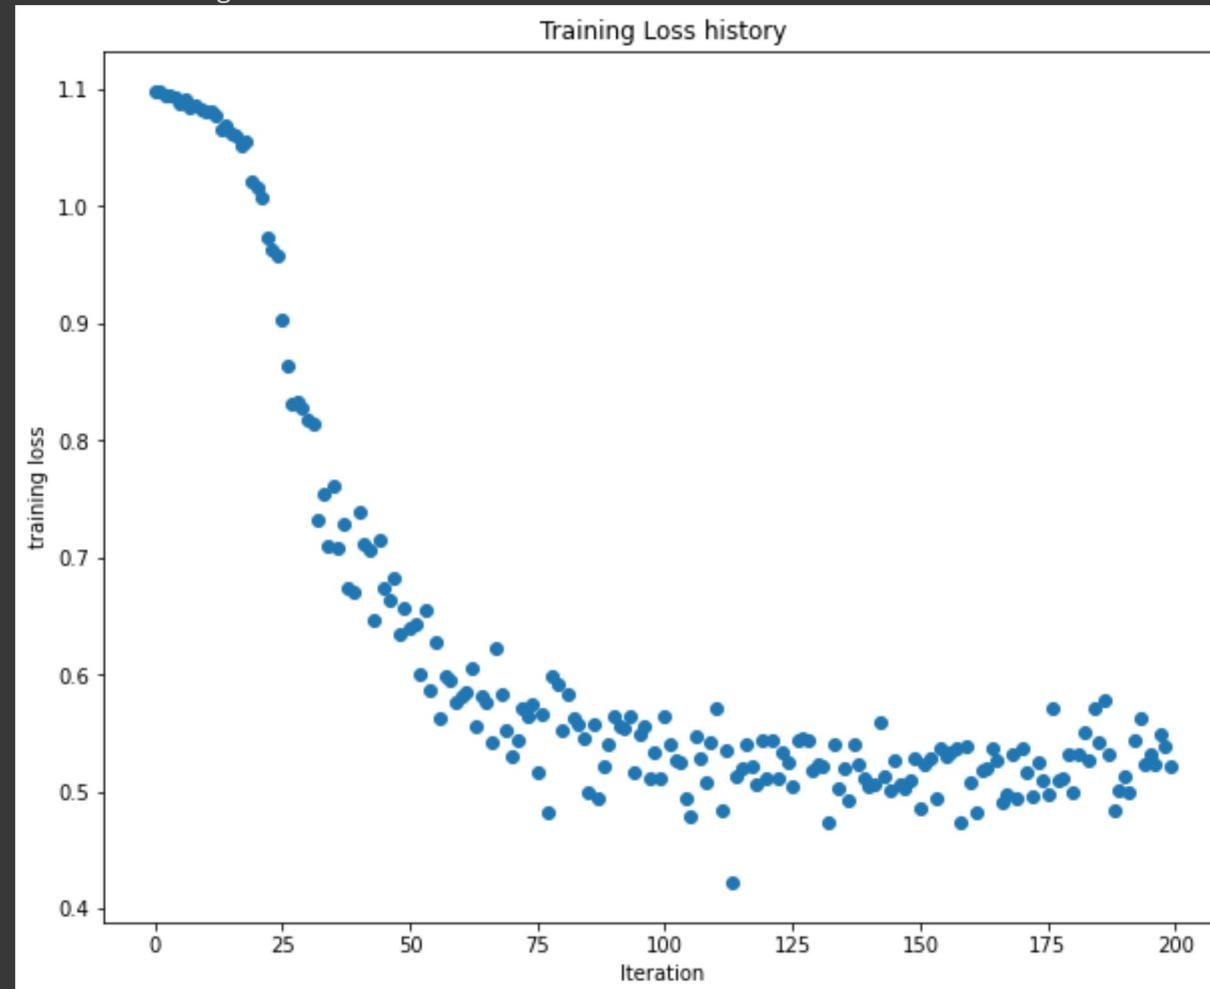
```
1 def nn_predict(params, loss_func, X):
2     """
3     Use the trained weights of this two-layer network to predict labels for
4     data points. For each data point we predict scores for each of the C
5     classes, and assign each data point to the class with the highest score.
6
7     Inputs:
8     - params: a dictionary of PyTorch Tensor that store the weights of a model.
9     It should have following keys with shape
10    W1: First layer weights; has shape (D, H)
11    b1: First layer biases; has shape (H,)
12    W2: Second layer weights; has shape (H, C)
13    b2: Second layer biases; has shape (C,)
14    - loss_func: a loss function that computes the loss and the gradients
15    - X: A PyTorch tensor of shape (N, D) giving N D-dimensional data points to
16      classify.
17
18    Returns:
19    - y_pred: A PyTorch tensor of shape (N,) giving predicted labels for each of
20      the elements of X. For all i, y_pred[i] = c means that X[i] is predicted
21      to have class c, where 0 <= c < C.
22    """
23    y_pred = None
24
25    #####
26    # TODO: Implement this function; it should be VERY simple! #
27    #####
28    # Replace "pass" statement with your code
29    scores = loss_func(params, X, y=None)
30    y_pred = torch.argmax(scores, dim=1)
31    #####
32    #           END OF YOUR CODE #
33    #####
34
35    return y_pred
```

Once you have implemented the method, run the code below to train a two-layer network on toy data. Your final training loss should be less than 1.05.

```
1 toy_X, toy_y, params = get_toy_data()
2
```

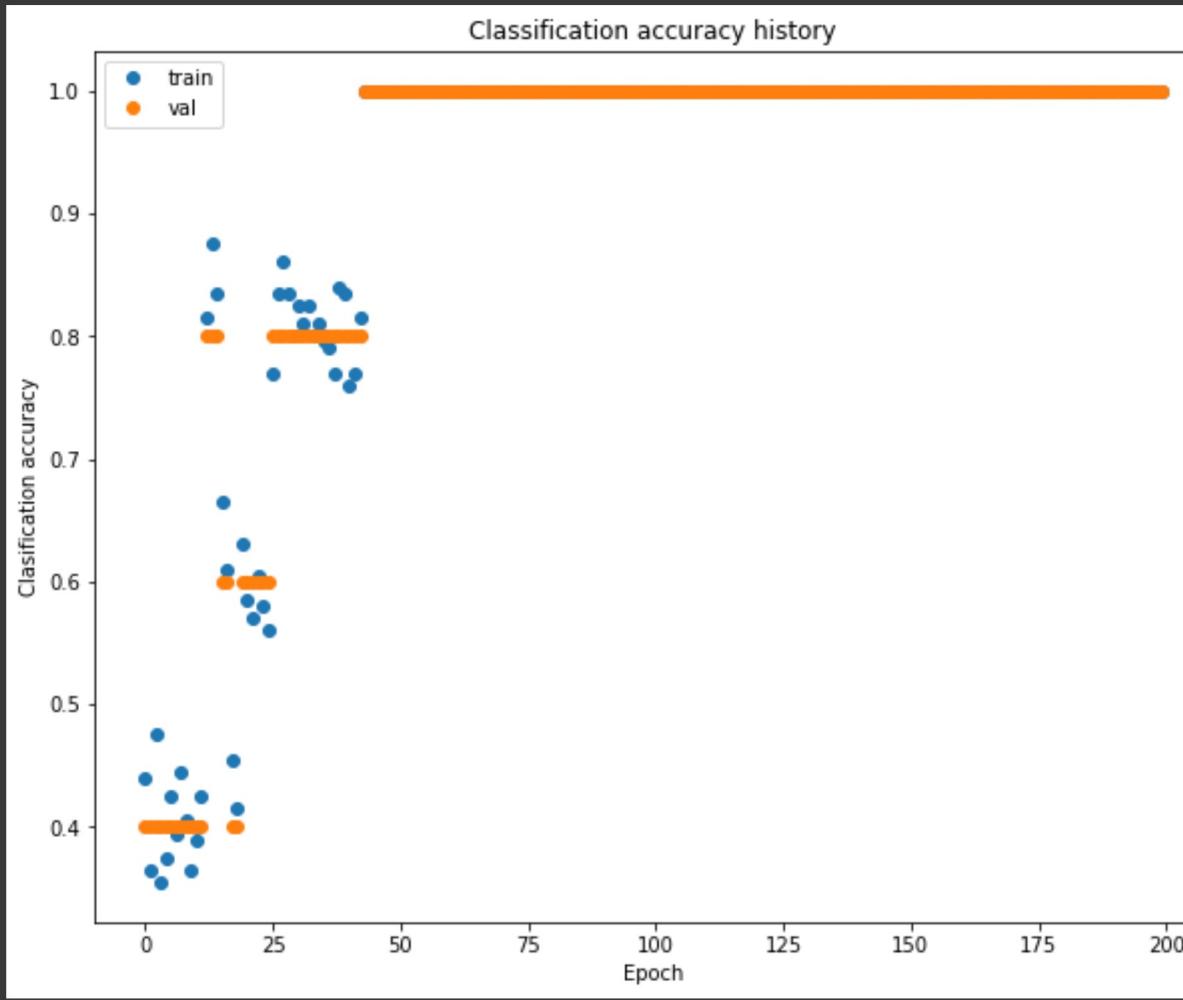
```
3 stats = nn_train(params, nn_loss_part2, nn_predict, toy_X, toy_y, toy_X, toy_y,
4                   learning_rate=1e-1, reg=1e-6,
5                   num_iters=200, verbose=False)
6
7 print('Final training loss: ', stats['loss_history'][-1])
8
9 # plot the loss history
10 plt.plot(stats['loss_history'], 'o')
11 plt.xlabel('Iteration')
12 plt.ylabel('training loss')
13 plt.title('Training Loss history')
14 plt.show()
```

Final training loss: 0.5211766362190247



```
1 # Plot the loss function and train / validation accuracies  
2 plt.figure(figsize=(10, 6))
```

```
2 plt.plot(stats['train_acc_history'], 'o', label='train')
3 plt.plot(stats['val_acc_history'], 'o', label='val')
4 plt.title('Classification accuracy history')
5 plt.xlabel('Epoch')
6 plt.ylabel('Classification accuracy')
7 plt.legend()
8 plt.show()
```



Wrap all function into a Class

We will use the class `TwoLayerNet` to represent instances of our network. The network parameters are stored in the instance variable `self.params` where keys are string parameter names and values are PyTorch tensors.

```
1 class TwoLayerNet(object):
2     def __init__(self, input_size, hidden_size, output_size, device='cuda',
3                  std=1e-4):
4         """
5             Initialize the model. Weights are initialized to small random values and
6             biases are initialized to zero. Weights and biases are stored in the
7             variable self.params, which is a dictionary with the following keys:
8
9             W1: First layer weights; has shape (D, H)
10            b1: First layer biases; has shape (H, )
11            W2: Second layer weights; has shape (H, C)
12            b2: Second layer biases; has shape (C, )
13
14        Inputs:
15        - input_size: The dimension D of the input data.
16        - hidden_size: The number of neurons H in the hidden layer.
17        - output_size: The number of classes C.
18        """
19
20        # fix random seed before we generate a set of parameters
21        cutils.utils.fix_random_seed()
22
23        self.params = {}
24        self.params['W1'] = std * torch.randn(input_size, hidden_size, device=device)
25        self.params['b1'] = torch.zeros(hidden_size, device=device)
26        self.params['W2'] = std * torch.randn(hidden_size, output_size, device=device)
27        self.params['b2'] = torch.zeros(output_size, device=device)
28
29    def _loss(self, params, X, y=None, reg=0.0):
30        return nn_loss_part2(params, X, y, reg)
31
32    def loss(self, X, y=None, reg=0.0):
33        return self._loss(self.params, X, y, reg)
34
35    def _train(self, params, loss_func, pred_func, X, y, X_val, y_val,
36               learning_rate=1e-3, learning_rate_decay=0.95,
37               reg=5e-6, num_iters=100,
38               batch_size=200, verbose=False):
39        return nn_train(params, loss_func, pred_func, X, y, X_val, y_val,
40                       learning_rate, learning_rate_decay,
41                       reg, num_iters, batch_size, verbose)
42
43    def train(self, X, y, X_val, y_val,
44              learning_rate=1e-3, learning_rate_decay=0.95,
45              reg=5e-6, num_iters=100,
46              batch_size=200, verbose=False):
47        return self._train(self.params, self._loss, self._predict,
```

```

48             X, y, X_val, y_val,
49             learning_rate, learning_rate_decay,
50             reg, num_iters, batch_size, verbose)
51
52     def _predict(self, params, loss_func, X):
53         return nn_predict(params, loss_func, X)
54
55     def predict(self, X):
56         return self._predict(self.params, self._loss, X)

```

Load CIFAR-10 data

Now that you have implemented a two-layer network that passes gradient checks and works on toy data, it's time to load up our favorite CIFAR-10 data so we can use it to train a classifier on a real dataset.

```

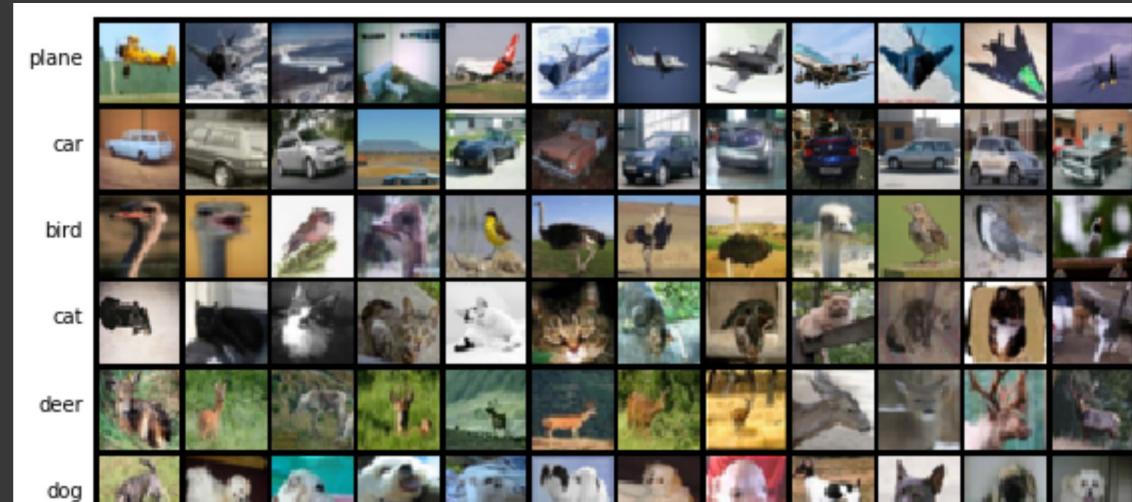
1 def get_CIFAR10_data(validation_ratio = 0.05):
2     """
3     Load the CIFAR-10 dataset from disk and perform preprocessing to prepare
4     it for the linear classifier. These are the same steps as we used for the
5     SVM, but condensed to a single function.
6     """
7     X_train, y_train, X_test, y_test = coutils.data.cifar10()
8
9     # load every data on cuda
10    X_train = X_train.cuda()
11    y_train = y_train.cuda()
12    X_test = X_test.cuda()
13    y_test = y_test.cuda()
14
15    # 0. Visualize some examples from the dataset.
16    class_names = [
17        'plane', 'car', 'bird', 'cat', 'deer',
18        'dog', 'frog', 'horse', 'ship', 'truck'
19    ]
20    img = coutils.utils.visualize_dataset(X_train, y_train, 12, class_names)
21    plt.imshow(img)
22    plt.axis('off')
23    plt.show()
24
25    # 1. Normalize the data: subtract the mean RGB (zero mean)
26    mean_image = X_train.mean(dim=0, keepdim=True).mean(dim=2, keepdim=True).mean(dim=3, keepdim=True)
27    X_train -= mean_image
28    X_test -= mean_image

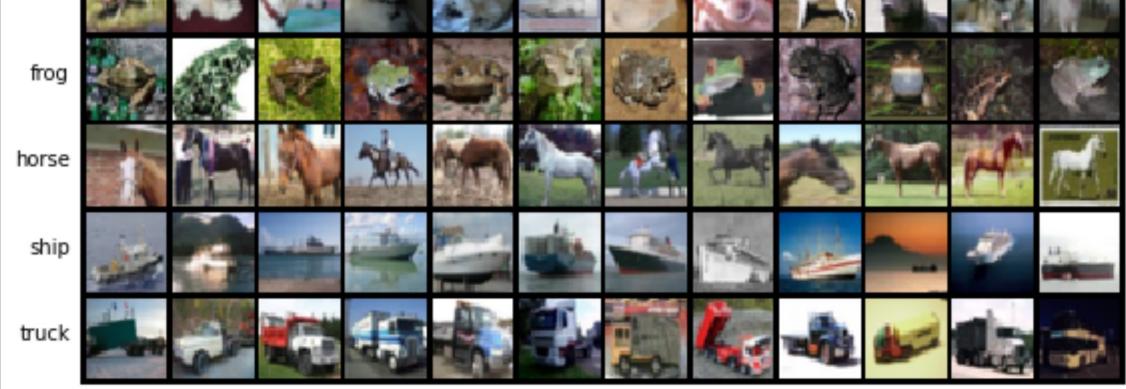
```

```

29 # 2. Reshape the image data into rows
30 X_train = X_train.reshape(X_train.shape[0], -1)
31 X_test = X_test.reshape(X_test.shape[0], -1)
32
33
34 # 3. take the validation set from the training set
35 # Note: It should not be taken from the test set
36 # For random permumation, you can use torch.randperm or torch.randint
37 # But, for this homework, we use slicing instead.
38 num_training = int( X_train.shape[0] * (1.0 - validation_ratio) )
39 num_validation = X_train.shape[0] - num_training
40
41 # return the dataset
42 data_dict = {}
43 data_dict['X_val'] = X_train[num_training:num_training + num_validation]
44 data_dict['y_val'] = y_train[num_training:num_training + num_validation]
45 data_dict['X_train'] = X_train[0:num_training]
46 data_dict['y_train'] = y_train[0:num_training]
47
48 data_dict['X_test'] = X_test
49 data_dict['y_test'] = y_test
50 return data_dict
51
52 # Invoke the above function to get our data.
53 data_dict = get_CIFAR10_data()
54 print('Train data shape: ', data_dict['X_train'].shape)
55 print('Train labels shape: ', data_dict['y_train'].shape)
56 print('Validation data shape: ', data_dict['X_val'].shape)
57 print('Validation labels shape: ', data_dict['y_val'].shape)
58 print('Test data shape: ', data_dict['X_test'].shape)
59 print('Test labels shape: ', data_dict['y_test'].shape)

```





```
Train data shape: torch.Size([47500, 3072])
Train labels shape: torch.Size([47500])
Validation data shape: torch.Size([2500, 3072])
Validation labels shape: torch.Size([2500])
Test data shape: torch.Size([10000, 3072])
Test labels shape: torch.Size([10000])
```

Train a network

To train our network we will use SGD. In addition, we will adjust the learning rate with an exponential learning rate schedule as optimization proceeds; after each epoch, we will reduce the learning rate by multiplying it by a decay rate.

```
1 input_size = 3 * 32 * 32
2 hidden_size = 36
3 num_classes = 10
4 net = TwoLayerNet(input_size, hidden_size, num_classes)
5
6 # Train the network
7 stats = net.train(data_dict['X_train'], data_dict['y_train'],
8                   data_dict['X_val'], data_dict['y_val'],
9                   num_iters=500, batch_size=1000,
10                  learning_rate=1e-2, learning_rate_decay=0.95,
11                  reg=0.25, verbose=True)
12
13 # Predict on the validation set
14 y_val_pred = net.predict(data_dict['X_val'])
15 val_acc = 100.0 * (y_val_pred == data_dict['y_val']).float().mean().item()
16 print('Validation accuracy: %.2f%%' % val_acc)
```

```
iteration 0 / 500: loss 2.303139
iteration 100 / 500: loss 2.302684
```

```
iteration 200 / 500: loss 2.302641
iteration 300 / 500: loss 2.302566
iteration 400 / 500: loss 2.302623
Validation accuracy: 8.76%
```

Debug the training

With the default parameters we provided above, you should get a validation accuracy of about 8.76% on the validation set. This isn't very good.

One strategy for getting insight into what's wrong is to plot the loss function and the accuracies on the training and validation sets during optimization.

Another strategy is to visualize the weights that were learned in the first layer of the network. In most neural networks trained on visual data, the first layer weights typically show some visible structure when visualized.

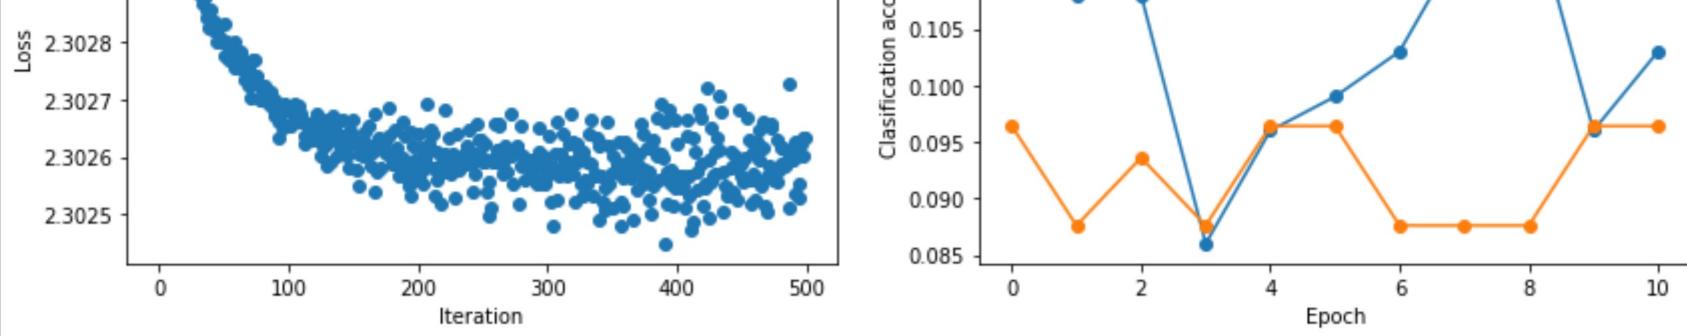
```
1 # Plot the loss function and train / validation accuracies
2 def plot_stats(stat_dict):
3     plt.subplot(1, 2, 1)
4     plt.plot(stat_dict['loss_history'], 'o')
5     plt.title('Loss history')
6     plt.xlabel('Iteration')
7     plt.ylabel('Loss')
8
9     plt.subplot(1, 2, 2)
10    plt.plot(stat_dict['train_acc_history'], 'o-', label='train')
11    plt.plot(stat_dict['val_acc_history'], 'o-', label='val')
12    plt.title('Classification accuracy history')
13    plt.xlabel('Epoch')
14    plt.ylabel('Classification accuracy')
15    plt.legend()
16
17    plt.gcf().set_size_inches(14, 4)
18    plt.show()
19
20 plot_stats(stats)
```

Loss history



Classification accuracy history





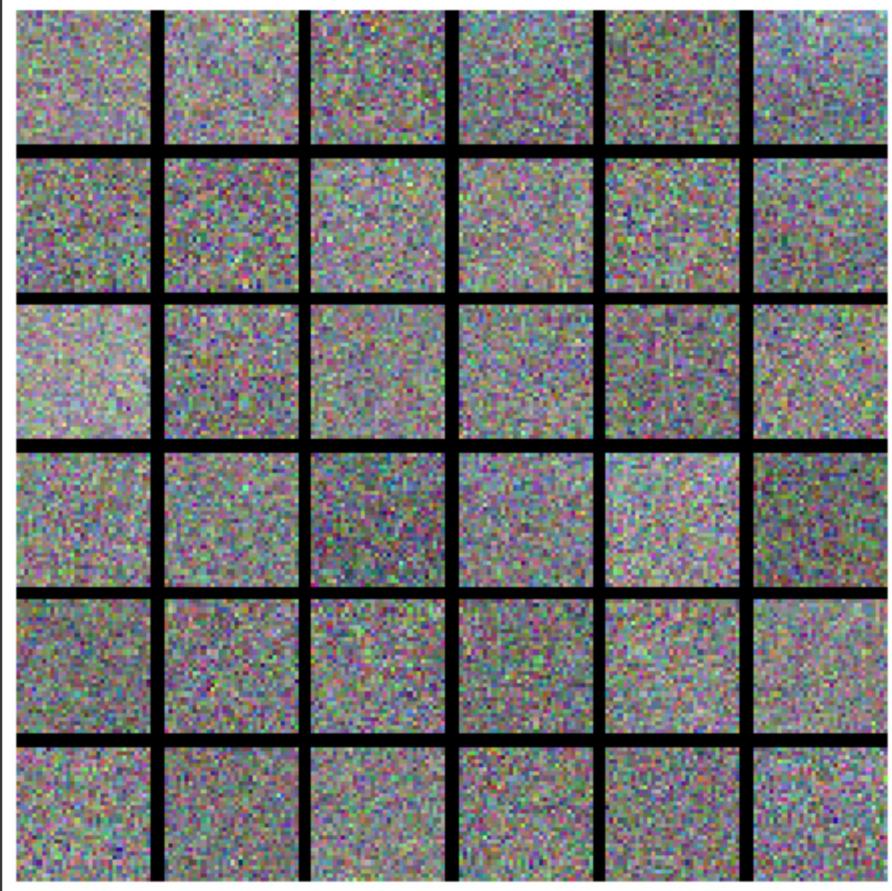
Similar to SVM and Softmax classifier, let's visualize the weights.

```

1 def visualize_grid(Xs, ubound=255.0, padding=1):
2     """
3         Reshape a 4D tensor of image data to a grid for easy visualization.
4
5     Inputs:
6     - Xs: Data of shape (N, H, W, C)
7     - ubound: Output grid will have values scaled to the range [0, ubound]
8     - padding: The number of blank pixels between elements of the grid
9     """
10    (N, H, W, C) = Xs.shape
11    # print(Xs.shape)
12    grid_size = int(math.ceil(math.sqrt(N)))
13    grid_height = H * grid_size + padding * (grid_size - 1)
14    grid_width = W * grid_size + padding * (grid_size - 1)
15    grid = torch.zeros((grid_height, grid_width, C), device=Xs.device)
16    next_idx = 0
17    y0, y1 = 0, H
18    for y in range(grid_size):
19        x0, x1 = 0, W
20        for x in range(grid_size):
21            if next_idx < N:
22                img = Xs[next_idx]
23                low, high = torch.min(img), torch.max(img)
24                grid[y0:y1, x0:x1] = ubound * (img - low) / (high - low)
25                # grid[y0:y1, x0:x1] = Xs[next_idx]
26                next_idx += 1
27            x0 += W + padding
28            x1 += W + padding
29            y0 += H + padding
30            y1 += H + padding
31    # print(grid.shape)
32    return grid

```

```
32     return grid
33
34
35 # Visualize the weights of the network
36 def show_net_weights(net):
37     W1 = net.params['W1']
38     W1 = W1.reshape(3, 32, 32, -1).transpose(0, 3)
39     plt.imshow(visualize_grid(W1, padding=3).type(torch.uint8).cpu())
40     plt.gca().axis('off')
41     plt.show()
42
43 show_net_weights(net)
```



What's wrong?

Looking at the visualizations above, we see that the loss is decreasing more or less linearly, which seems to suggest that the learning rate may be too low. Moreover, there is no gap between the training and validation accuracy, suggesting that the model we used has low capacity, and

that we should increase its size. On the other hand, with a very large model we would expect to see more overfitting, which would manifest itself as a very large gap between the training and validation accuracy.

Capacity?

Our initial model has very similar performance on the training and validation sets. This suggests that the model is underfitting, and that its performance might improve if we were to increase its capacity.

One way we can increase the capacity of a neural network model is to increase the size of its hidden layer. Here we investigate the effect of increasing the size of the hidden layer. The performance (as measured by validation-set accuracy) should increase as the size of the hidden layer increases; however it may show diminishing returns for larger layer sizes.

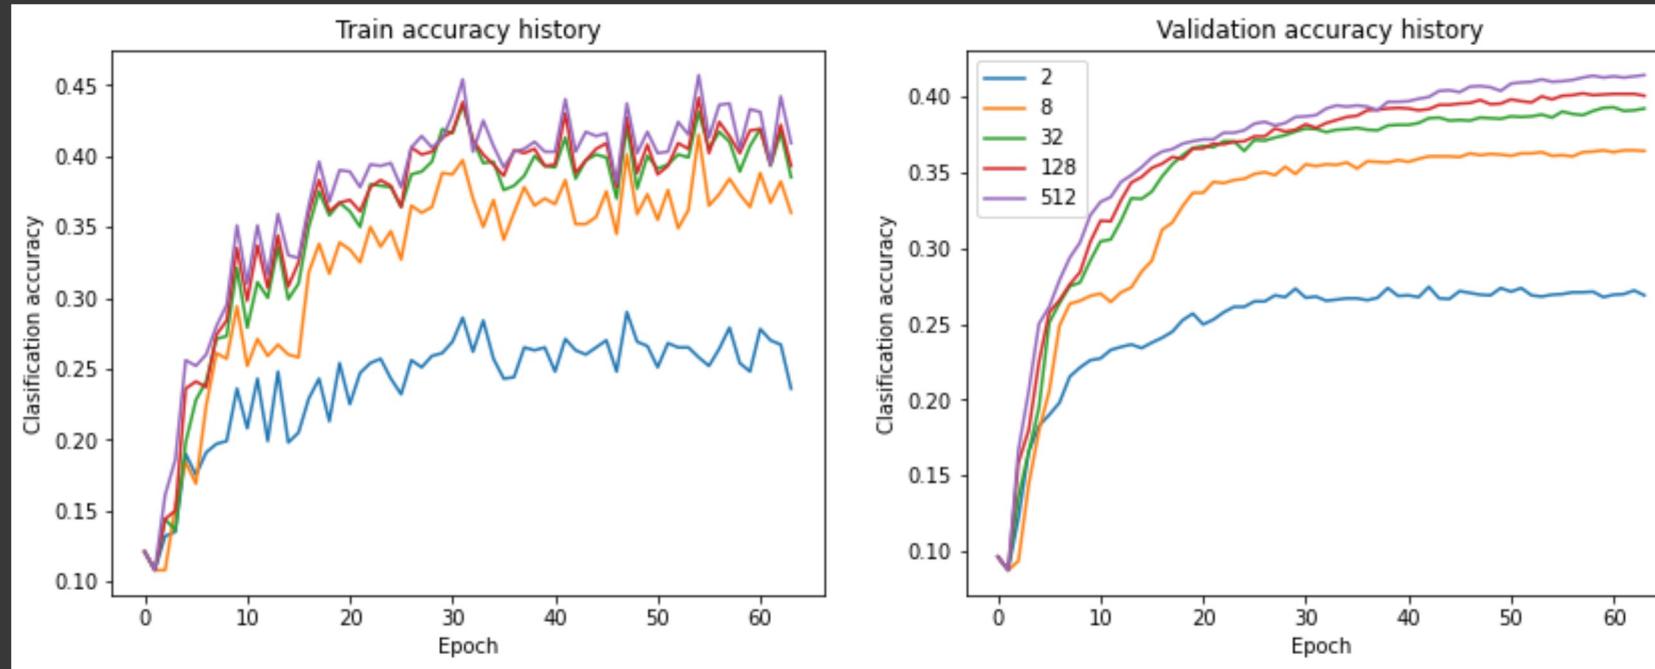
```
1 def plot_acc_curves(stat_dict):
2     plt.subplot(1, 2, 1)
3     for key, single_stats in stat_dict.items():
4         plt.plot(single_stats['train_acc_history'], label=str(key))
5     plt.title('Train accuracy history')
6     plt.xlabel('Epoch')
7     plt.ylabel('Classification accuracy')
8
9     plt.subplot(1, 2, 2)
10    for key, single_stats in stat_dict.items():
11        plt.plot(single_stats['val_acc_history'], label=str(key))
12    plt.title('Validation accuracy history')
13    plt.xlabel('Epoch')
14    plt.ylabel('Classification accuracy')
15    plt.legend()
16
17    plt.gcf().set_size_inches(14, 5)
18    plt.show()
```

```

11     learning_rate=lr, learning_rate_decay=0.95,
12     reg=reg, verbose=False)
13 stat_dict[hs] = stats
14
15 plot_acc_curves(stat_dict)

```

train with hidden size: 2
 train with hidden size: 8
 train with hidden size: 32
 train with hidden size: 128
 train with hidden size: 512



Regularization?

Another possible explanation for the small gap we saw between the train and validation accuracies of our model is regularization. In particular, if the regularization coefficient were too high then the model may be unable to fit the training data.

We can investigate the phenomenon empirically by training a set of models with varying regularization strengths while fixing other hyperparameters.

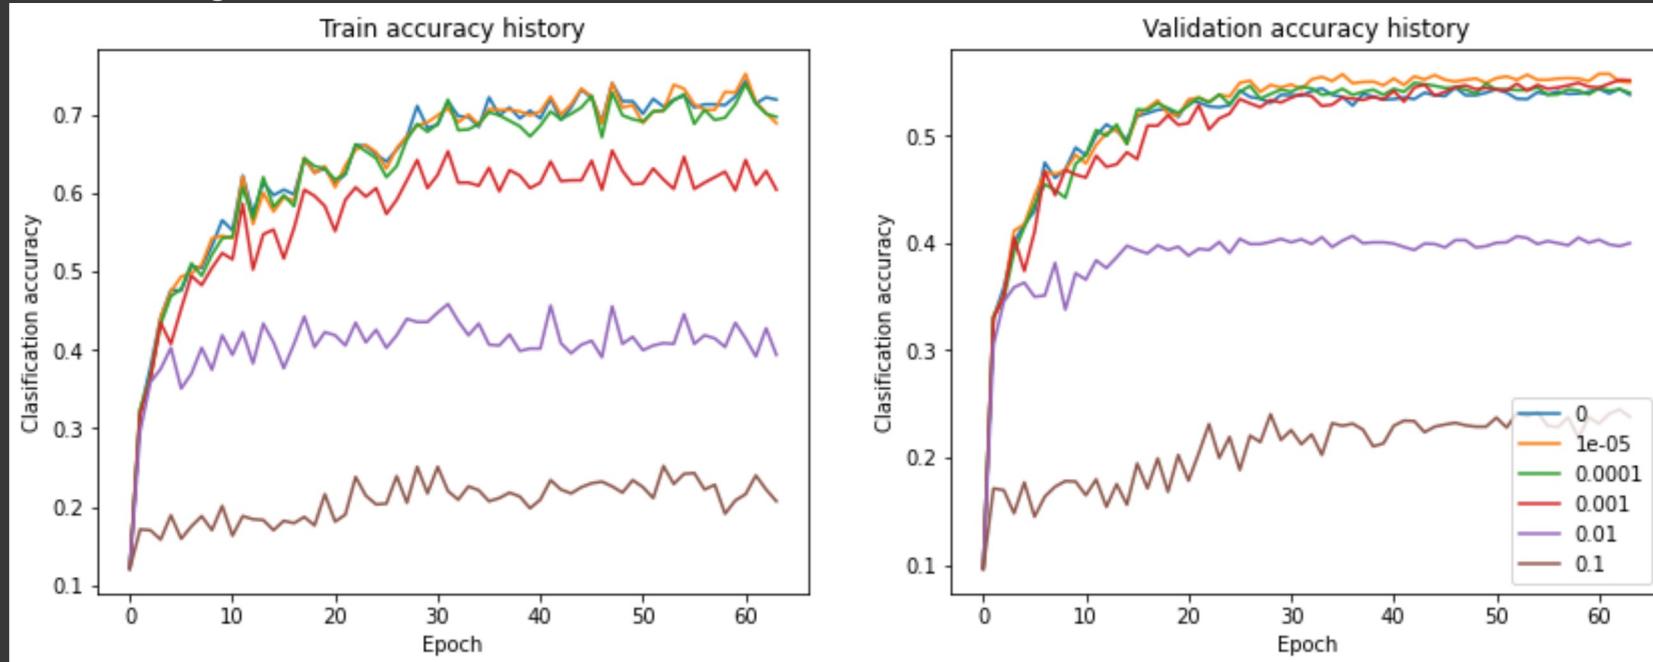
You should see that setting the regularization strength too high will harm the validation-set performance of the model:

```

2 lr = 1.0
3 regs = [0, 1e-5, 1e-4, 1e-3, 1e-2, 1e-1]
4
5 stat_dict = {}
6 for reg in regs:
7     print('train with regularization: {}'.format(reg))
8     net = TwoLayerNet(3 * 32 * 32, hs, 10, device=data_dict['X_train'].device)
9     stats = net.train(data_dict['X_train'], data_dict['y_train'], data_dict['X_val'], data_dict['y_val'],
10         num_iters=3000, batch_size=1000,
11         learning_rate=lr, learning_rate_decay=0.95,
12         reg=reg, verbose=False)
13     stat_dict[reg] = stats
14
15 plot_acc_curves(stat_dict)

```

train with regularization: 0
 train with regularization: 1e-05
 train with regularization: 0.0001
 train with regularization: 0.001
 train with regularization: 0.01
 train with regularization: 0.1



Learning Rate?

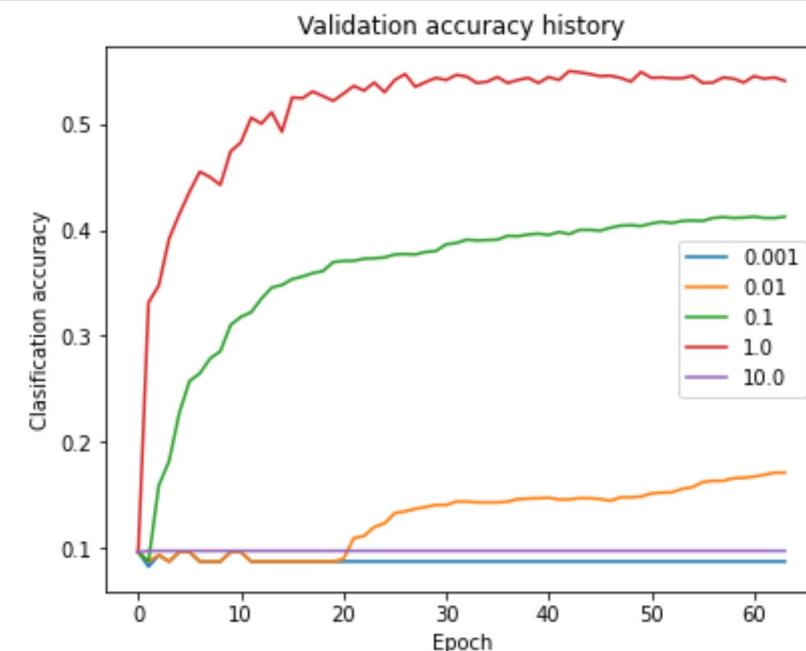
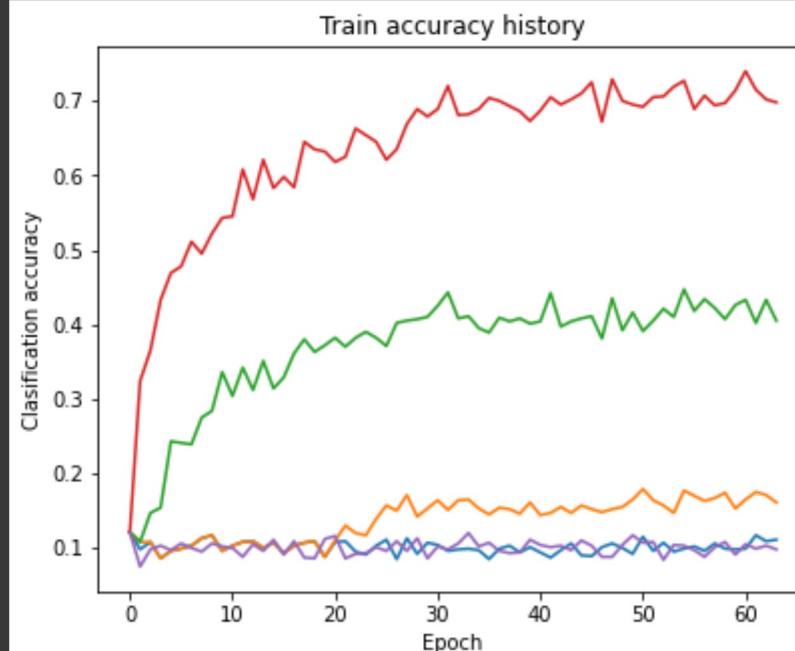
Last but not least, we also want to see the effect of learning rate with respect to the performance.

```

1 hs = 128
2 lrs = [1e-3, 1e-2, 1e-1, 1e0, 1e1]
3 reg = 1e-4
4
5 stat_dict = {}
6 for lr in lrs:
7     print('train with learning rate: {}'.format(lr))
8     net = TwoLayerNet(3 * 32 * 32, hs, 10, device=data_dict['X_train'].device)
9     stats = net.train(data_dict['X_train'], data_dict['y_train'], data_dict['X_val'], data_dict['y_val'],
10                     num_iters=3000, batch_size=1000,
11                     learning_rate=lr, learning_rate_decay=0.95,
12                     reg=reg, verbose=False)
13     stat_dict[lr] = stats
14
15 plot_acc_curves(stat_dict)

```

train with learning rate: 0.001
 train with learning rate: 0.01
 train with learning rate: 0.1
 train with learning rate: 1.0
 train with learning rate: 10.0



Tuning. Tuning the hyperparameters and developing intuition for how they affect the final performance is a large part of using Neural Networks, so we want you to get a lot of practice. Below, you should experiment with different values of the various hyperparameters, including hidden layer size, learning rate, number of training epochs, and regularization strength. You might also consider tuning the learning rate decay, but you should be able to get good performance using the default value.

Plots. To guide your hyperparameter search, you might consider making auxiliary plots of training and validation performance as above, or plotting the results arising from different hyperparameter combinations as we did in the Linear Classifier notebook. You should feel free to plot any auxiliary results you need in order to find a good network, but we don't require any particular plots from you.

Approximate results. To get full credit for the assignment, you should achieve a classification accuracy above 50% on the validation set.

(Our best model gets a validation-set accuracy above 58% -- did you beat us?)

```
1 best_net = None # store the best model into this
2
3 ##### TODO: Tune hyperparameters using the validation set. Store your best trained #####
4 # TODO: Tune hyperparameters using the validation set. Store your best trained #
5 # model in best_net. #
6 #
7 # To help debug your network, it may help to use visualizations similar to the #
8 # ones we used above; these visualizations will have significant qualitative #
9 # differences from the ones we saw above for the poorly tuned network. #
10 #
11 # Tweaking hyperparameters by hand can be fun, but you might find it useful to #
12 # write code to sweep through possible combinations of hyperparameters #
13 # automatically like we did on the previous exercises. #
14 #####
15 # Replace "pass" statement with your code
16
17 from scipy.stats import loguniform
18 lrs = [1.2-e for e in loguniform(0.1, 1.1).rvs(size=5)]
19 rgs = loguniform(1e-5, 0.01).rvs(size=5)
20 epochsnum = [1_000, 5_000]
21 hidden_sizes = [128, 512]
22 results = {}
23 best_acc = -float('inf')
24 # Idea: Consider training quicker (1000 epochs, 128 hs) with diff lr, rg, other.
25 for lr in lrs:
26     for rg in rgs:
27         for epochs in epochsnum:
28             for hs in hidden_sizes:
29                 net = TwoLayerNet(3 * 32 * 32, hs, 10, device=data_dict['X train'].device)
```

```

30     stats = net.train(data_dict['X_train'], data_dict['y_train'], data_dict['X_val'], data_dict['y_val'],
31                         num_iters=epochs, batch_size=1000,
32                         learning_rate=lr, learning_rate_decay=0.95,
33                         reg=reg, verbose=False)
34     # Predict on the validation set
35     y_val_pred = net.predict(data_dict['X_val'])
36     val_acc = 100.0 * (y_val_pred == data_dict['y_val']).float().mean().item()
37     print(f'VA: {val_acc:.2f}%. Params: {lr} lr, {rg} rg, {epochs} epochs, {hs} hs.')
38     if best_acc < val_acc:
39         best_acc = val_acc
40         best_net = net
41
42 ######
43 #           END OF YOUR CODE
44 #####

```

--NORMAL--

<Esc>

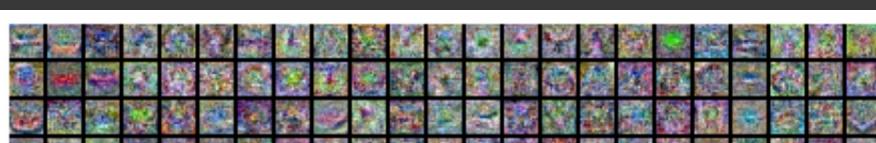
VA: 51.16%. Params: 0.4558211443830802 lr, 0.00014672076203121217 rg, 1000 epochs, 128 hs.
VA: 51.84%. Params: 0.4558211443830802 lr, 0.00014672076203121217 rg, 1000 epochs, 512 hs.
VA: 53.72%. Params: 0.4558211443830802 lr, 0.00014672076203121217 rg, 5000 epochs, 128 hs.
VA: 54.96%. Params: 0.4558211443830802 lr, 0.00014672076203121217 rg, 5000 epochs, 512 hs.
VA: 51.16%. Params: 0.4558211443830802 lr, 0.00028158439338971824 rg, 1000 epochs, 128 hs.
VA: 51.84%. Params: 0.4558211443830802 lr, 0.00028158439338971824 rg, 1000 epochs, 512 hs.
VA: 53.72%. Params: 0.4558211443830802 lr, 0.00028158439338971824 rg, 5000 epochs, 128 hs.
VA: 54.96%. Params: 0.4558211443830802 lr, 0.00028158439338971824 rg, 5000 epochs, 512 hs.
VA: 51.16%. Params: 0.4558211443830802 lr, 1.9774243200150764e-05 rg, 1000 epochs, 128 hs.
VA: 51.84%. Params: 0.4558211443830802 lr, 1.9774243200150764e-05 rg, 1000 epochs, 512 hs.
VA: 53.72%. Params: 0.4558211443830802 lr, 1.9774243200150764e-05 rg, 5000 epochs, 128 hs.
VA: 54.96%. Params: 0.4558211443830802 lr, 1.9774243200150764e-05 rg, 5000 epochs, 512 hs.
VA: 51.16%. Params: 0.4558211443830802 lr, 0.0008050414015595968 rg, 1000 epochs, 128 hs.
VA: 51.84%. Params: 0.4558211443830802 lr, 0.0008050414015595968 rg, 1000 epochs, 512 hs.
VA: 53.72%. Params: 0.4558211443830802 lr, 0.0008050414015595968 rg, 5000 epochs, 128 hs.
VA: 54.96%. Params: 0.4558211443830802 lr, 0.0008050414015595968 rg, 5000 epochs, 512 hs.
VA: 51.16%. Params: 0.4558211443830802 lr, 9.275585988409237e-05 rg, 1000 epochs, 128 hs.
VA: 51.84%. Params: 0.4558211443830802 lr, 9.275585988409237e-05 rg, 1000 epochs, 512 hs.
VA: 53.72%. Params: 0.4558211443830802 lr, 9.275585988409237e-05 rg, 5000 epochs, 128 hs.
VA: 54.96%. Params: 0.4558211443830802 lr, 9.275585988409237e-05 rg, 5000 epochs, 512 hs.
VA: 53.08%. Params: 0.7524673316088673 lr, 0.00014672076203121217 rg, 1000 epochs, 128 hs.
VA: 53.72%. Params: 0.7524673316088673 lr, 0.00014672076203121217 rg, 1000 epochs, 512 hs.
VA: 54.92%. Params: 0.7524673316088673 lr, 0.00014672076203121217 rg, 5000 epochs, 128 hs.
VA: 56.96%. Params: 0.7524673316088673 lr, 0.00014672076203121217 rg, 5000 epochs, 512 hs.
VA: 53.08%. Params: 0.7524673316088673 lr, 0.00028158439338971824 rg, 1000 epochs, 128 hs.
VA: 53.72%. Params: 0.7524673316088673 lr, 0.00028158439338971824 rg, 1000 epochs, 512 hs.
VA: 54.92%. Params: 0.7524673316088673 lr, 0.00028158439338971824 rg, 5000 epochs, 128 hs.
VA: 56.96%. Params: 0.7524673316088673 lr, 0.00028158439338971824 rg, 5000 epochs, 512 hs.
VA: 53.08%. Params: 0.7524673316088673 lr, 1.9774243200150764e-05 rg, 1000 epochs, 128 hs.

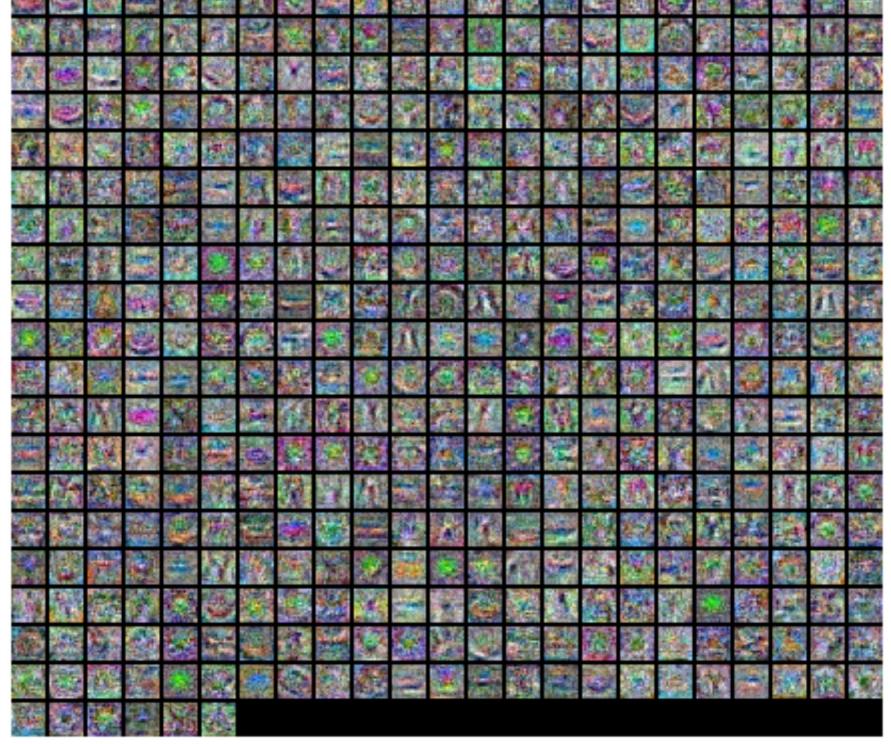
VA: 53.72%. Params: 0.7524673316088673 lr, 1.9774243200150764e-05 rg, 1000 epochs, 512 hs.
VA: 54.92%. Params: 0.7524673316088673 lr, 1.9774243200150764e-05 rg, 5000 epochs, 128 hs.
VA: 56.96%. Params: 0.7524673316088673 lr, 1.9774243200150764e-05 rg, 5000 epochs, 512 hs.
VA: 53.08%. Params: 0.7524673316088673 lr, 0.0008050414015595968 rg, 1000 epochs, 128 hs.
VA: 53.72%. Params: 0.7524673316088673 lr, 0.0008050414015595968 rg, 1000 epochs, 512 hs.
VA: 54.92%. Params: 0.7524673316088673 lr, 0.0008050414015595968 rg, 5000 epochs, 128 hs.
VA: 56.96%. Params: 0.7524673316088673 lr, 0.0008050414015595968 rg, 5000 epochs, 512 hs.
VA: 53.08%. Params: 0.7524673316088673 lr, 9.275585988409237e-05 rg, 1000 epochs, 128 hs.
VA: 53.72%. Params: 0.7524673316088673 lr, 9.275585988409237e-05 rg, 1000 epochs, 512 hs.
VA: 54.92%. Params: 0.7524673316088673 lr, 9.275585988409237e-05 rg, 5000 epochs, 128 hs.
VA: 56.96%. Params: 0.7524673316088673 lr, 9.275585988409237e-05 rg, 5000 epochs, 512 hs.
VA: 53.60%. Params: 1.0867604971202163 lr, 0.00014672076203121217 rg, 1000 epochs, 128 hs.
VA: 55.28%. Params: 1.0867604971202163 lr, 0.00014672076203121217 rg, 1000 epochs, 512 hs.
VA: 54.88%. Params: 1.0867604971202163 lr, 0.00014672076203121217 rg, 5000 epochs, 128 hs.
VA: 56.84%. Params: 1.0867604971202163 lr, 0.00014672076203121217 rg, 5000 epochs, 512 hs.
VA: 53.60%. Params: 1.0867604971202163 lr, 0.00028158439338971824 rg, 1000 epochs, 128 hs.
VA: 55.28%. Params: 1.0867604971202163 lr, 0.00028158439338971824 rg, 1000 epochs, 512 hs.
VA: 54.88%. Params: 1.0867604971202163 lr, 0.00028158439338971824 rg, 5000 epochs, 128 hs.
VA: 56.84%. Params: 1.0867604971202163 lr, 0.00028158439338971824 rg, 5000 epochs, 512 hs.
VA: 53.60%. Params: 1.0867604971202163 lr, 1.9774243200150764e-05 rg, 1000 epochs, 128 hs.
VA: 55.28%. Params: 1.0867604971202163 lr, 1.9774243200150764e-05 rg, 1000 epochs, 512 hs.
VA: 54.88%. Params: 1.0867604971202163 lr, 1.9774243200150764e-05 rg, 5000 epochs, 128 hs.
VA: 56.84%. Params: 1.0867604971202163 lr, 1.9774243200150764e-05 rg, 5000 epochs, 512 hs.
VA: 53.60%. Params: 1.0867604971202163 lr, 0.0008050414015595968 rg, 1000 epochs, 128 hs.
VA: 55.28%. Params: 1.0867604971202163 lr, 0.0008050414015595968 rg, 1000 epochs, 512 hs.
VA: 54.88%. Params: 1.0867604971202163 lr, 0.0008050414015595968 rg, 5000 epochs, 128 hs.
VA: 56.84%. Params: 1.0867604971202163 lr, 0.0008050414015595968 rg, 5000 epochs, 512 hs.
VA: 53.60%. Params: 1.0867604971202163 lr, 9.275585988409237e-05 rg, 1000 epochs, 128 hs.
VA: 55.28%. Params: 1.0867604971202163 lr, 9.275585988409237e-05 rg, 1000 epochs, 512 hs.

```
1 # Check the validation-set accuracy of your best model
2 y_val_preds = best_net.predict(data_dict['X_val'])
3 val_acc = 100 * (y_val_preds == data_dict['y_val']).float().mean().item()
4 print('Best val-set accuracy: %.2f%%' % val_acc)
```

Best val-set accuracy: 57.96%

```
1 # visualize the weights of the best network
2 show_net_weights(best_net)
```





Run on the test set

When you are done experimenting, you should evaluate your final trained network on the test set. To get full credit for the assignment, you should achieve over 50% classification accuracy on the test set.

(Our best model gets 54.1% test-set accuracy -- did you beat us?)

```
1 y_test_preds = best_net.predict(data_dict['X_test'])  
2 test_acc = 100 * (y_test_preds == data_dict['y_test']).float().mean().item()  
3 print('Test accuracy: %.2f%%' % test_acc)
```

Test accuracy: 55.65%

